

A complex, glowing blue neural network diagram is set against a dark blue background. The network consists of numerous circular nodes of varying sizes, interconnected by a dense web of glowing blue lines representing connections. Some nodes are highlighted with a red glow, indicating active neurons or specific points of interest. The overall effect is a futuristic and dynamic representation of a machine learning model.

## ARTIFICIAL NEURAL NETWORK

# Introduction to DL

Unit 4

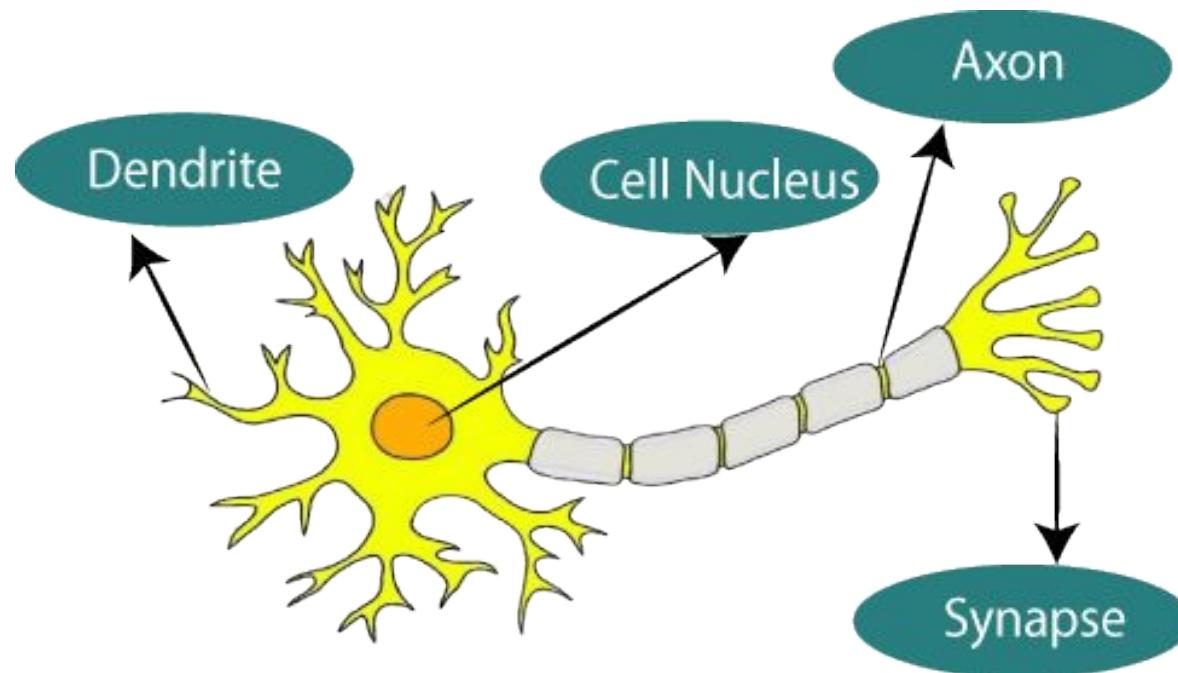
# Syllabus

1. Introduction to Deep Learning, ANN, Machine Learning Vs Deep Learning, Working of Deep Learning;
2. Convolutional Neural Network: Introduction, Components of CNN Architecture, Properties of CNN, Architectures of CNN, Applications of CNN,
3. Recurrent Neural Network: Introduction, Simple RNN, LSTM Implementation, Deep RNN,
4. Autoencoder: Introduction, Features, Types, Applications of Deep Learning. Self-learning Topics: Restricted Boltzmann Machine (RBM).

# DEEP LEARNING AND ARTIFICIAL NEURAL NETWORK

- Deep Learning is a type of machine learning that imitates the way humans gain certain type of knowledge.
- Deep Learning is a subset of Artificial Intelligence also an important element of data science, which includes statistics and predictive modelling.
- Deep Learning networks are mathematical models that are used to represent human brains in order to solve tasks with unstructured data. These mathematical models are created in form of neural network that consists of neurons.
- An Artificial neural network is usually a computational network based on biological neural networks that construct the structure of the human brain.
- Similar to a human brain has neurons interconnected to each other, artificial neural networks also have neurons (nodes) that are linked to each other in various layers of the networks.

# BIOLOGICAL NEURAL NETWORK

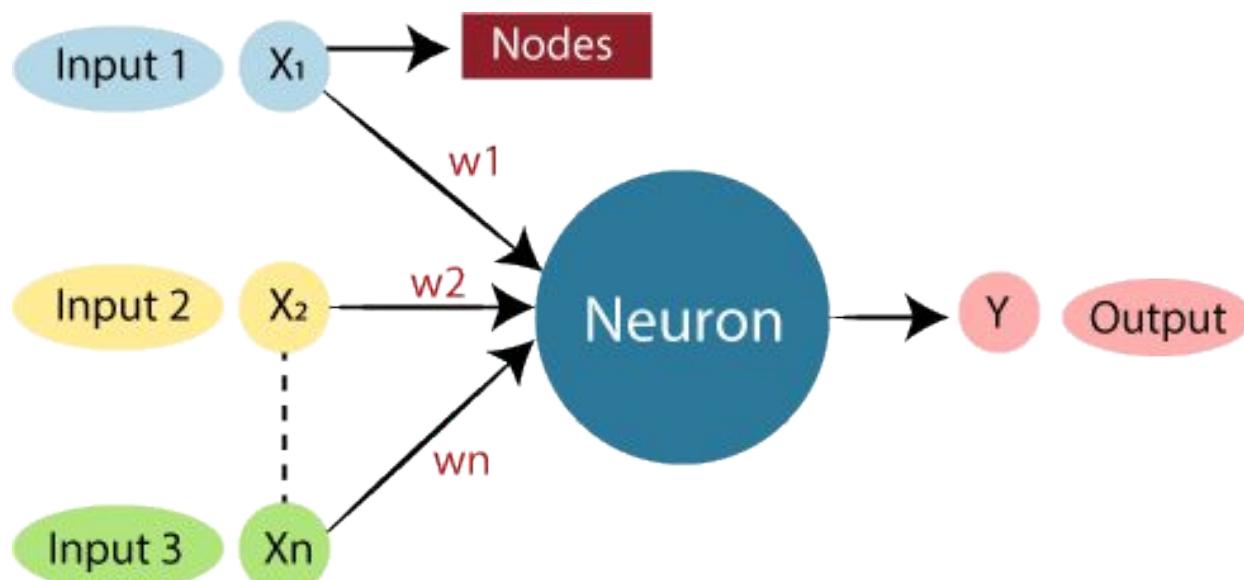


**Soma or cell body:** Here the cell nucleus is located.

**Dendrites:** Where the nerve is connected to the cell body.

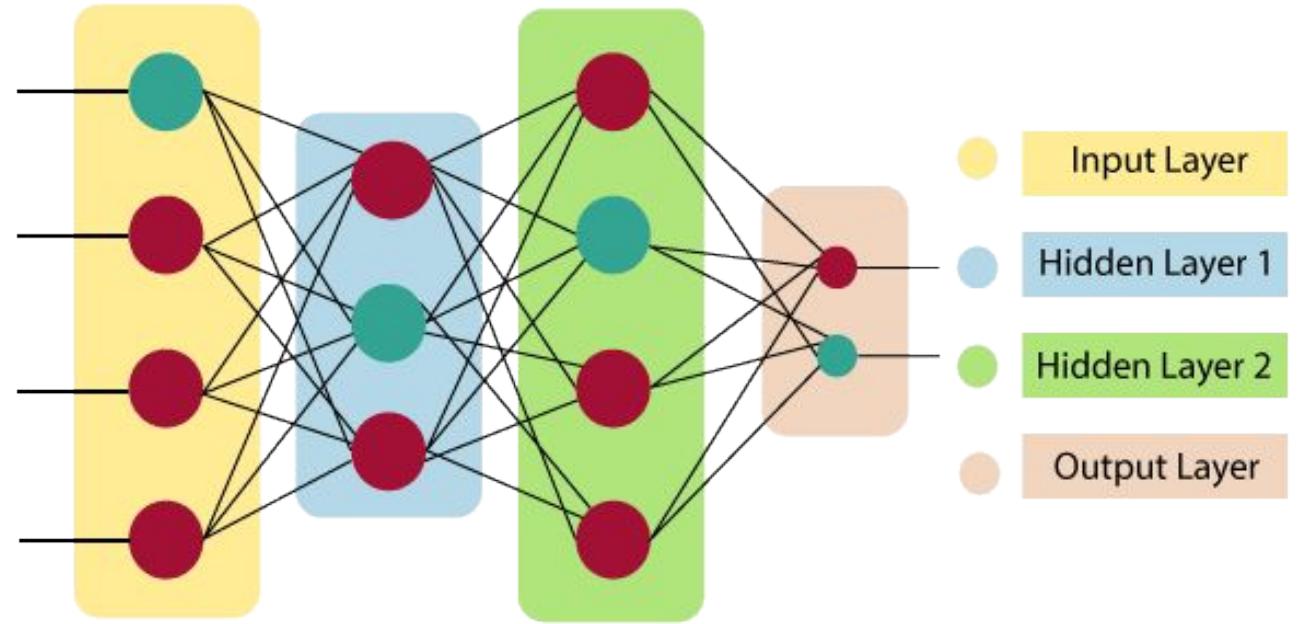
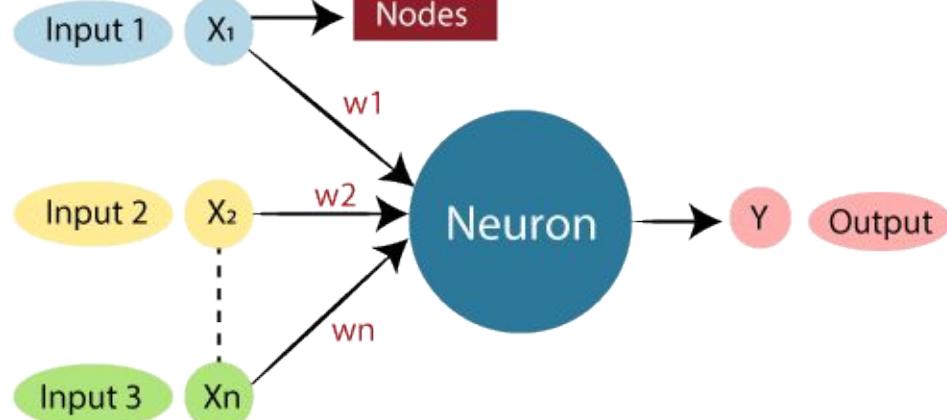
**Axon:** It carries the impulses of the neuron.

# COMPARISON OF ANN AND BNN



Sr. No.	Biological neural network	Artificial neural network
1.	Cell nucleus	Neuron / Nodes
2.	Synapse	Weights or interconnections
3.	Dendrites	Inputs
4.	Axon	Output

# ARCHITECTURE OF ANN



# ARCHITECTURE OF ANN

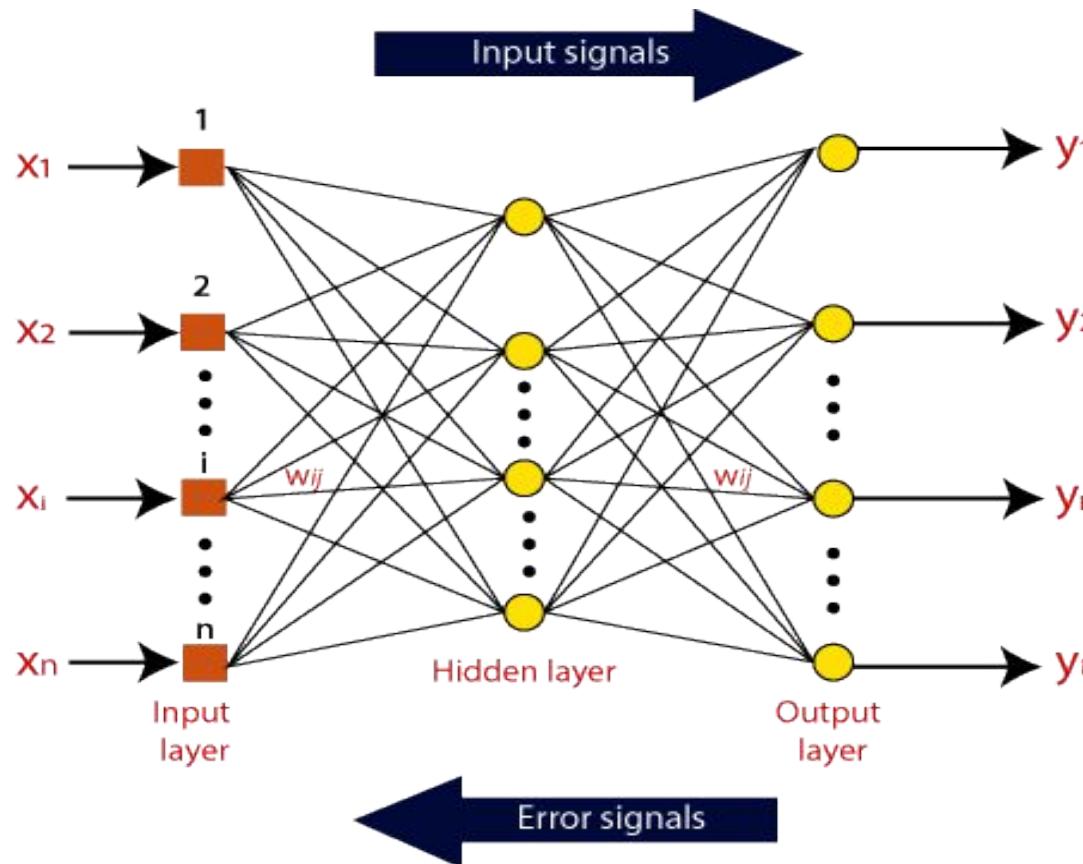
1     **Input Layer:** Input layer accepts inputs in several formats provided by the programmer.

2.     **Hidden Layer:** It performs the calculations to find hidden features and patterns between input and output layers.

3.     **Output Layer:** The input goes through a series of transformations using the hidden layer,

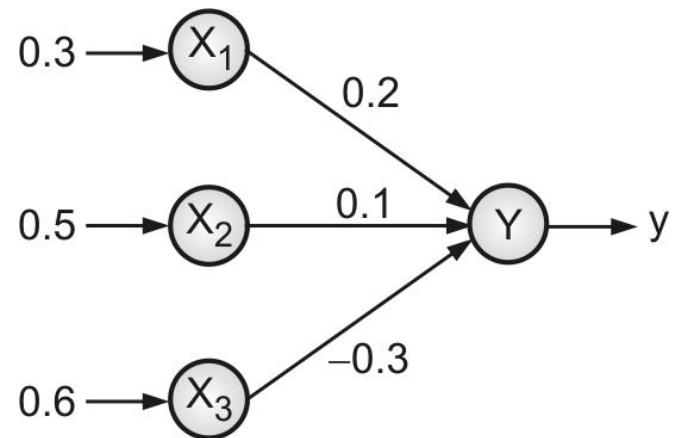
which finally results in output that is conveyed using this layer. An **artificial neural network** takes input and computes the weighted sum of the inputs and includes a **bias**. This computation is represented in the form of a **transfer function**.

# WORKING OF ANN



$$y = x_1 w_1 + x_2 w_2 + \dots + x_n w_n = x_i w_i$$

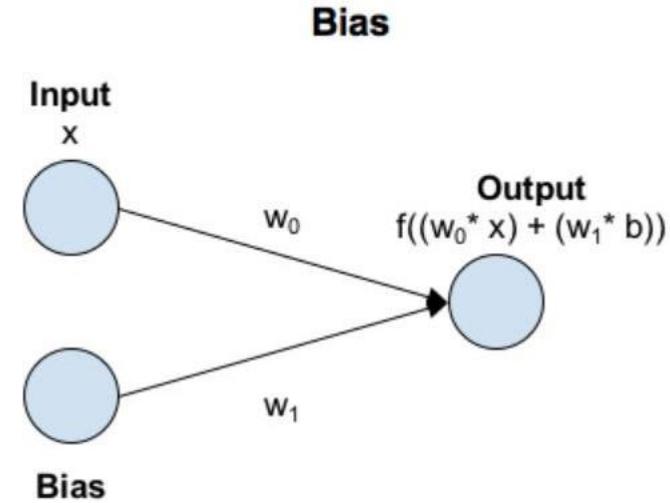
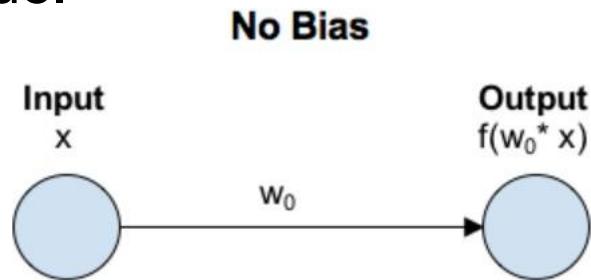
# EXAMPLE 1



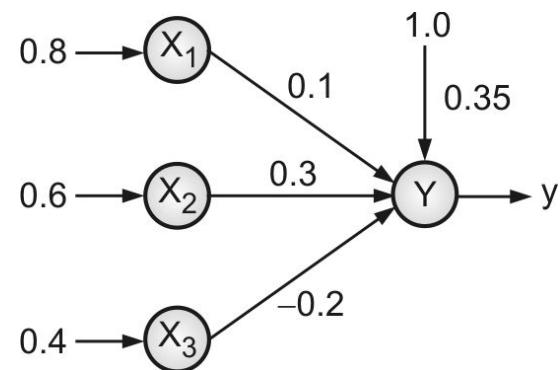
$$\begin{aligned} \text{yin} &= x_1 w_1 + x_2 w_2 + x_3 w_3 = (0.3)(0.2) + (0.5)(0.1) + (0.6)(-0.3) \\ \text{yin} &= 0.06 + 0.05 - 0.18 = -0.07 \end{aligned}$$

# BIAS

- Bias in Neural Networks can be thought of as analogous to the role of a constant in a linear function, whereby the line is effectively transposed by the constant value.



# EXAMPLE 2



- $y_{in} = b + x_i w_i$

- $= b + x_1 w_1 + x_2 w_2 + x_3 w_3 = 0.35 + (0.8)(0.1) + (0.6)(0.3) + (0.4)(-0.2)$

- $= 0.35 + 0.08 + 0.18 - 0.08$

- $= 0.53$

# ACTIVATION/THRESHOLD FUNCTIONS

- To obtain exact output, the activation function is applied over the net input of an ANN.
- Activation Function helps the neural network to use important information while suppressing irrelevant data points.
- The purpose of an activation function is to add non-linearity to the neural network.
- What if a neural network doesn't use activation function?  
Linear Regression Model.....Can't learn for complex tasks

# ACTIVATION FUNCTIONS TYPES

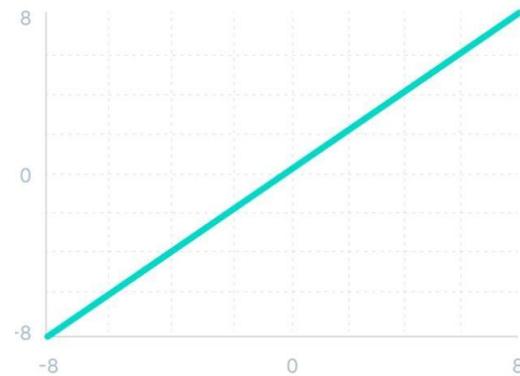
- Binary step function
- Linear Activation Function
- Sigmoid/Logistic Activation Function
- The derivative of the sigmoid Activation Function
- Tanh Function (Hyperbolic Tangent)
- Gradient of the Tanh Activation function
- ReLu Activation Function
- SoftMax

# SOME LINEAR ACTIVATION FUNCTIONS EXPLAINED...

- **Linear function:** It is defined as

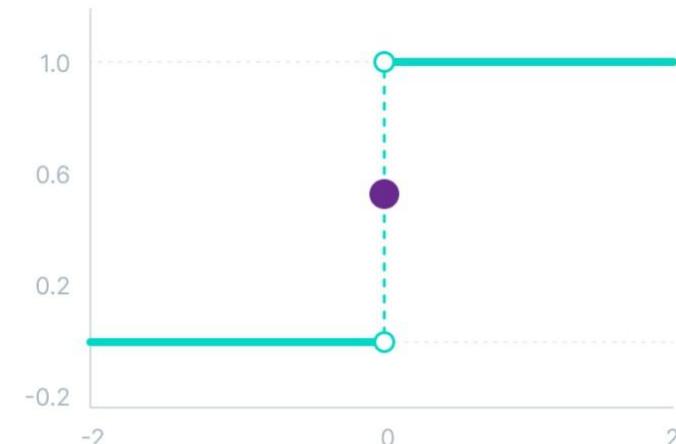
$$f(x) = x \text{ for all } x$$

Here, input=output



- **Binary Step function:** The function is defined as :

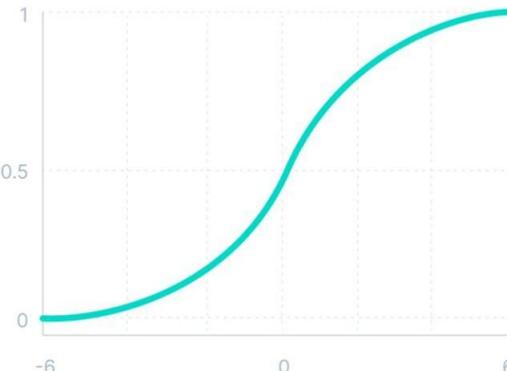
$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$$



# SOME LINEAR ACTIVATION FUNCTIONS EXPLAINED.....

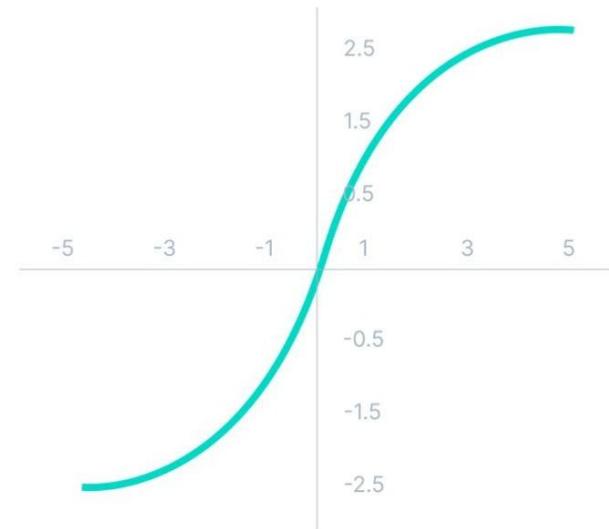
- Sigmoid function:

$$f(x) = \frac{1}{1 + e^{-x}}$$



- Tanh function:

$$f(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$$

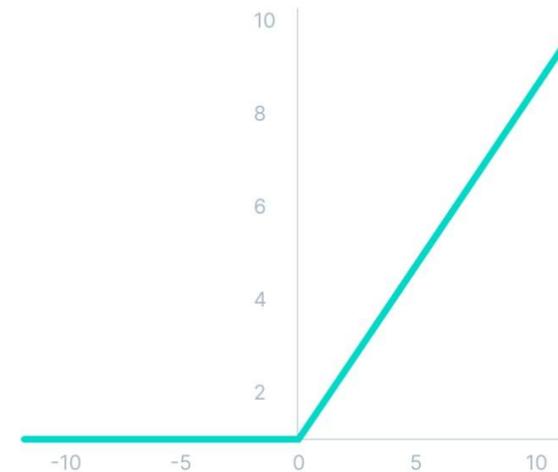


# SOME LINEAR ACTIVATION FUNCTIONS EXPLAINED....

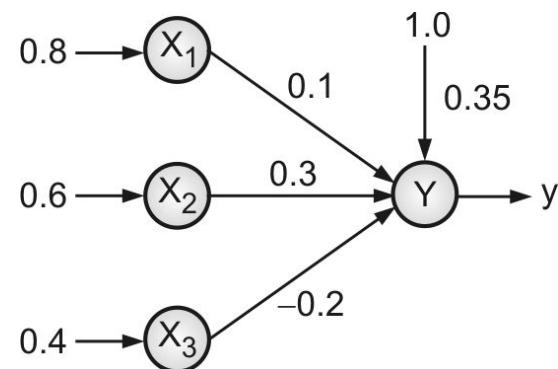
**ReLU:**

**U:**

$$f(x) = \max(0, x)$$



# EXAMPLE 3



- $y_{in} = b + \sum x_i w_i$

- $= b + x_1 w_1 + x_2 w_2 + x_3 w_3 = 0.35 + (0.8)(0.1) + (0.6)(0.3) + (0.4)(-0.2)$

- $= 0.35 + 0.08 + 0.18 - 0.08$

- $= 0.53$

**Binary sigmoidal activation function :**

$$y = f(y_{in}) = \frac{1}{1 + e^{-y_{in}}} = \frac{1}{1 + e^{-0.53}} = 0.625$$

**Bipolar sigmoidal activation function :**

$$y = f(y_{in}) = \frac{2}{1 + e^{-y_{in}}} - 1 = \left( \frac{2}{1 + e^{-0.53}} \right) - 1 = 0.259$$

# VANISHING GRADIENTS

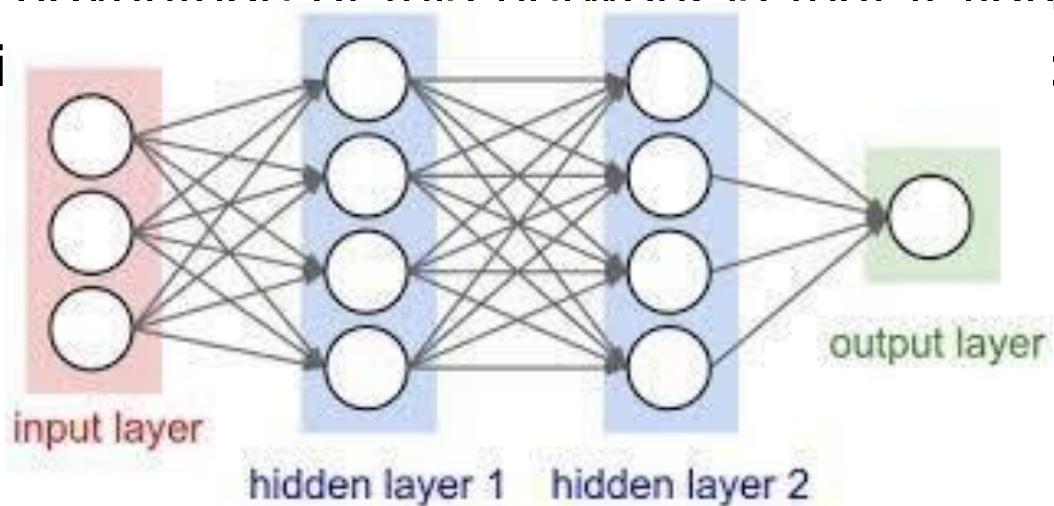
- Like the sigmoid function, certain activation functions squish an ample input space into a small output space between 0 and 1.
- Therefore, a large change in the input of the sigmoid function will cause a small change in the output. Hence, the derivative becomes small.
- However, when more layers are used, it can cause the gradient to be too small for training to work effectively.

# EXPLODING GRADIENTS

- Exploding gradients are problems where significant error gradients accumulate and result in very large updates to neural network model weights during training.
- An unstable network can result when there are exploding gradients, and the learning cannot be completed.
- The values of the weights can also become so large as to overflow and result in something called NaN values.

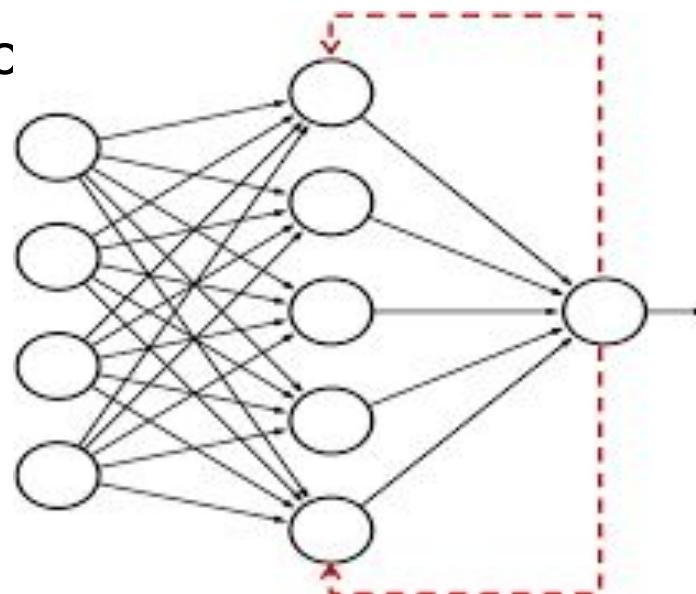
# FEED FORWARD ANN

- A feed-forward network is a basic neural network comprising of an input layer, an output layer, and at least one layer of neurons.
- Through assessment of its output by reviewing its input, the intensity of the network can be noticed based on group behavior of the associated neurons, and the output is decided.
- The primary advantage of this network is that it figures out how to evaluate and recognize patterns without being explicitly programmed to do so.
  - not form a loop.

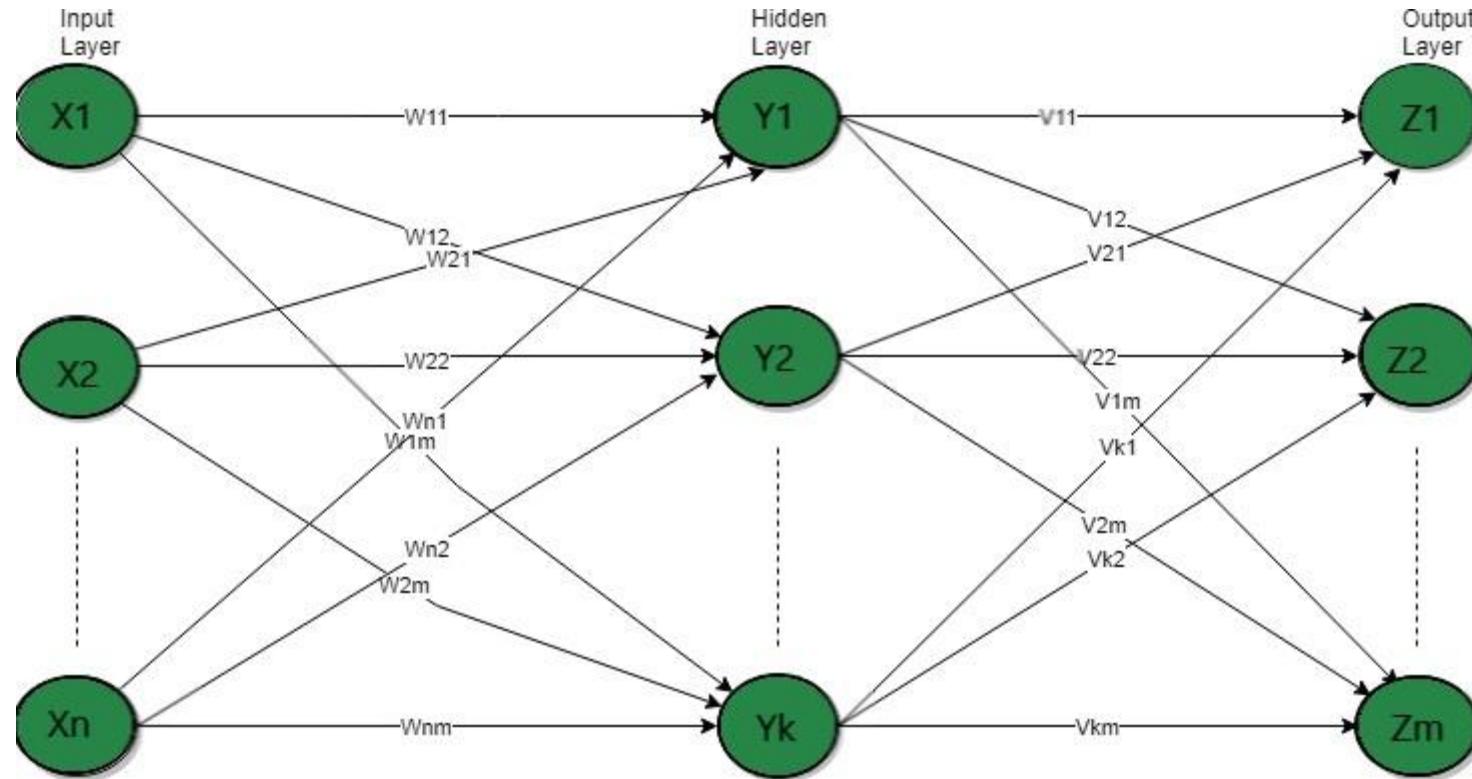


# FEEDBACK ANN

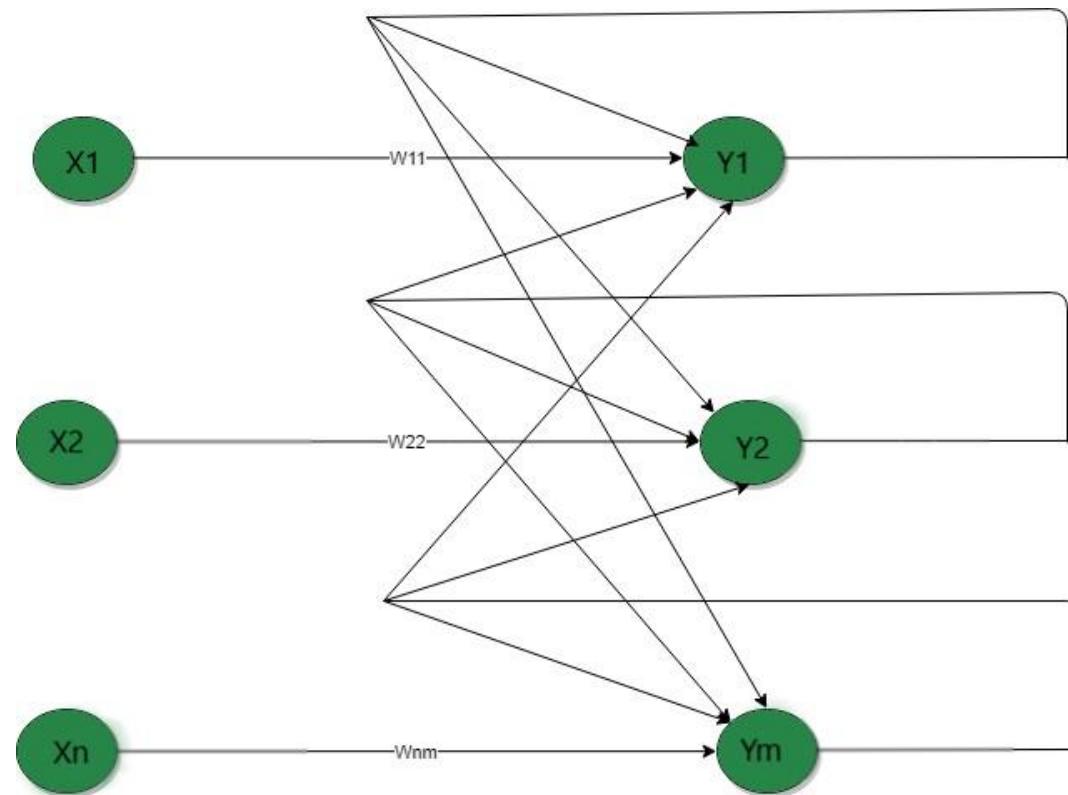
- In this type of ANN, the output returns into the network to accomplish the best-evolved results internally.
- The feedback networks feed information back into itself and are well suited to solve optimization problems error corrections utilize feedback ANNs.



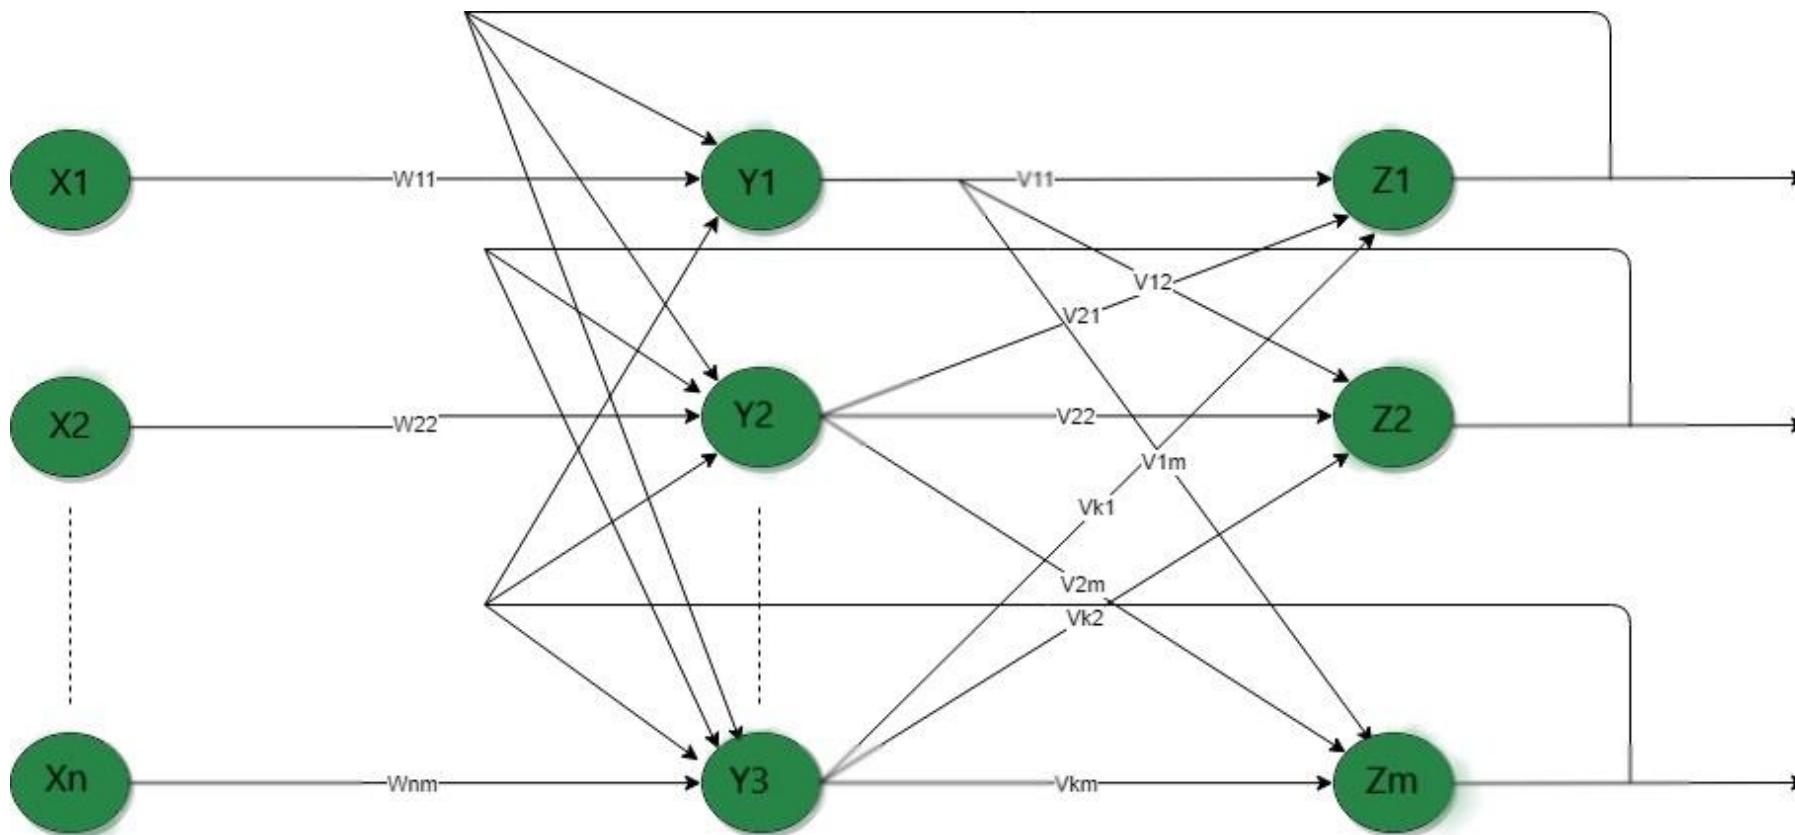
# MULTILAYER FEED-FORWARD NETWORK



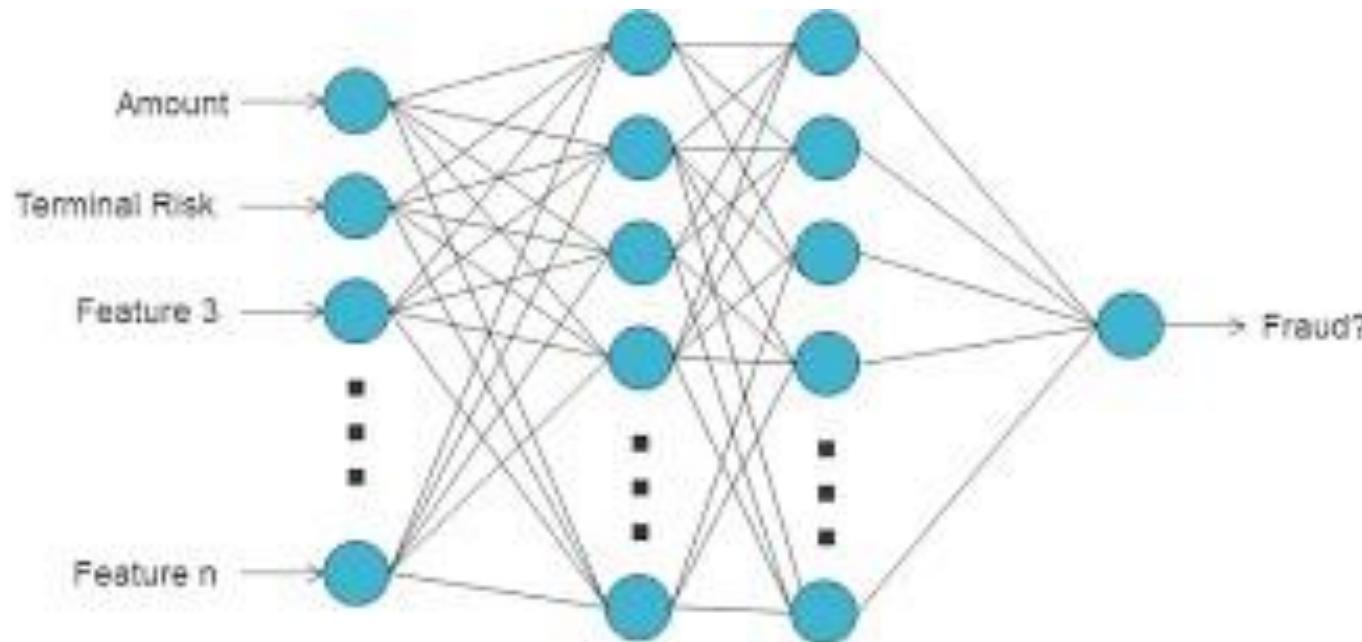
# SINGLE LAYER RECURRENT NETWORK



# MULTI-LAYER RECURRENT NETWORK



# CREDIT CARD FRAUD DETECTION USING ANN



# APPLICATIONS OF NEURAL NETWORKS

- Image Processing – Face Recognition, Signature verification, Handwriting Analysis
- Stock Market Prediction
- Social Media
- Aerospace
- Pattern Recognition
- Defence
- Healthcare
- Optimization/Constraint Satisfactions
- Forecasting and Risk Assessment
- Control system

# Machine Learning Vs Deep Learning

Factors	Deep Learning	Machine Learning
Data Requirement	Requires large data	Can train on lesser data
Accuracy	Provides high accuracy	Gives lesser accuracy
Training Time	Takes longer to train	Takes less time to train
Hardware Dependency	Requires GPU to train properly	Trains on CPU
Hyperparameter Tuning	Can be tuned in various different ways.	Limited tuning capabilities

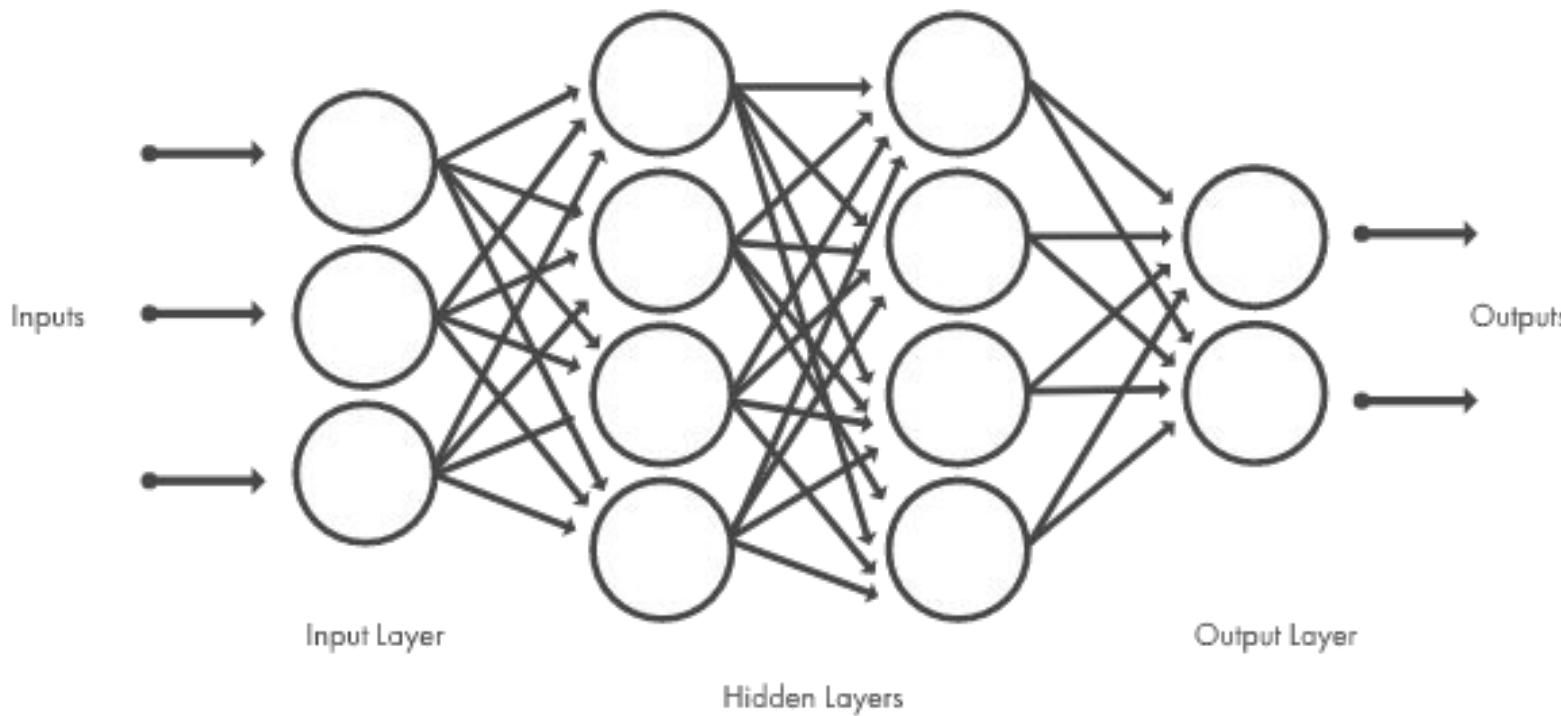
## **Machine Learning**

- Small data set with less feature
  - Feature Extraction
  - Feature Selection
- Manual feature engineering is tough & time consuming

## **Artificial Neural Network**

- Auto Feature Engineering
- Activation function , bias, weight can be added
- Capable of learning any non linear function hence referred as universal function approximation

# Working of Deep Learning



# Working of Deep Learning

- Most deep learning methods use **neural network** architectures, which is why deep learning models are often referred to as **deep neural networks**.
- The term “deep” usually refers to the number of hidden layers in the neural network. **Traditional neural networks** only contain 2-3 hidden layers, while deep networks can have as many as 150.
- Neurons are grouped into three different types of layers:
  - a) Input layer
  - b) Hidden layer
  - c) Output layer
- Application
  - Image recognition
  - Speech Recognition
  - Huge Input Parameter

# Deep Learning Layers

- 1. Input Layer: It receives the input data from the observation
- 2. Hidden Layer: It performs mathematical computations on input data. To decide the number of hidden layers and the number of neurons in each layer is challenging. It forms the hierarchy concepts from the learning. In the hierarchy, each level grasps to transform the input data into a more and more abstract and composite representation.
- Output Layer:It gives the desired result.

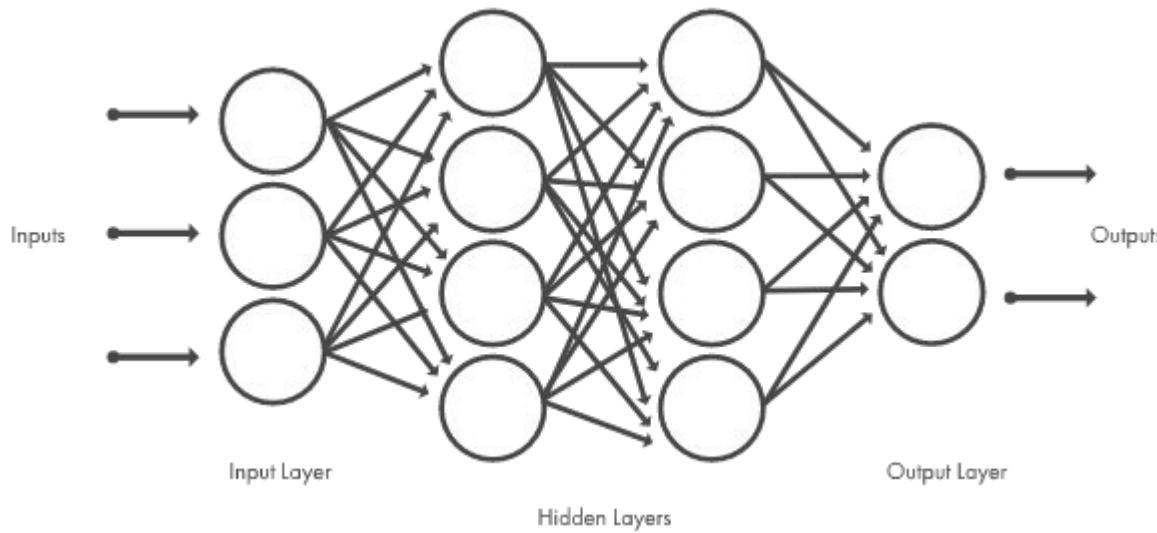
# Applications that limits the ANN but RNN can do

- Time series analysis or use case
- Share market
- Profit forecast for quarters
- Language Translation
- Image caption
- NLP use case
- Can't handle sequential data
- Does not have memory / considers current input

- <https://data-flair.training/blogs/deep-learning-vs-machine-learning/>
- <https://www.mathworks.com/discovery/deep-learning.html>

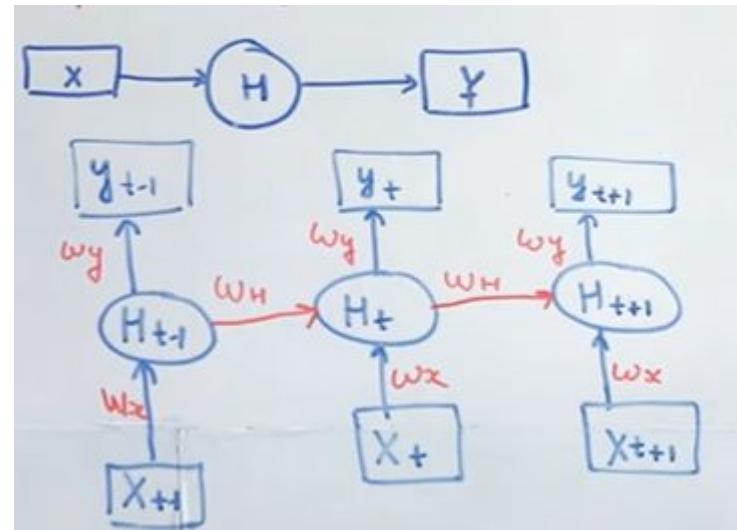


# Recurrent Neural Network: Introduction, Simple RNN

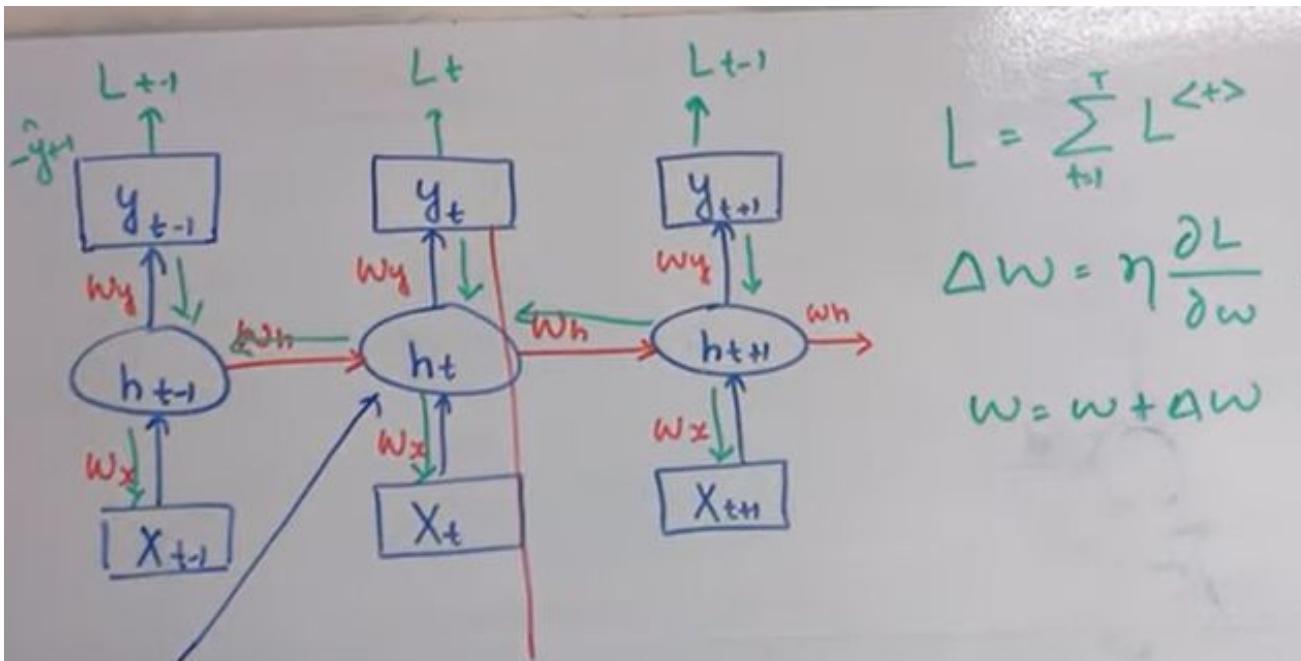


# RNN

- RNNs are a powerful and robust type of neural network, and belong to the most promising algorithms in use because it is the only one with an internal memory.
- NNs and feed-forward neural networks get their names from the way they channel information.
- Backpropagation is used for calculating the gradient of an error function with respect to a neural network's weights. The algorithm works its way backwards through the various layers of gradients to find the partial derivative of the errors with respect to the weights. Backprop then uses these weights to decrease error margins when training.
- Applications
  - Image Recognition
  - Speech Recognition
- **TYPES OF RNNS**
  - One to One
  - One to Many
  - Many to One
  - Many to Many



# RNN



$$L = \sum_{t=1}^T L^{<+>}$$

$$\Delta w = \eta \frac{\partial L}{\partial w}$$

$$w = w + \Delta w$$

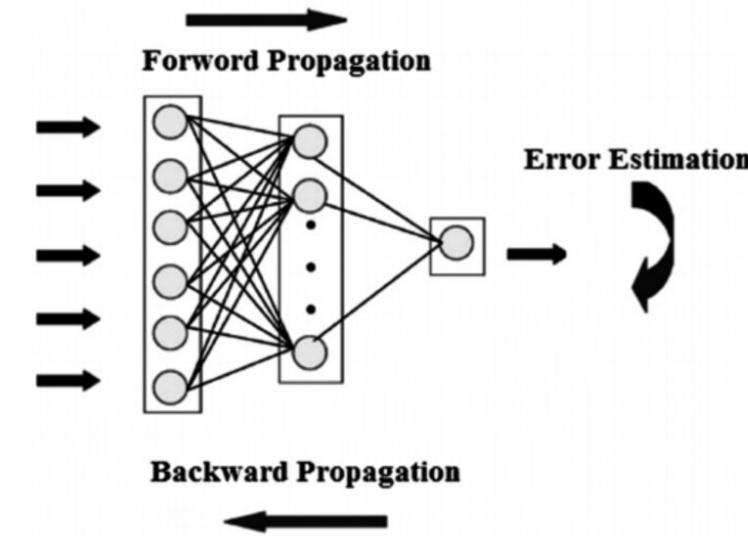
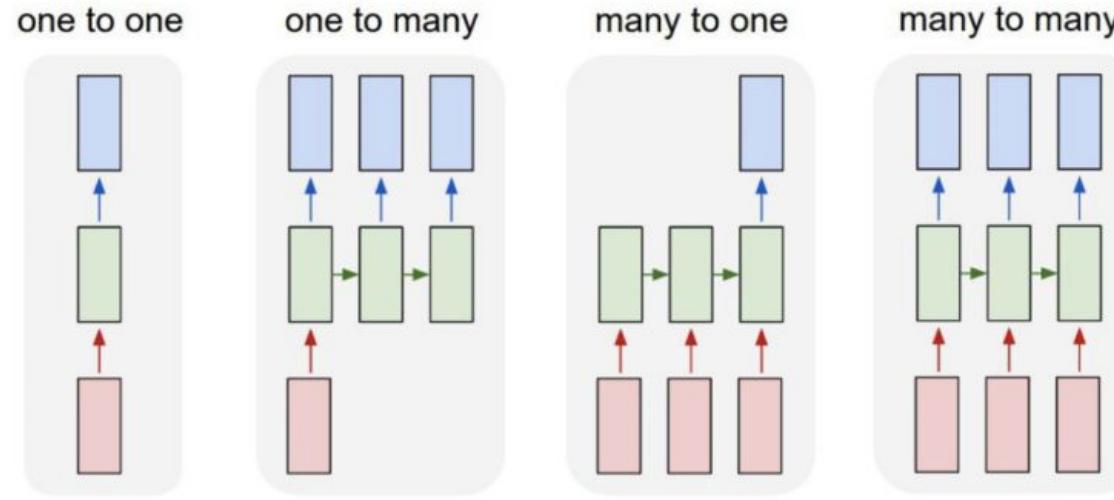
$$z_h^{<+>} = [x_t w_x + h_{t-1} w_h] + b_h$$

$$h_t = \sigma(z_h^{<+>})$$

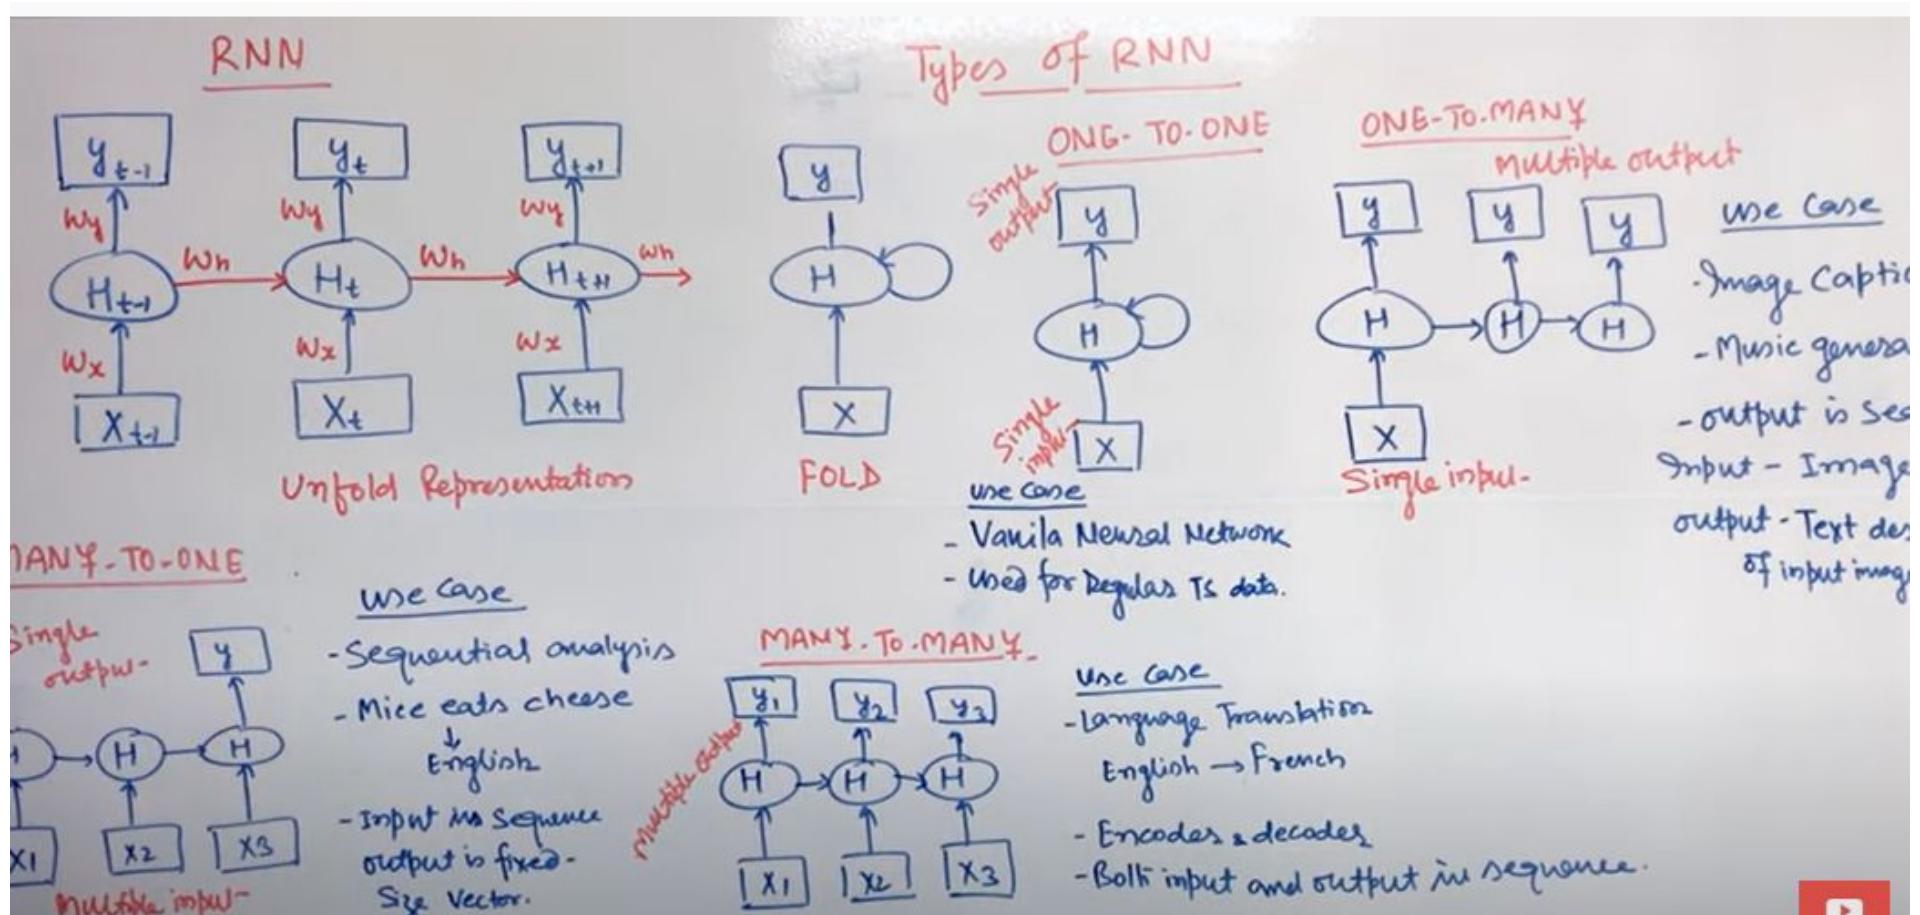
$$z_y^{<+>} = [h_t w_y] + b_y$$

$$y_t = \sigma(z_y^{<+>})$$

# Types of Recurrent NN & Backpropagation NN



# Types of RNN with application



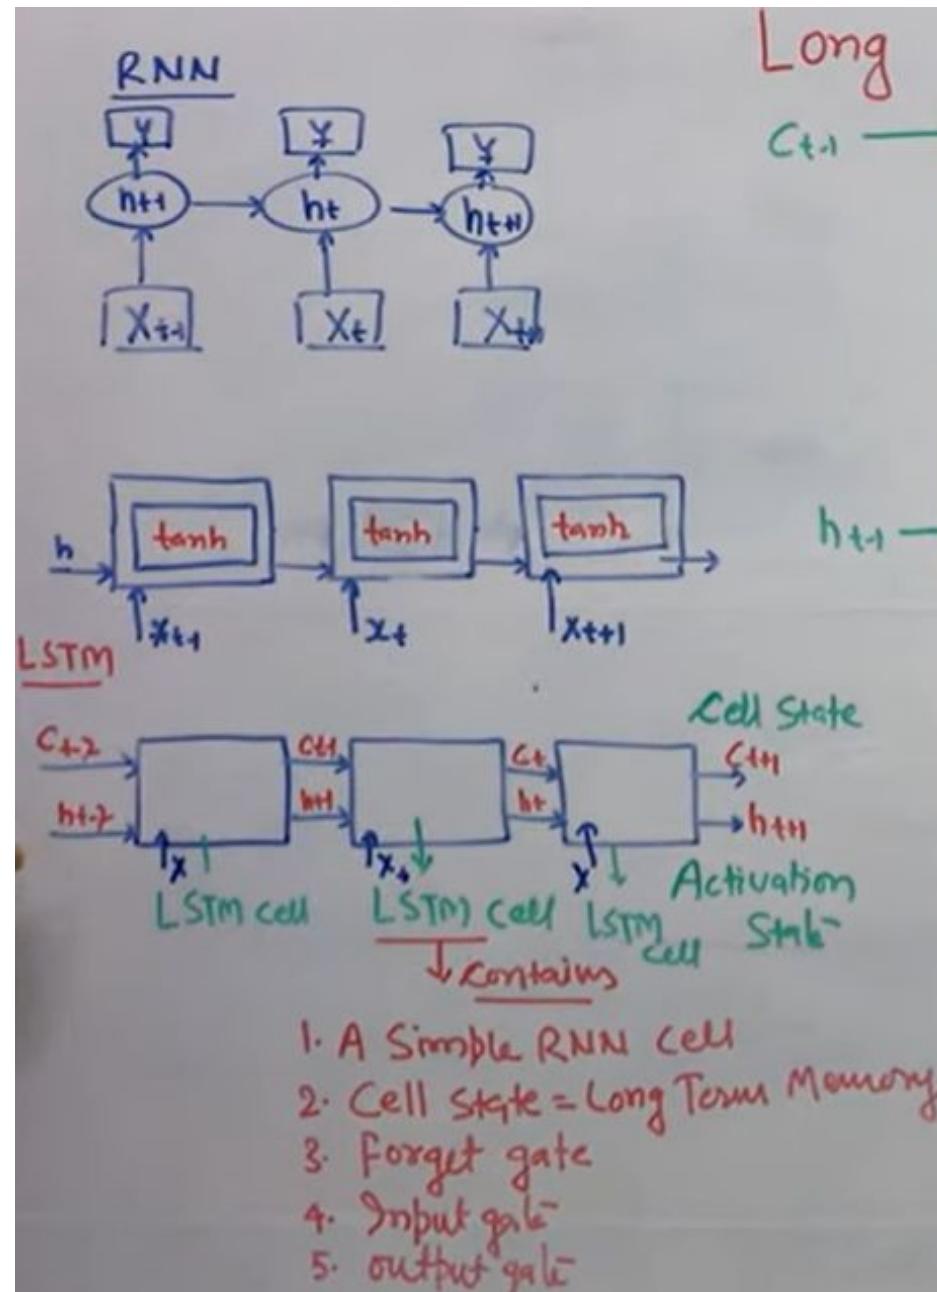
# LSTM Implementation

- Long short-term memory (LSTM) networks are an extension of RNN that extend the memory. LSTM are used as the building blocks for the layers of a RNN. LSTMs assign data “weights” which helps RNNs to either let new information in, forget information or give it importance enough to impact the output.
- The units of an LSTM are used as building units for the layers of a RNN, often called an LSTM network.
- This memory can be seen as a gated cell, with gated meaning the cell decides whether or not to store or delete information

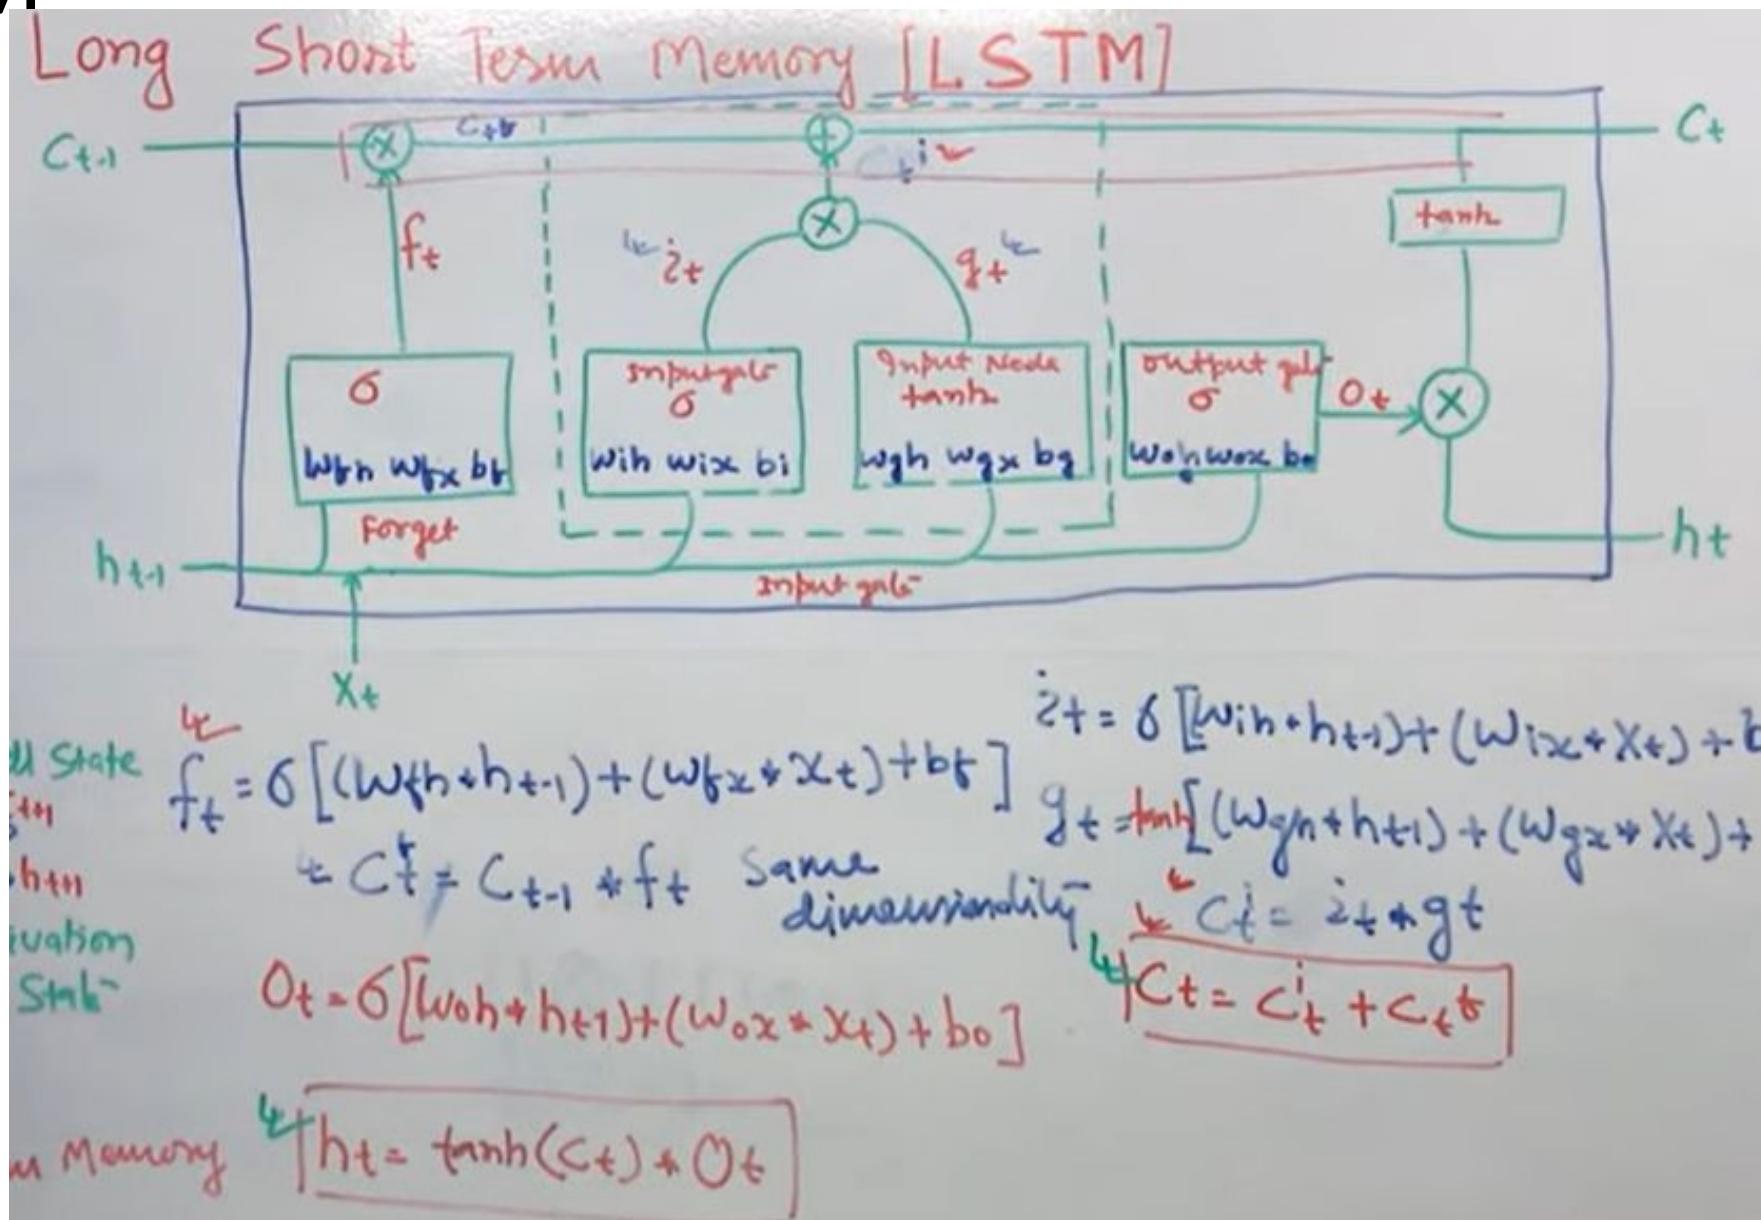
# LSTM

- The units of an LSTM are used as building units for the layers of a RNN, often called an LSTM network.
- LSTMs enable RNNs to remember inputs over a long period of time. This is because LSTMs contain information in a memory, much like the memory of a computer. The LSTM can read, write and delete information from its memory.
- This memory can be seen as a gated cell, with gated meaning the cell decides whether or not to store or delete information.
- This simply means that it learns over time what information is important and what is not.
- In a long short-term memory cell you have three gates: input, forget and output gate. These gates determine whether or not to let new input in (input gate), delete the information because it isn't important (forget gate), or let it impact the output at the current timestep (output gate)
- The problematic issue of vanishing gradients is solved through LSTM because it keeps the gradients steep enough, which keeps the training relatively short and the accuracy high.

# LSTM



# LSTM



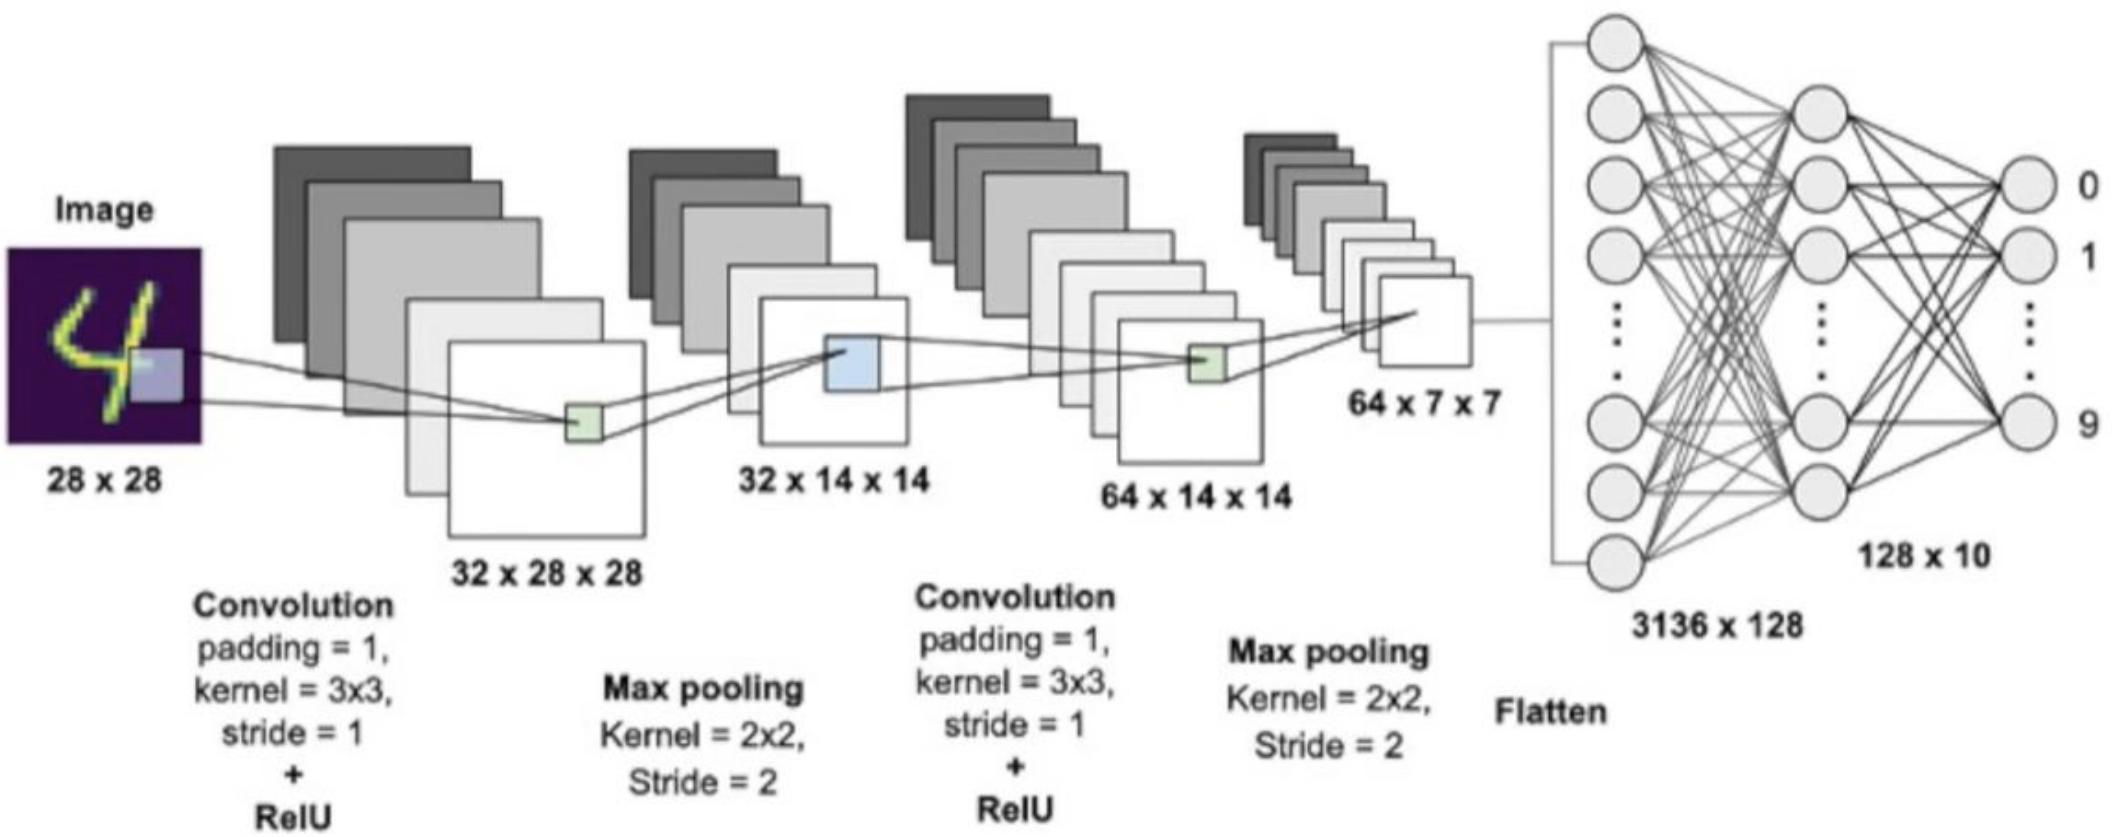
# Reference

- <https://www.youtube.com/watch?v=Opj2ATOiYCw&t=416s>

# Deep RNN

- Fortunately, by learning from the multitude of already existing compositions, a neural network does have the ability to generate a new kind of music. Generating music using computers is an exciting application of what a neural network can do.
- In DNN inputs are independent of each other
- All inputs are given at same time to the network
- Use Different parameters at any layers of the network
- Application but not useful for sequential information
- Is to use sequential information for predicting next word in sentence
  - Handwriting recognition
  - speech recognition
  - machine translation

# Convolutional Neural Network: Introduction



# Advantage of CNN



Image size = 1920 x 1080 X 3

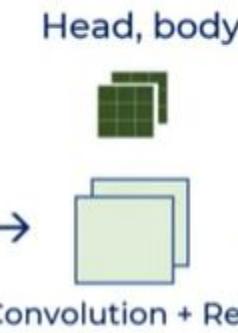
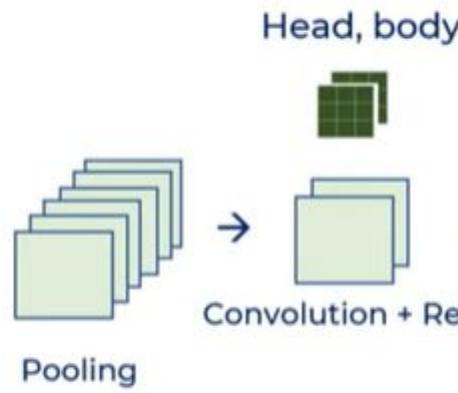
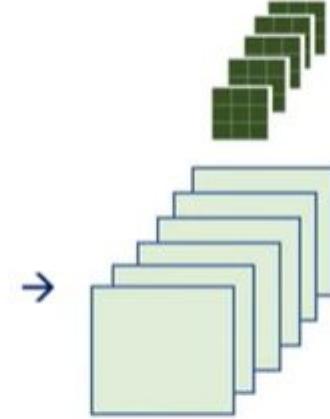
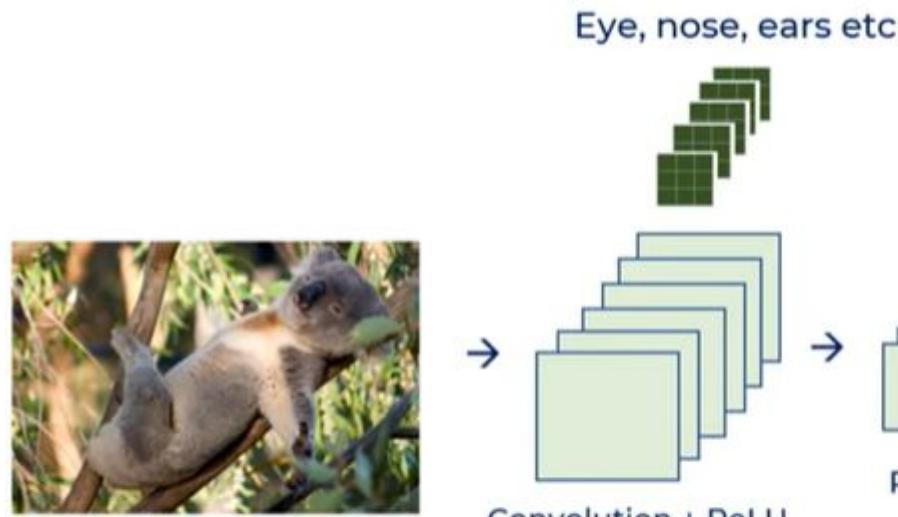
First layer neurons =  $1920 \times 1080 \times 3 \sim 6 \text{ million}$



Hidden layer neurons = Let's say you keep it  $\sim 4 \text{ million}$

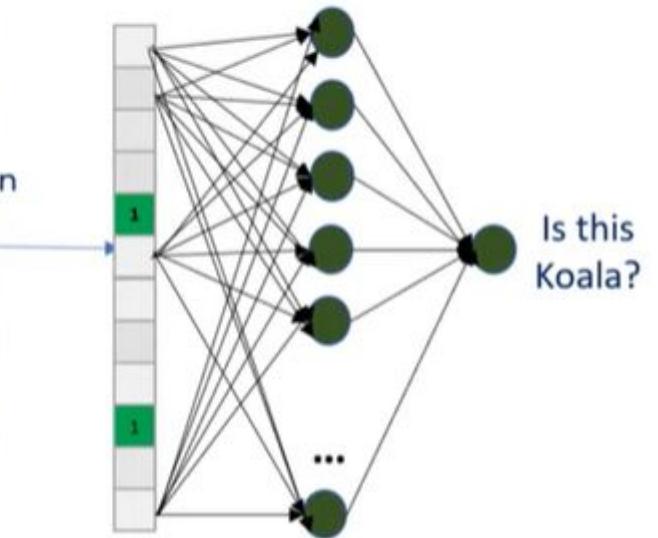
Weights between input and hidden layer =  $6 \text{ mil} * 4 \text{ mil}$   
 $= 24 \text{ million}$

## Feature Extraction



→ Convolution + ReLU  
→ Pooling

flattened



## Classification

Rotation

A large, thin, black cursive-style letter 'g'.

Thickness

A large, thick, black cursive-style letter 'g'.

## CNN by itself doesn't take care of rotation and scale

- You need to have rotated, scaled samples in training dataset
- If you don't have such samples than use **data augmentation** methods to generate new rotated/scaled samples from existing training samples

# Input image to recognize handwriting from images

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1



# hardcoded

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1



# Location shifted

1	1	1	-1	-1
1	-1	1	-1	-1
1	1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1
1	-1	-1	-1	-1



1	1	1	-1	-1
1	-1	1	-1	-1
1	1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1
1	-1	-1	-1	-1



-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	1	-1	-1	-1

# Variation 1

-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1



-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1
-1	-1	1	-1	-1



-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

# Representation of numbers

## Variation 2

-1	-1	1	1	-1
-1	1	-1	1	-1
-1	1	-1	1	-1
-1	-1	1	1	-1
-1	-1	-1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1



-1	-1	1	1	-1
-1	1	-1	1	-1
-1	1	-1	1	-1
-1	-1	1	1	-1
-1	-1	-1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1



-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	1	-1	-1	-1
-1	1	-1	-1	-1

To handle **variety** in digits we can use simple artificial neural network (ANN)



## Disadvantages of using ANN for image classification

1. Too much computation
2. Treats local pixels same as pixels far apart
3. Sensitive to location of an object in an image

**HOW CAN WE MAKE COMPUTERS  
RECOGNIZE THESE TINY FEATURES?**



-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

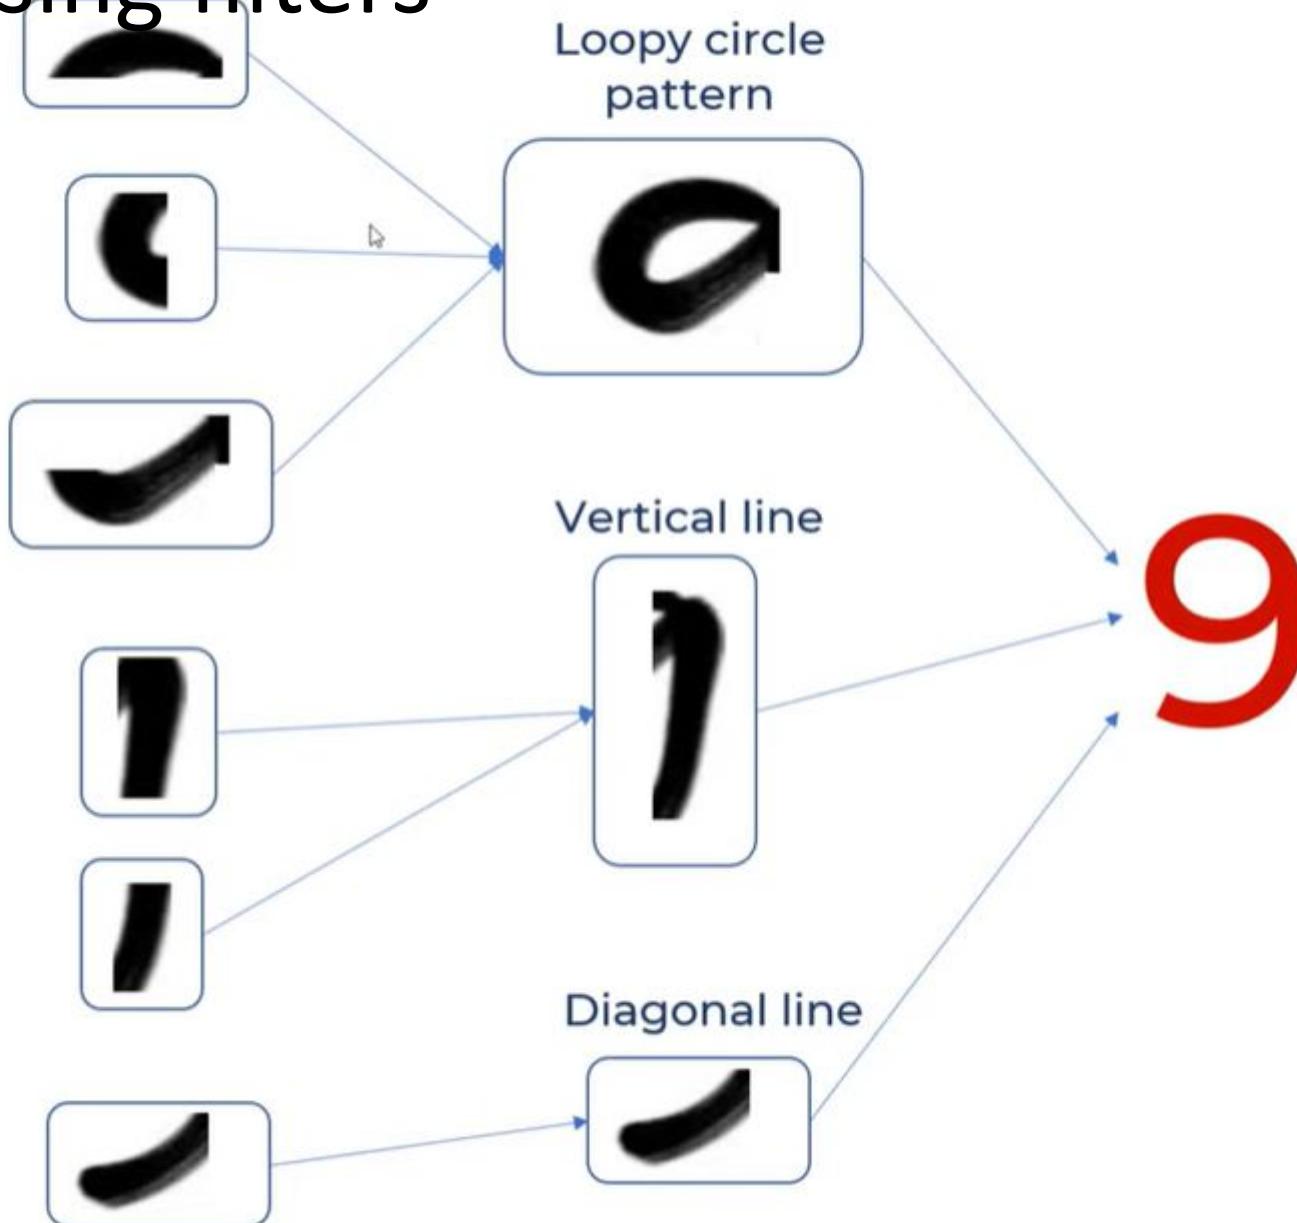
Loopy pattern  
filter

Vertical line  
filter

Diagonal line  
filter

# Edge detection using filters

A large, black, handwritten-style digit 'g' is displayed on the left side of the diagram.



$$-1+1+1-1-1-1+1+1 = -1 \rightarrow -1/9 = -0.11$$

-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	<b>1</b>	-1	<b>1</b>	-1
-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	<b>1</b>	-1	-1
-1	<b>1</b>	-1	-1	-1

\*

1	1	1
1	-1	1
1	1	1

-0.11		

-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	<b>1</b>	-1	<b>1</b>	-1
-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	<b>1</b>	-1	-1
-1	<b>1</b>	-1	-1	-1

\*

1	1	1
1	-1	1
1	1	1

-0.11	<b>1</b>	

-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	<b>1</b>	-1	<b>1</b>	-1
-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	<b>1</b>	-1	-1
-1	<b>1</b>	-1	-1	-1

\*

1	1	1
1	-1	1
1	1	1

-0.11	1	-0.11

-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	<b>1</b>	-1	<b>1</b>	-1
-1	<b>1</b>	<b>1</b>	<b>1</b>	-1
-1	-1	-1	<b>1</b>	-1
-1	-1	<b>-1</b>	<b>1</b>	-1
-1	-1	<b>1</b>	-1	-1
-1	<b>1</b>	<b>-1</b>	-1	-1

\*

1	1	1
1	-1	1
1	1	1

-0.11	<b>1</b>	-0.11
-0.55	0.11	-0.33
-0.33	0.33	-0.33
-0.22	-0.11	-0.22
-0.33	-0.33	-0.33

Feature Map

Loopy pattern detector

$$g * \begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \begin{array}{|c|c|c|} \hline & & 1 \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

6 \*  $\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & -1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$  =



$$8 * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & -1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline & & 1 \\ \hline & & \\ \hline & & \\ \hline \end{array}$$

$$96 * \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & -1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \begin{array}{|c|c|c|} \hline 1 & & \\ \hline & & \\ \hline & & \\ \hline \end{array}$$



$$\text{eye detector} * \begin{array}{|c|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} = \boxed{\begin{array}{|c|c|}\hline 1 & 1 \\ \hline \end{array}}$$

Location invariant: It can detect eyes in any location of the image



$$\text{eye detector} * \begin{array}{|c|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} = \boxed{\begin{array}{|c|c|}\hline 1 & 1 \\ \hline \end{array}}$$



$$\text{eye detector} * \begin{array}{|c|c|c|}\hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array} = \boxed{\begin{array}{|c|c|c|c|c|}\hline 1 & & & 1 & 1 \\ \hline & & & 1 & 1 \\ \hline \end{array}}$$



Loopy pattern  
detector

1	1	1
1	-1	1
1	1	1

\*

=

		1	

Vertical line  
detector

-1	1	-1
-1	1	-1
-1	1	-1

\*

		1	



Diagonal line  
detector

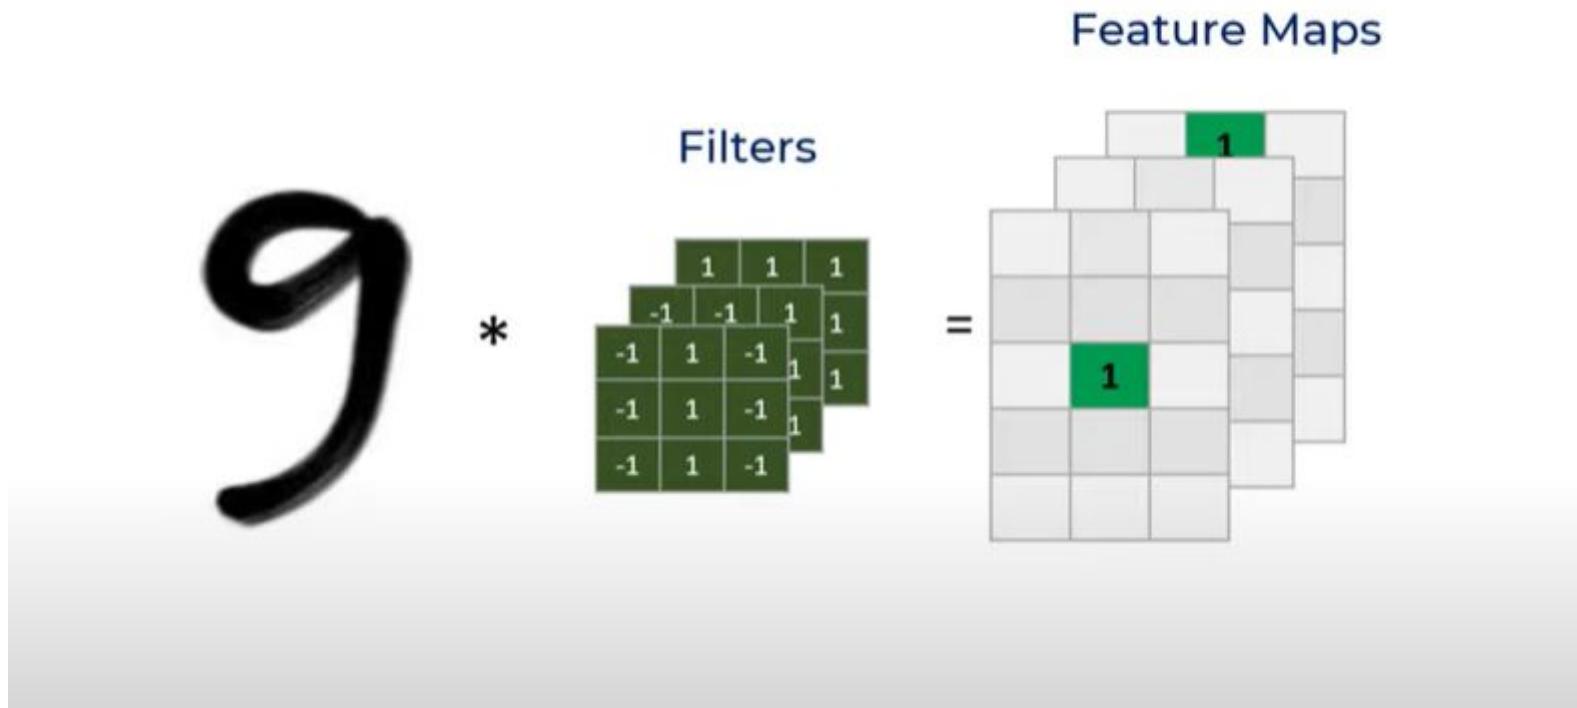
-1	-1	1
-1	1	-1
1	-1	-1

\*

=

		1	

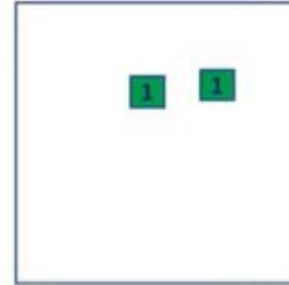
# 3 Feature Maps – volume



# Filter detecting the parts of Koala



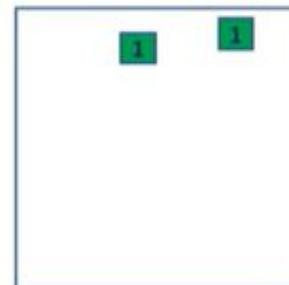
$$\text{eye} * \begin{matrix} & & \\ & & \\ & & \end{matrix} =$$



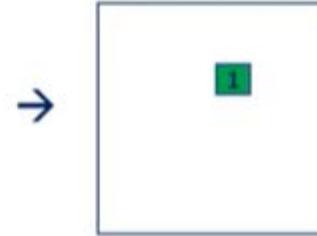
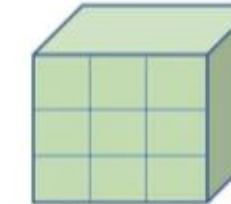
$$\text{nose} * \begin{matrix} & & \\ & & \\ & & \end{matrix} =$$



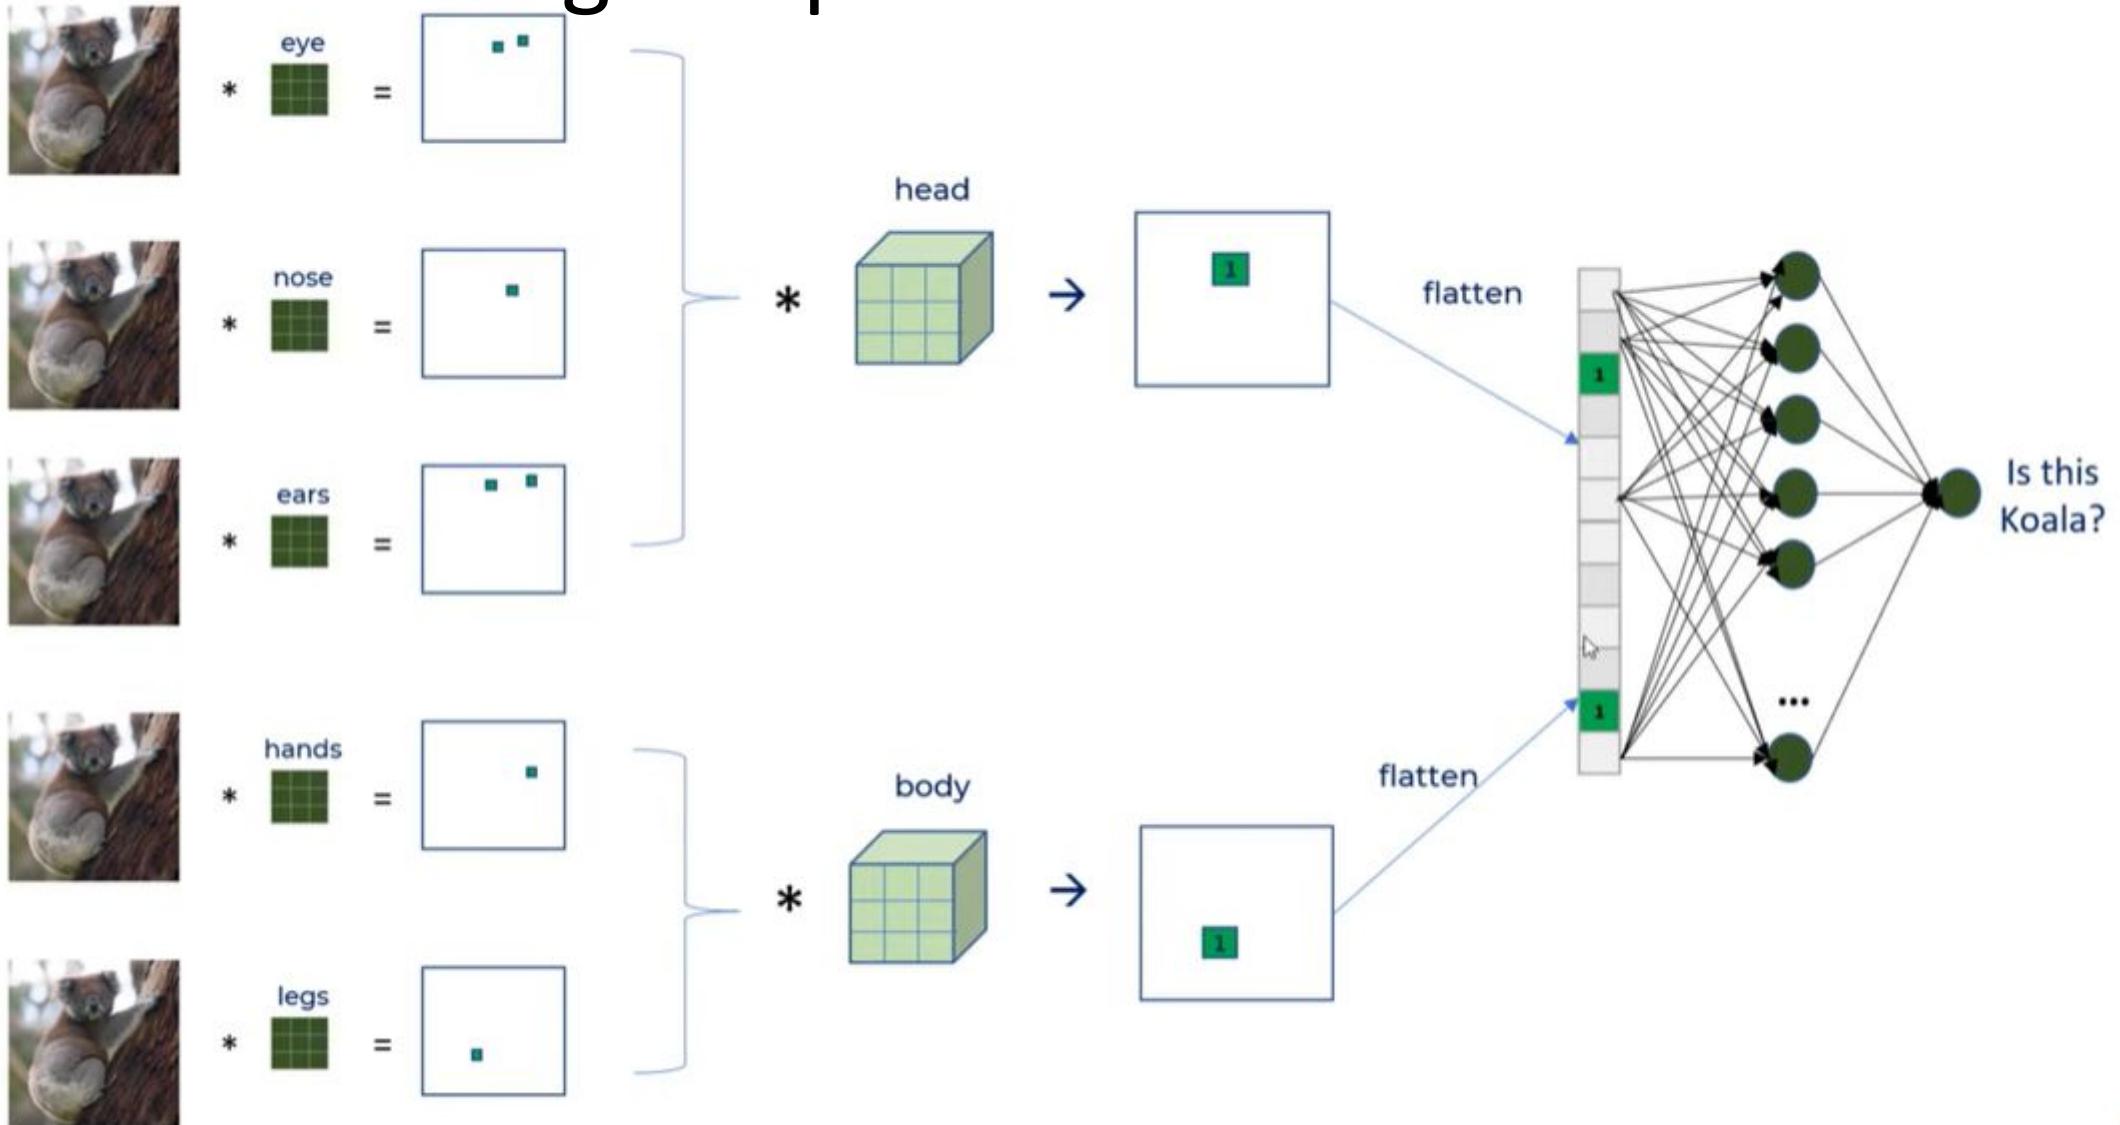
$$\text{ears} * \begin{matrix} & & \\ & & \\ & & \end{matrix} =$$



Filter for head



# Filter detecting the parts of Koala



# Activation Function – Relu

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

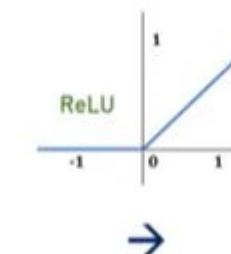
\*

Loopy pattern  
filter

1	1	1
1	-1	1
1	1	1



-0.11	1	-0.11
-0.55	0.11	-0.33
-0.33	0.33	-0.33
-0.22	-0.11	-0.22
-0.33	-0.33	-0.33



0	1	0
0	0.11	0
0	0.33	0
0	0	0
0	0	0

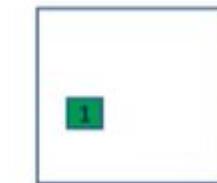


$$\begin{aligned} * \text{ eye} &= \boxed{\text{eye tensor}} \\ * \text{ nose} &= \boxed{\text{nose tensor}} \\ * \text{ ears} &= \boxed{\text{ears tensor}} \\ * \text{ hands} &= \boxed{\text{hands tensor}} \\ * \text{ legs} &= \boxed{\text{legs tensor}} \end{aligned}$$

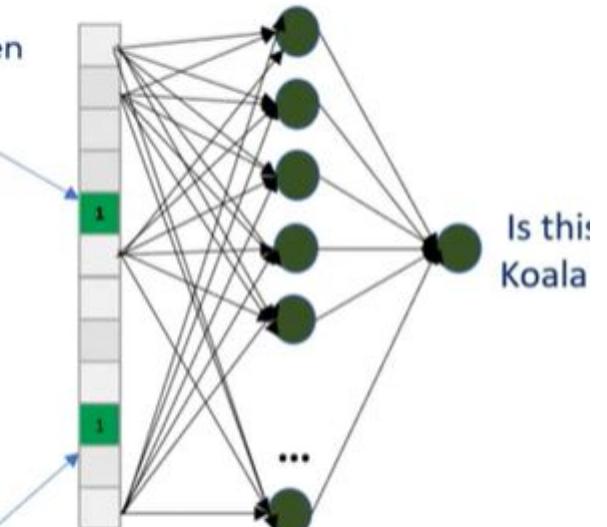
$$\left. \begin{aligned} & \text{head} \\ & \text{body} \end{aligned} \right\} \rightarrow$$



flatten



flatten



Stride = 2

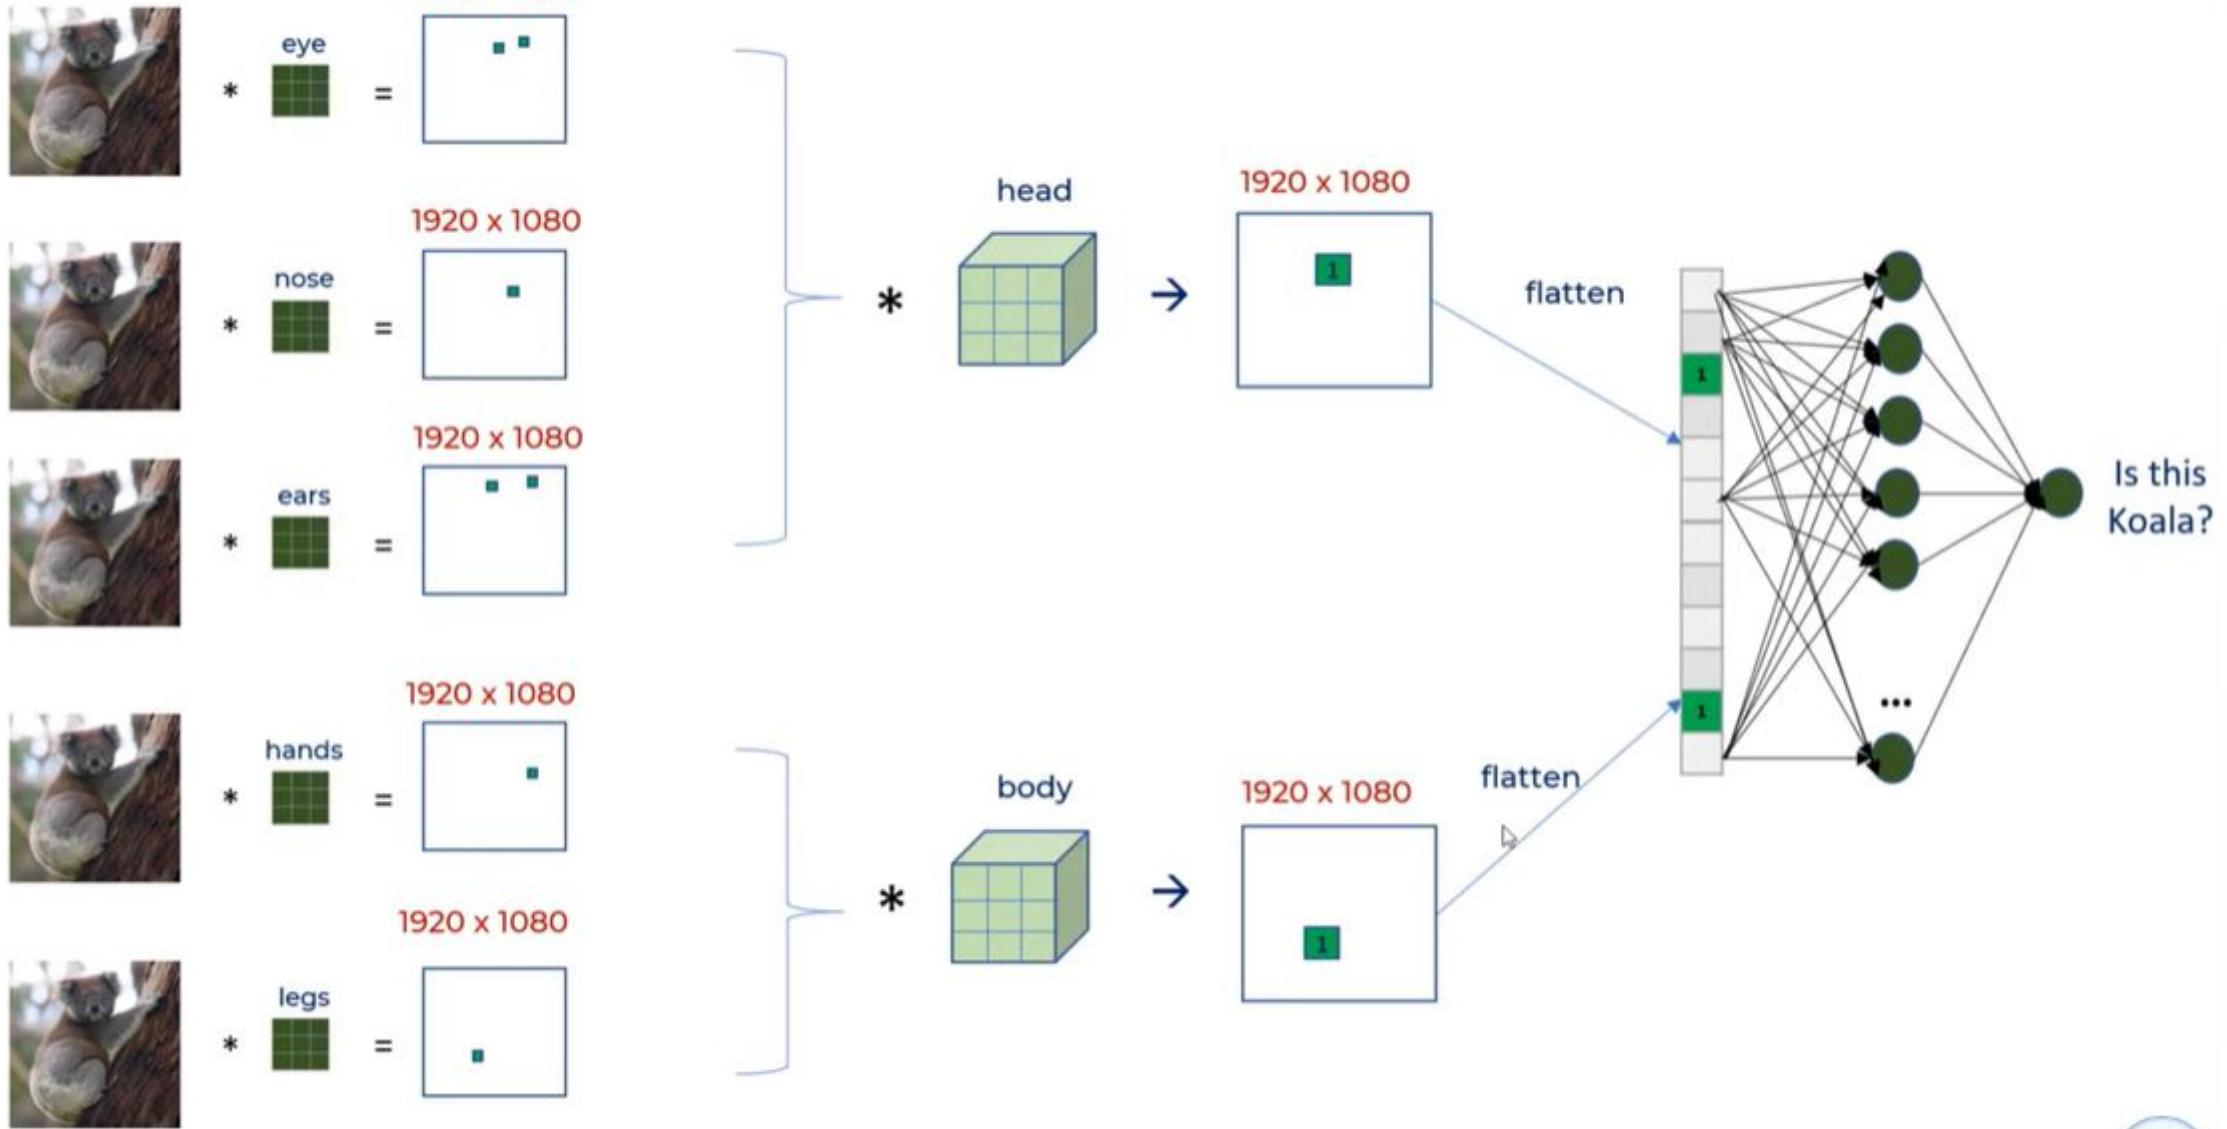
5	1	3	4
8	2	9	2
1	3	0	1
2	2	2	0

8	9
3	2



2 by 2 filter with stride = 2

# Pooling use- reduce dimension



# Max Pooling

-1	1	1	1	-1
-1	1	-1	1	-1
-1	1	1	1	-1
-1	-1	-1	1	-1
-1	-1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	-1	-1

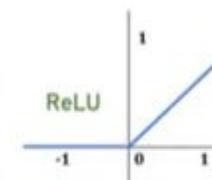
Loopy pattern  
filter

\*

1	1	1
1	-1	1
1	1	1



-0.11	<b>1</b>	-0.11
-0.55	0.11	-0.33
-0.33	0.33	-0.33
-0.22	-0.11	-0.22
-0.33	-0.33	-0.33



0	<b>1</b>	0
0	0.11	0
0	0.33	0
0	0	0
0	0	0

Max  
pooling  
→

<b>1</b>	<b>1</b>
0.33	0.33
0.33	0.33
0	0

There is average pooling also...

5	1	3	4
8	2	9	2
1	3	0	1
2	2	2	0

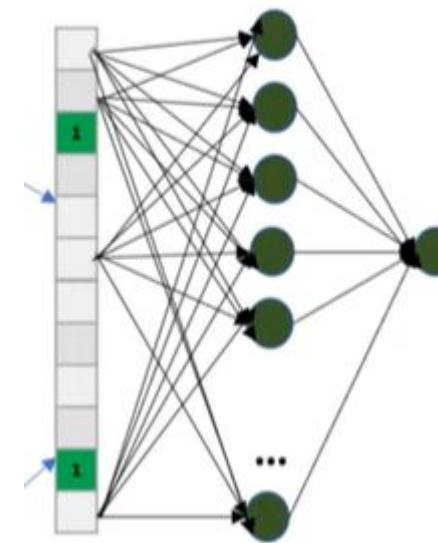
4	4.5
2	0.75

Lets get Serious

# CNN

Feature Extraction  $\Leftrightarrow$  Convolution Layer   Dense layers- Classification

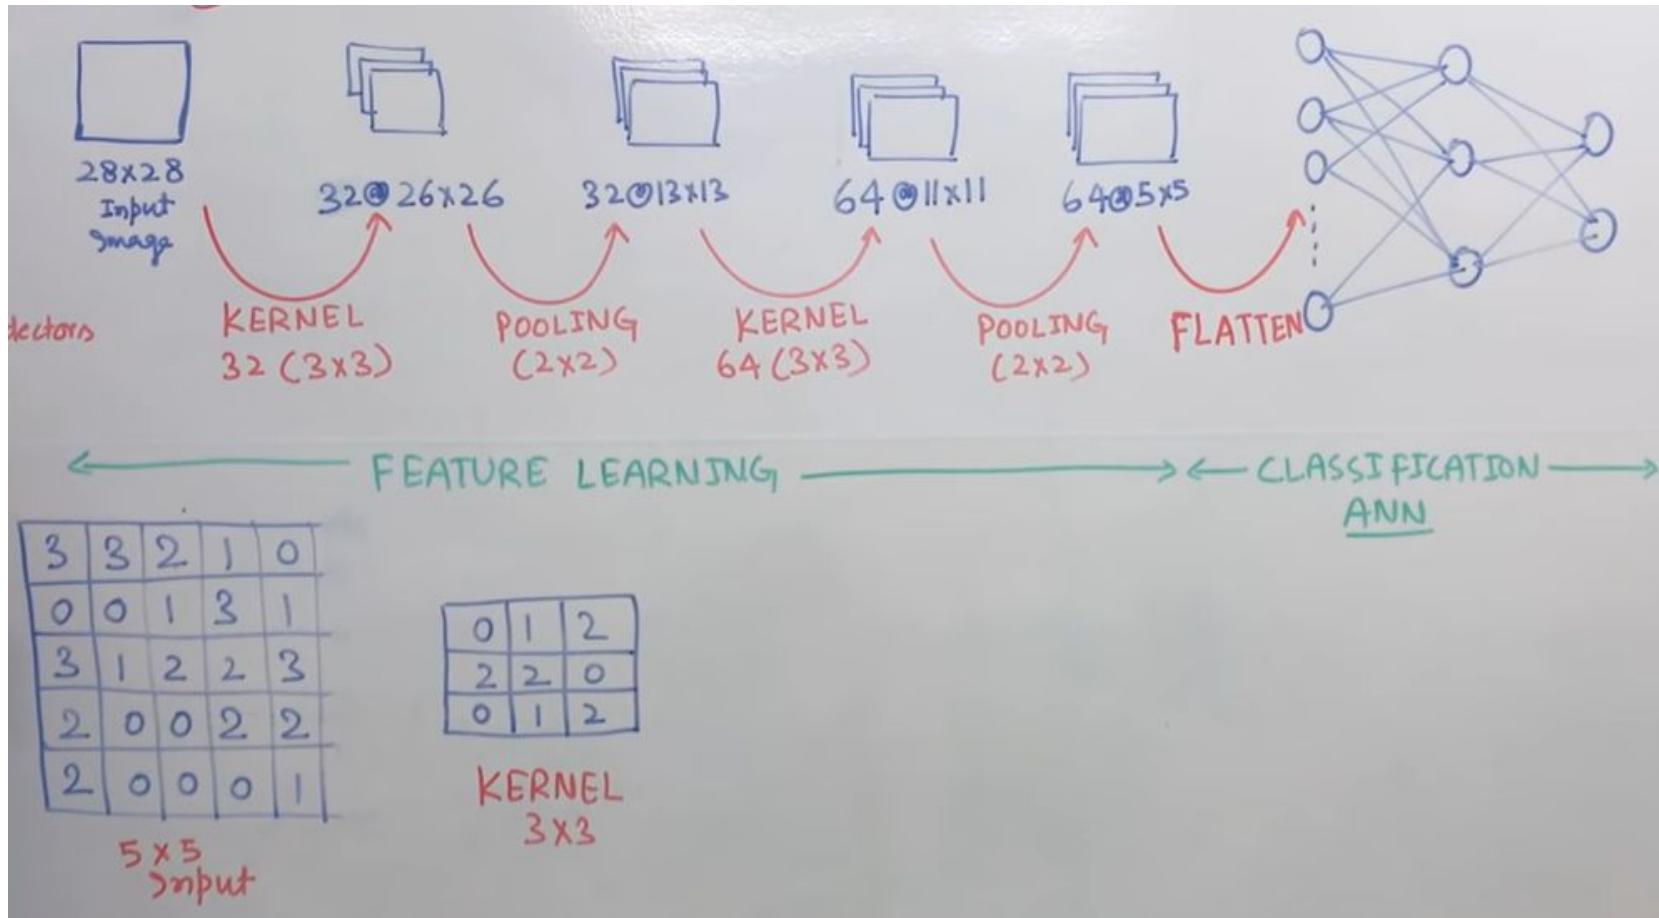
**Kernel = Filter  
Detector  
Stride  
Padding  
Pooling  
Flatten**

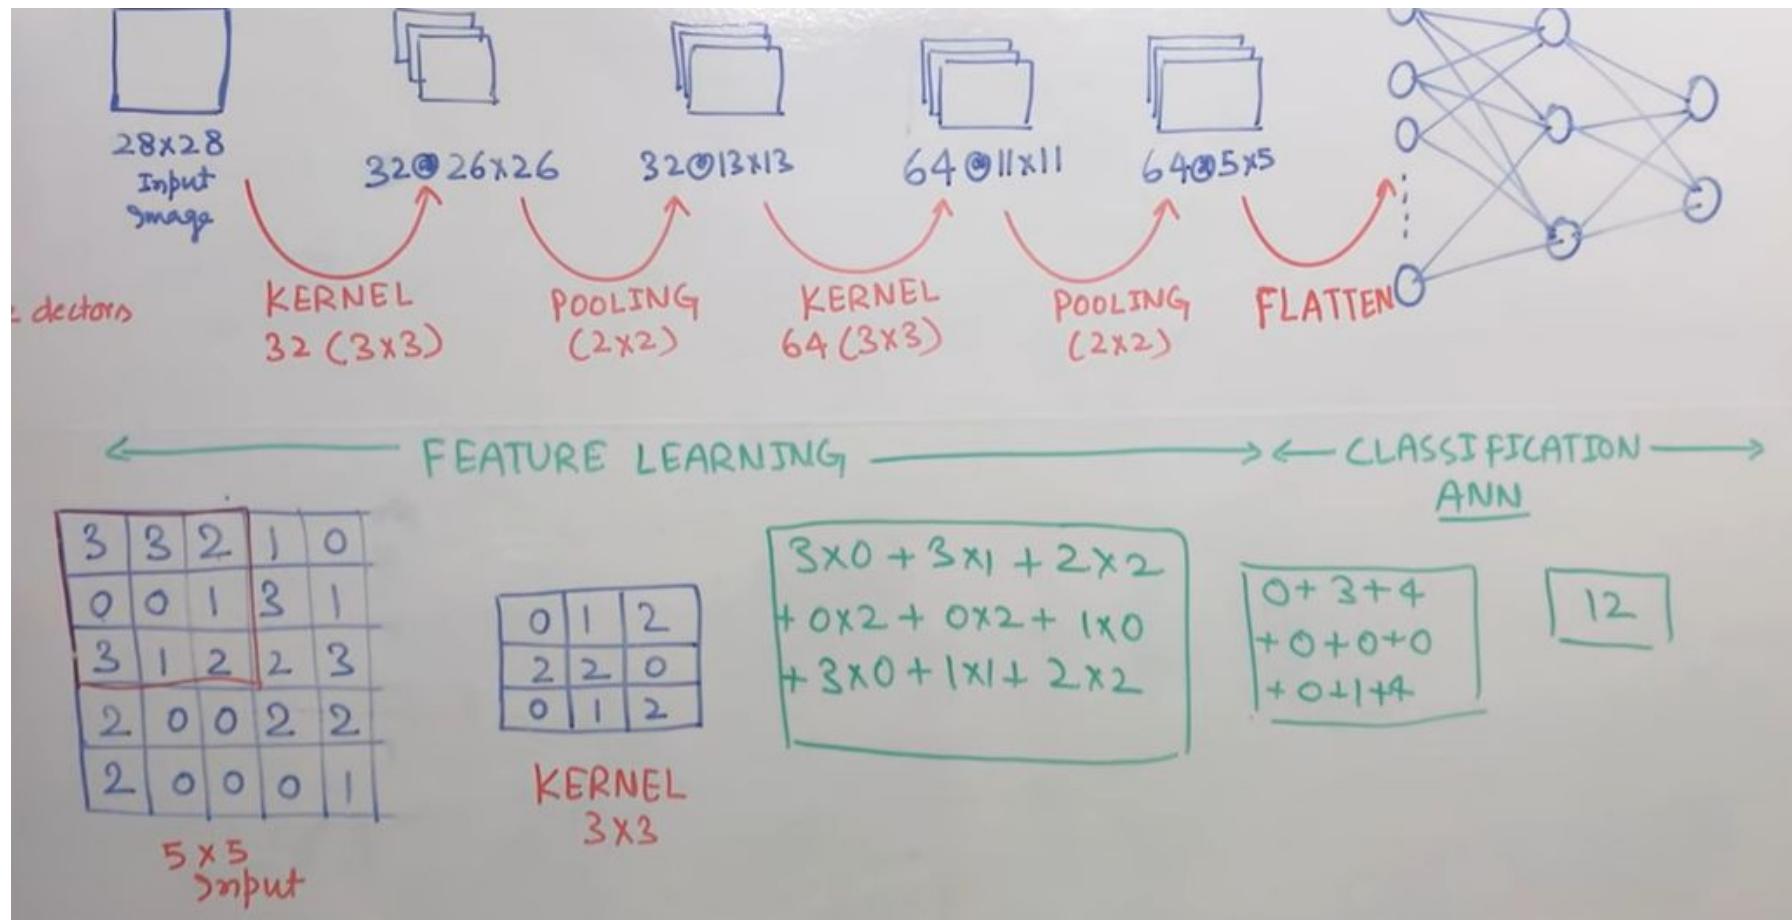


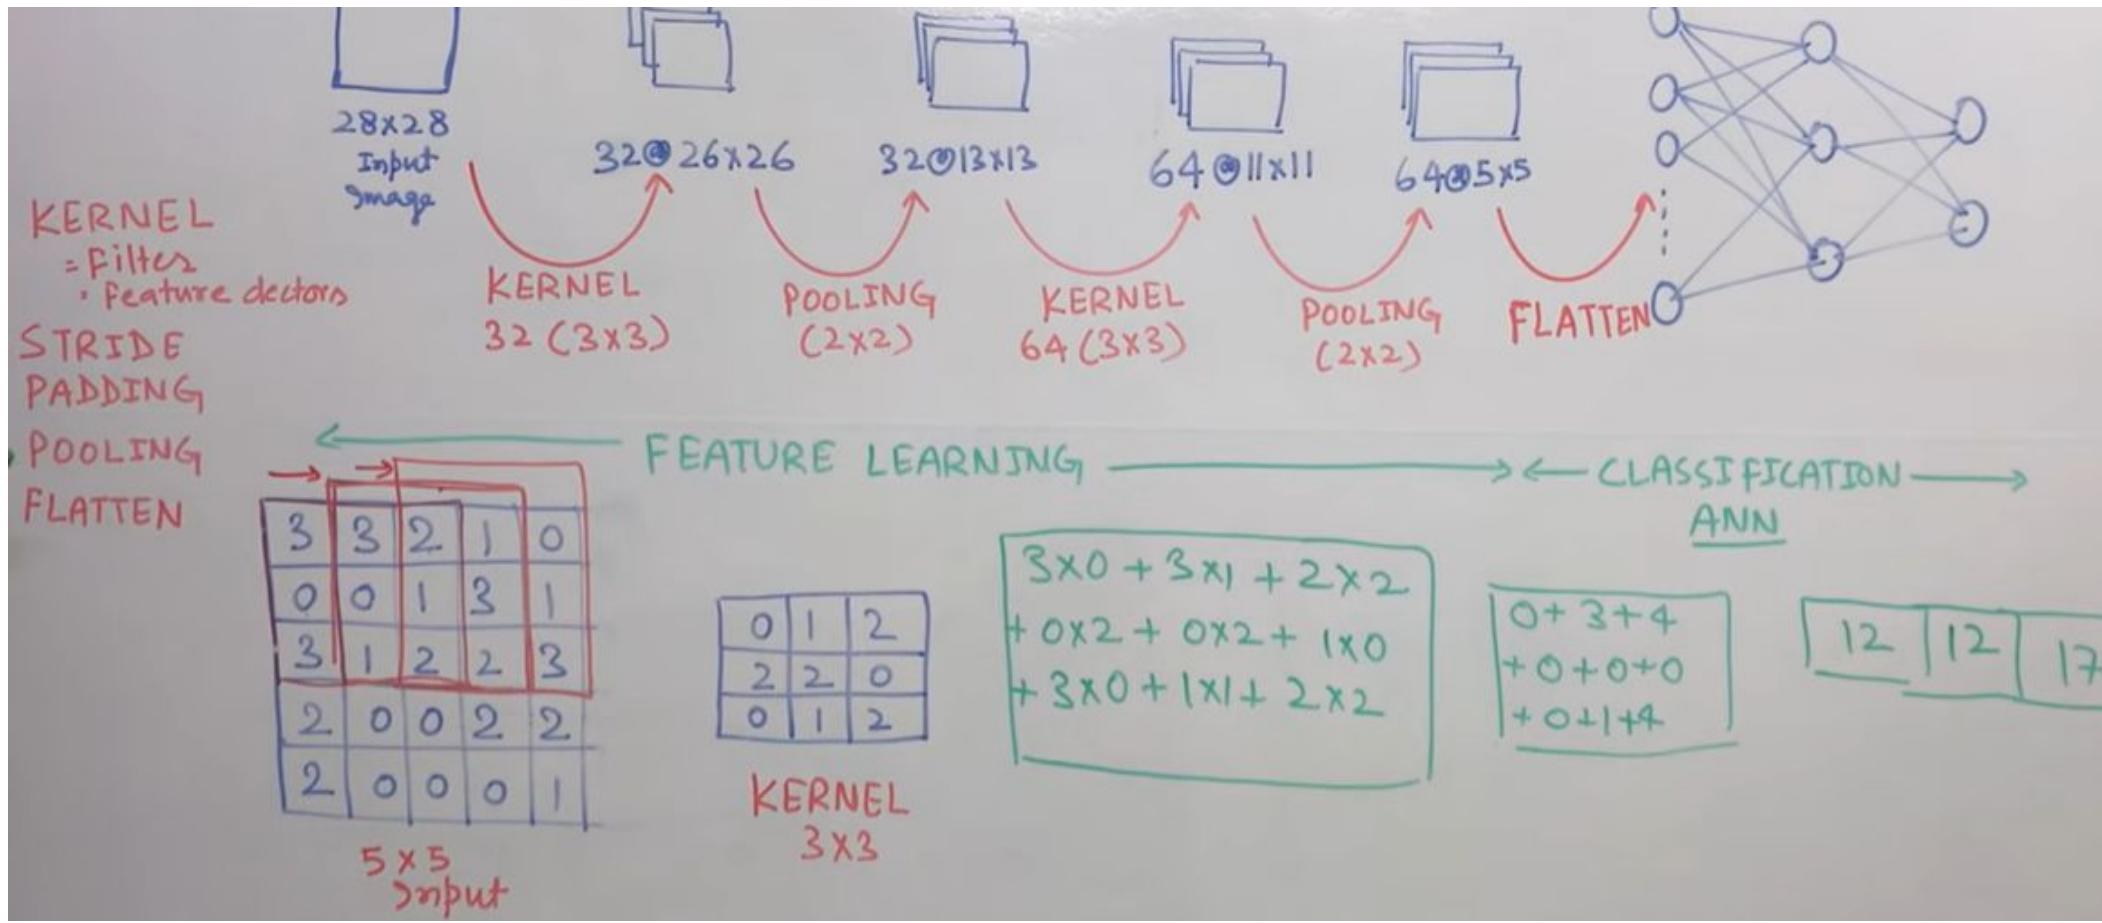
# Kernel

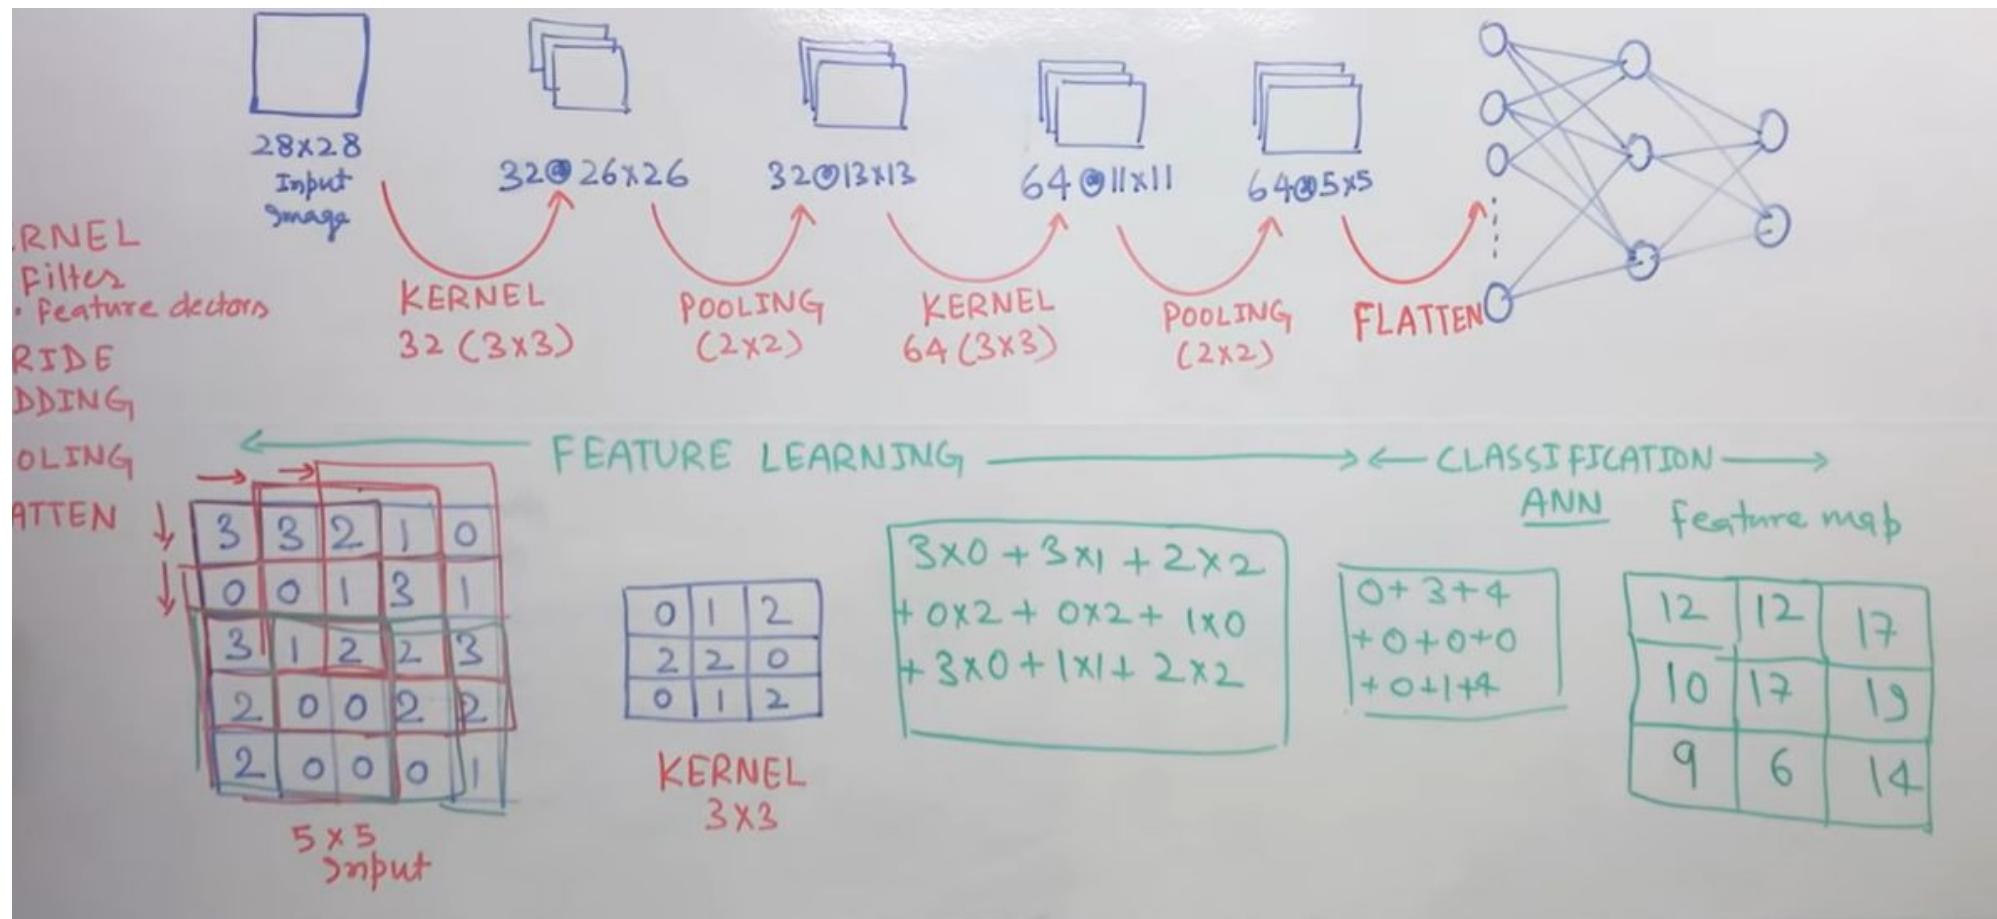
- **Kernel** is nothing but a filter that is used to extract the features from the images.
- The **kernel** is a matrix that moves over the input data, performs the dot product with the sub-region of input data, and gets the output as the matrix of dot products.
- **Kernel** moves on the input data by the stride value

# Kernel-> Filter+ Stride





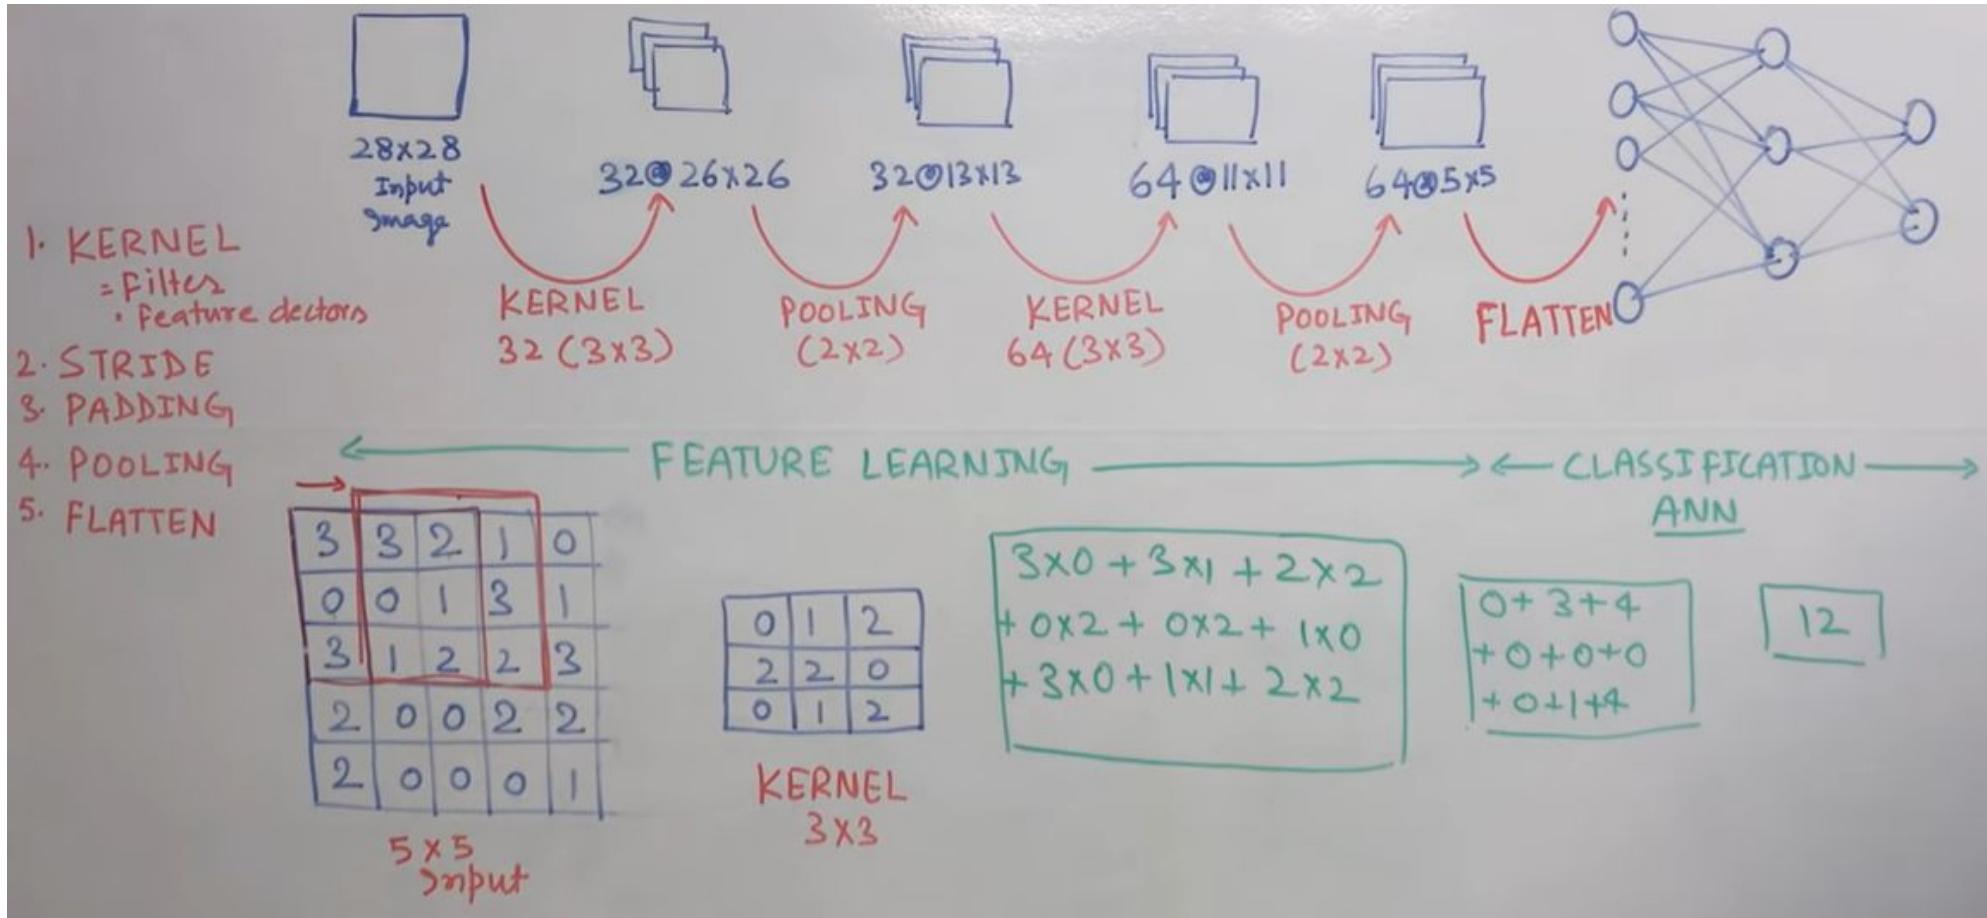




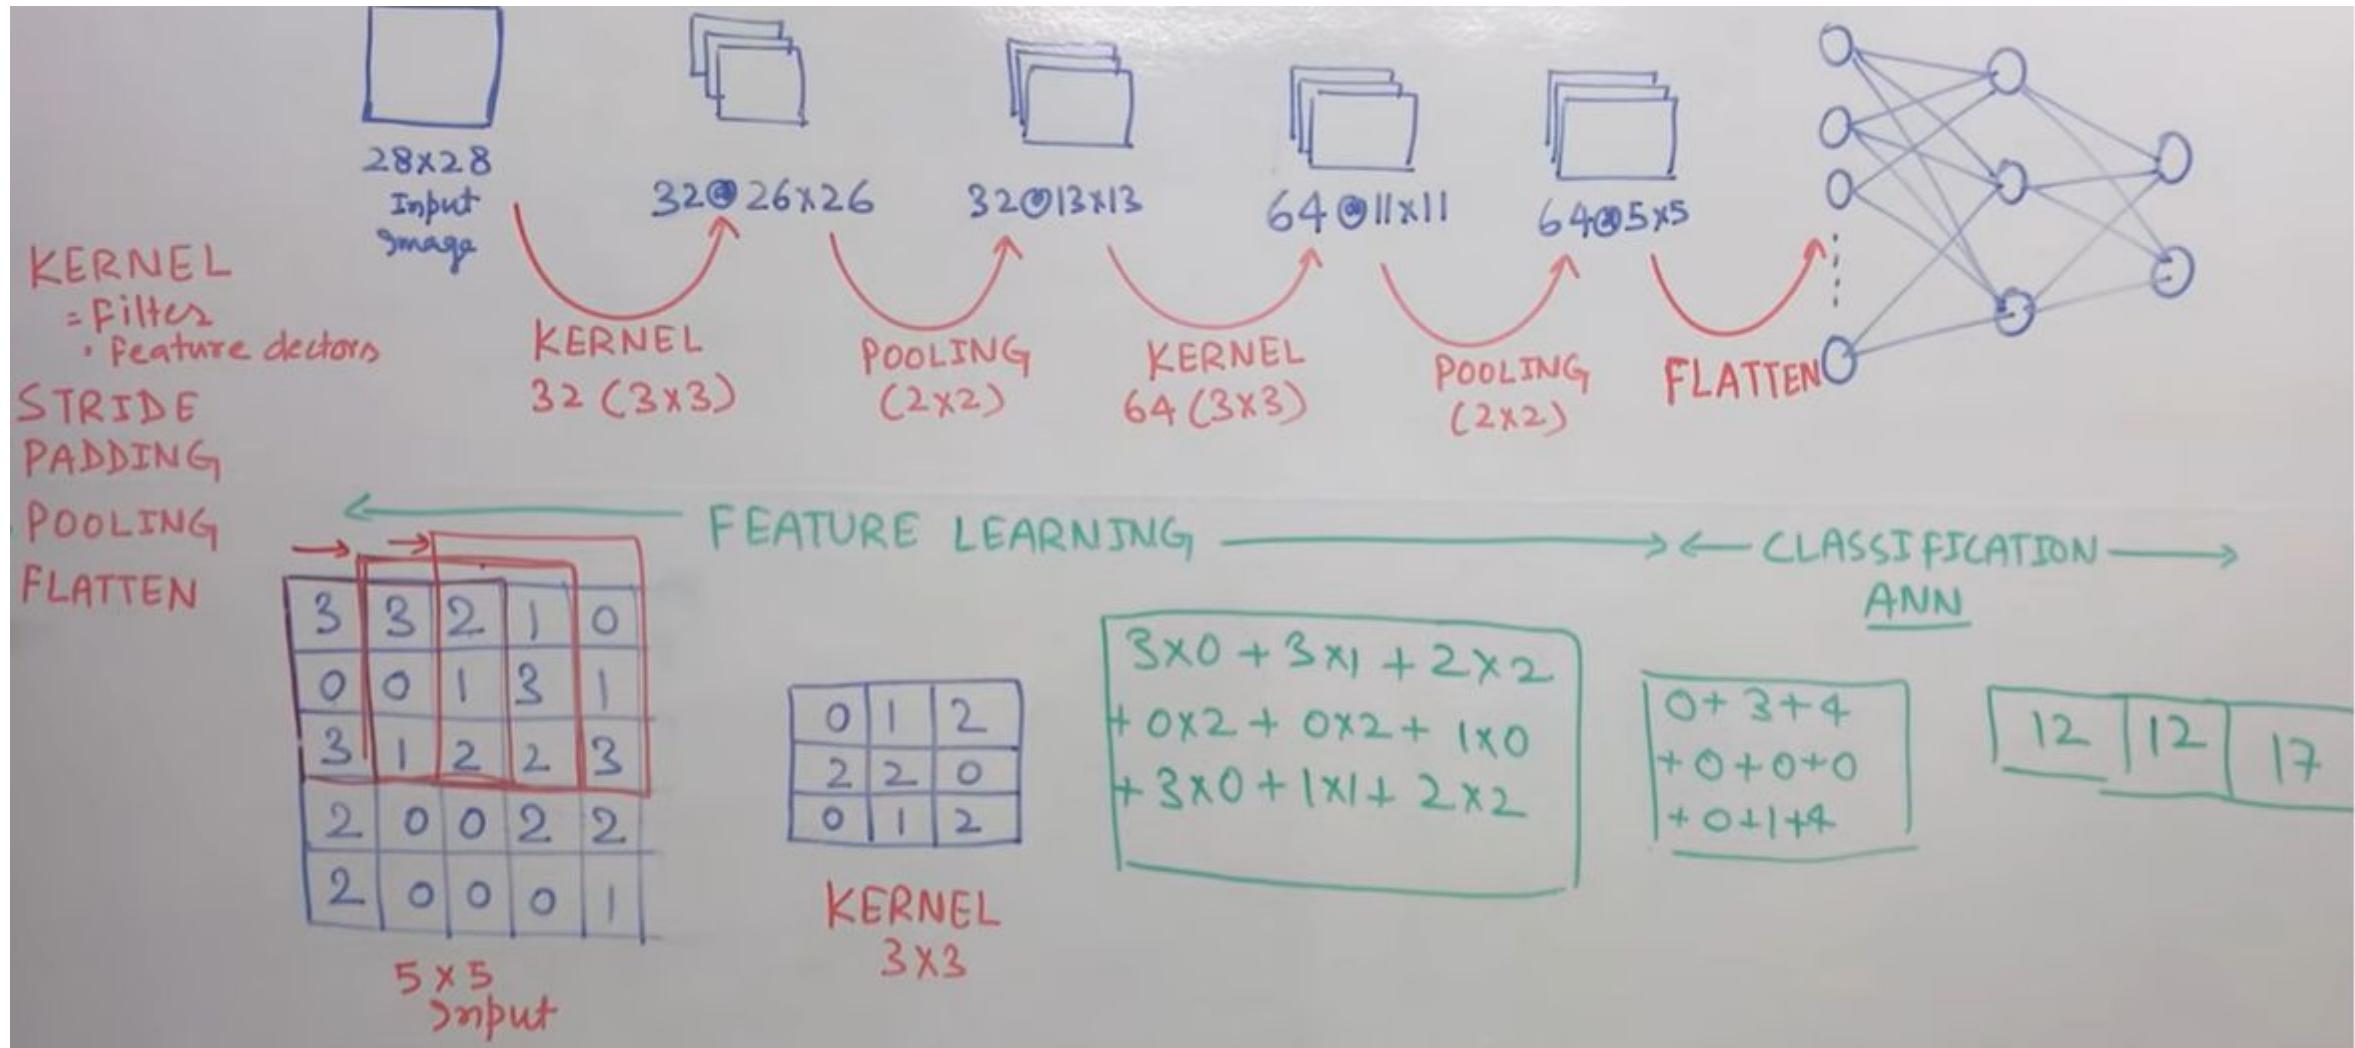
# Stride

- The filter is moved across the image left to right, top to bottom, with a one-pixel column change on the horizontal movements, then a one-pixel row change on the vertical movements.
- The amount of movement between applications of the filter to the input image is referred to as the stride, and it is almost always symmetrical in height and width dimensions.
- The default stride or strides in two dimensions is (1,1) for the height and the width movement.

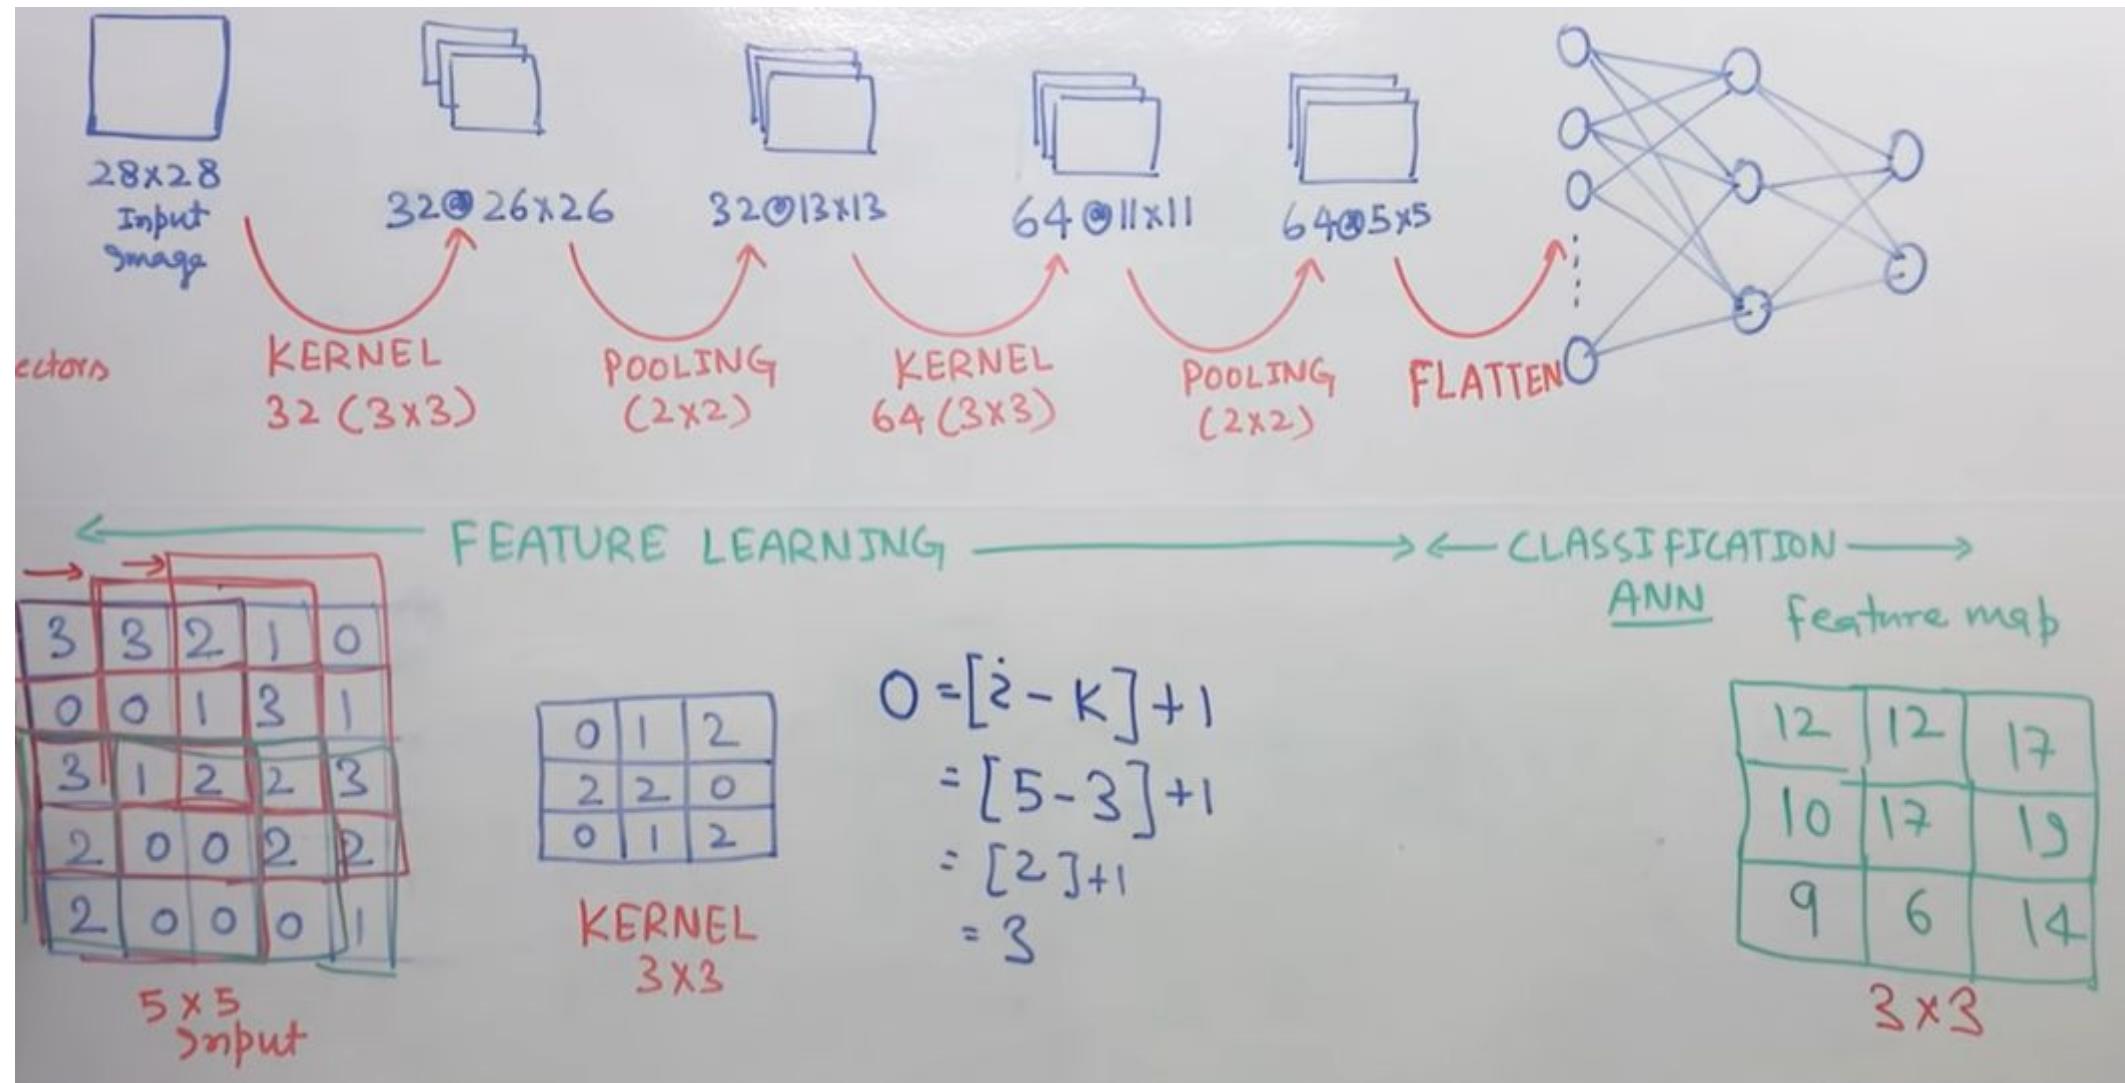
# Stride=1



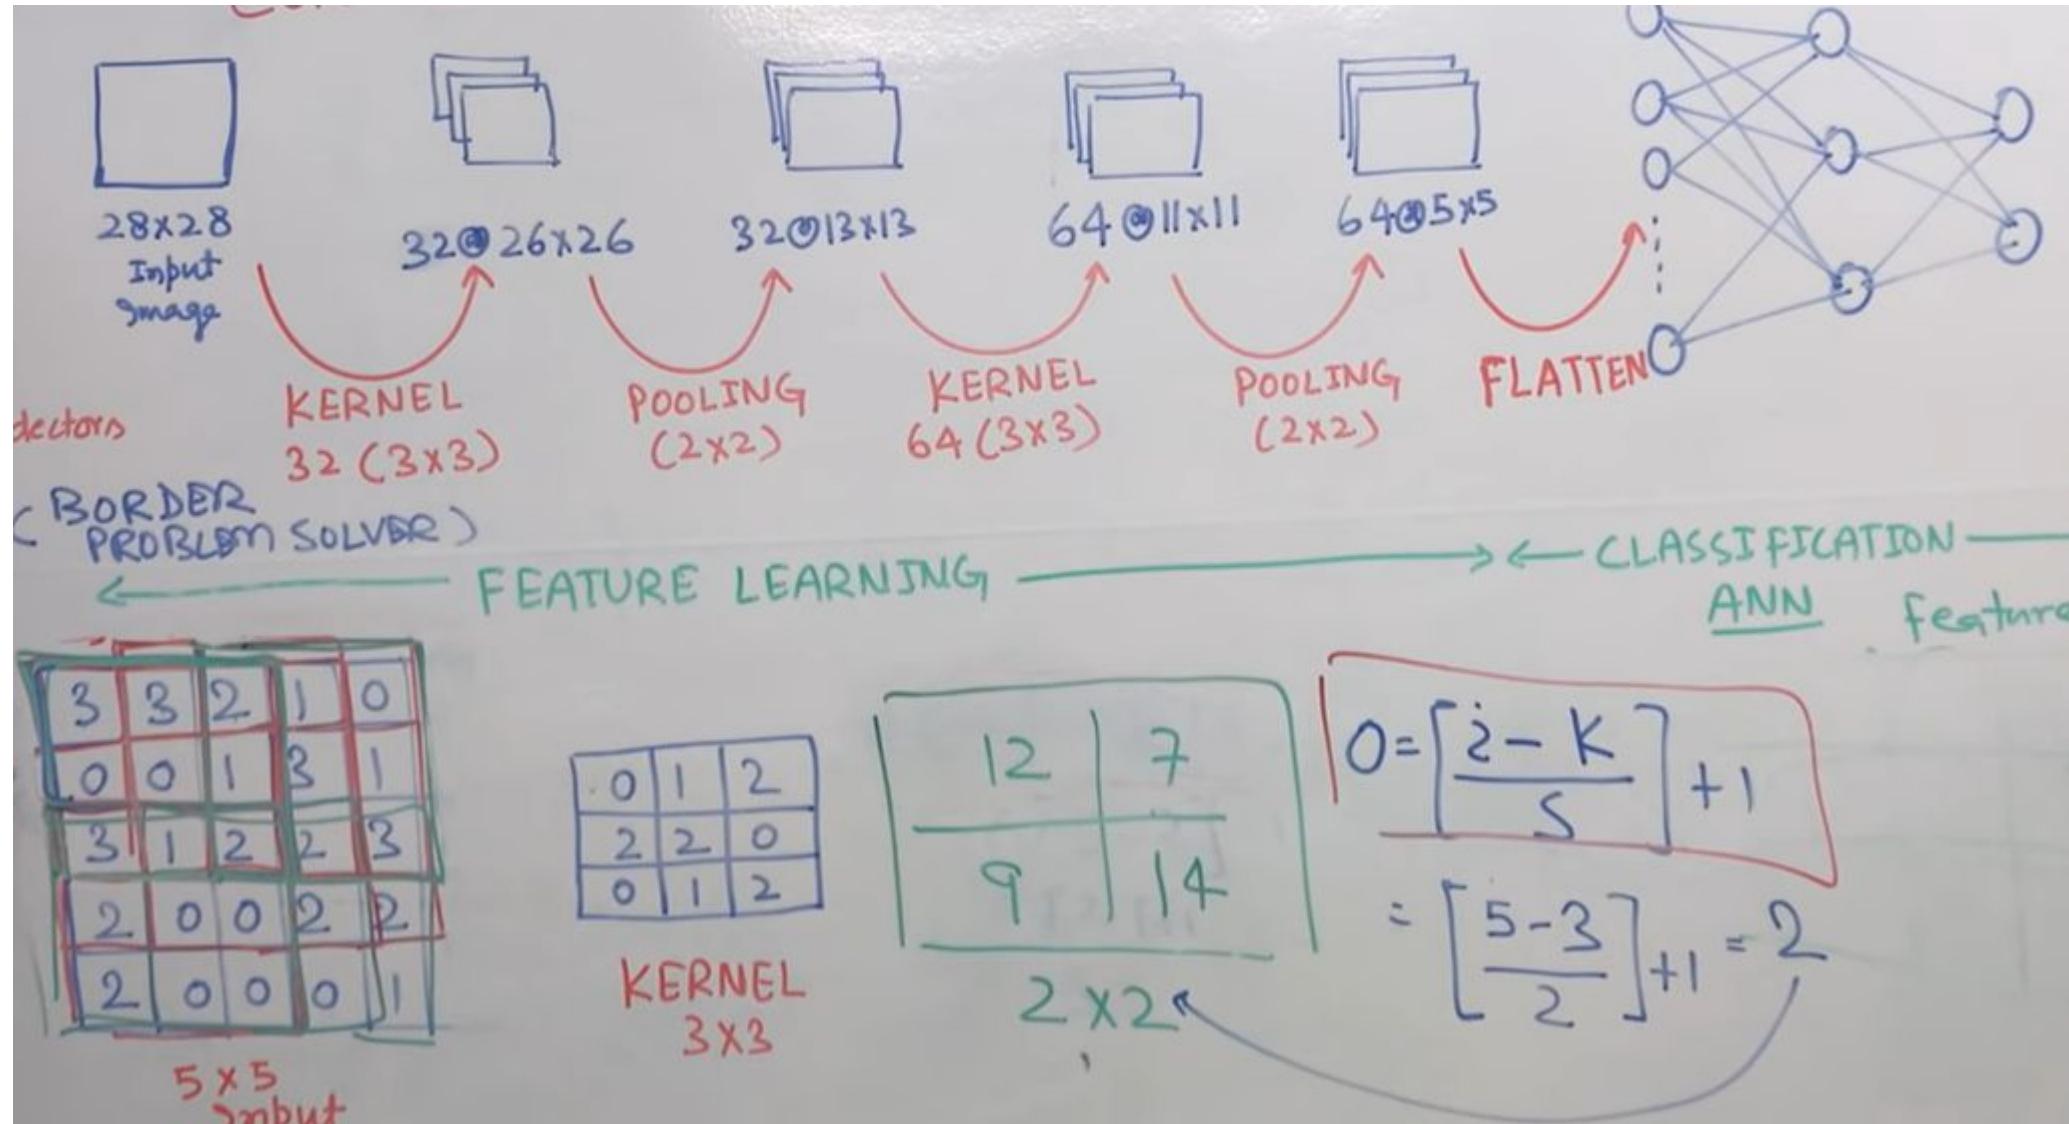
# Stride=1



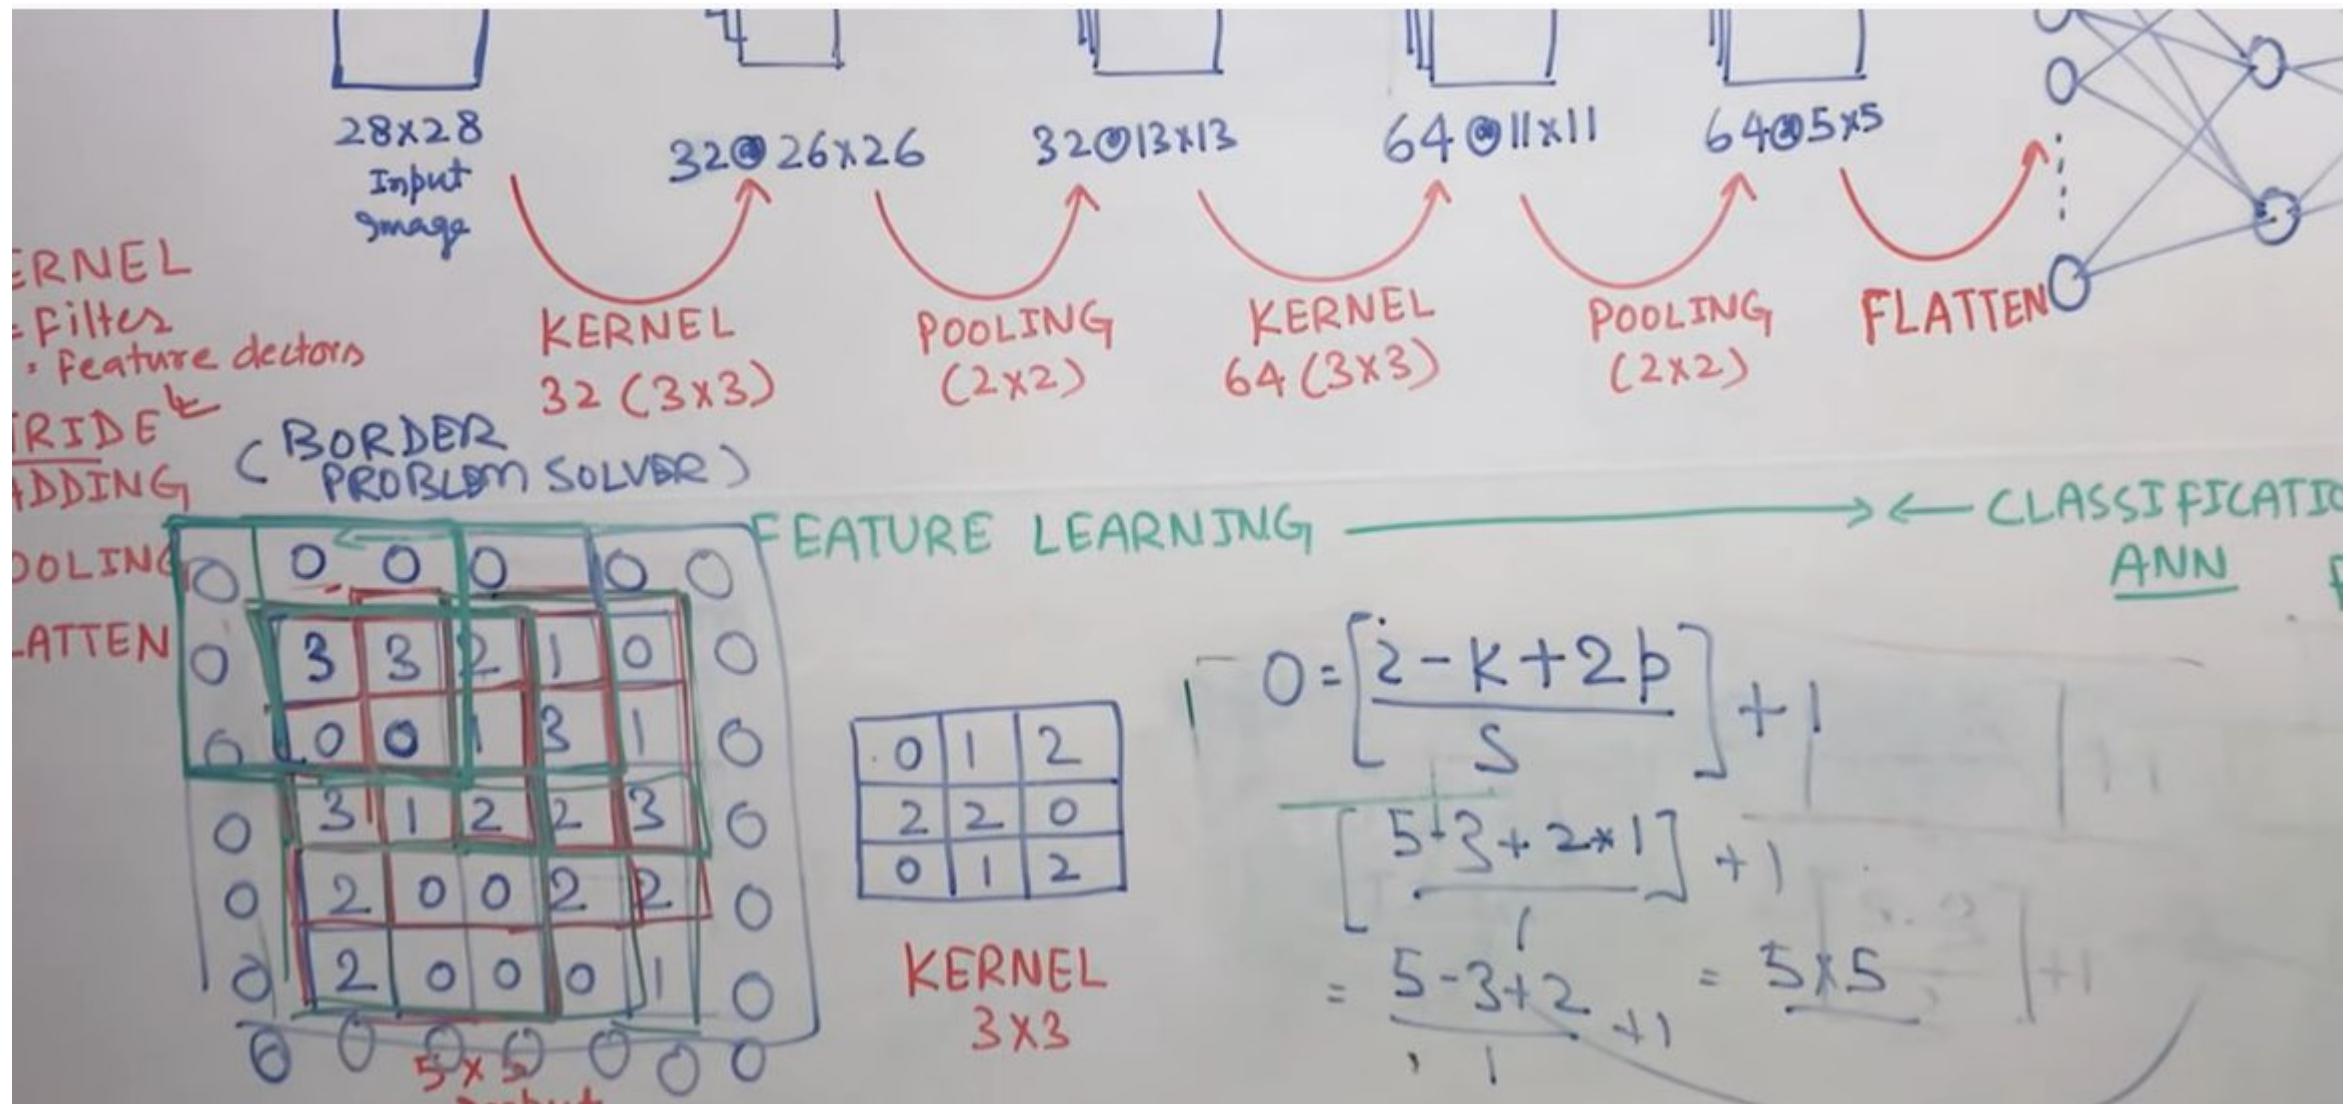
# Formula to calculate resultant matrix size, if stride = 1



# Formula to calculate resultant matrix size, if stride = 2



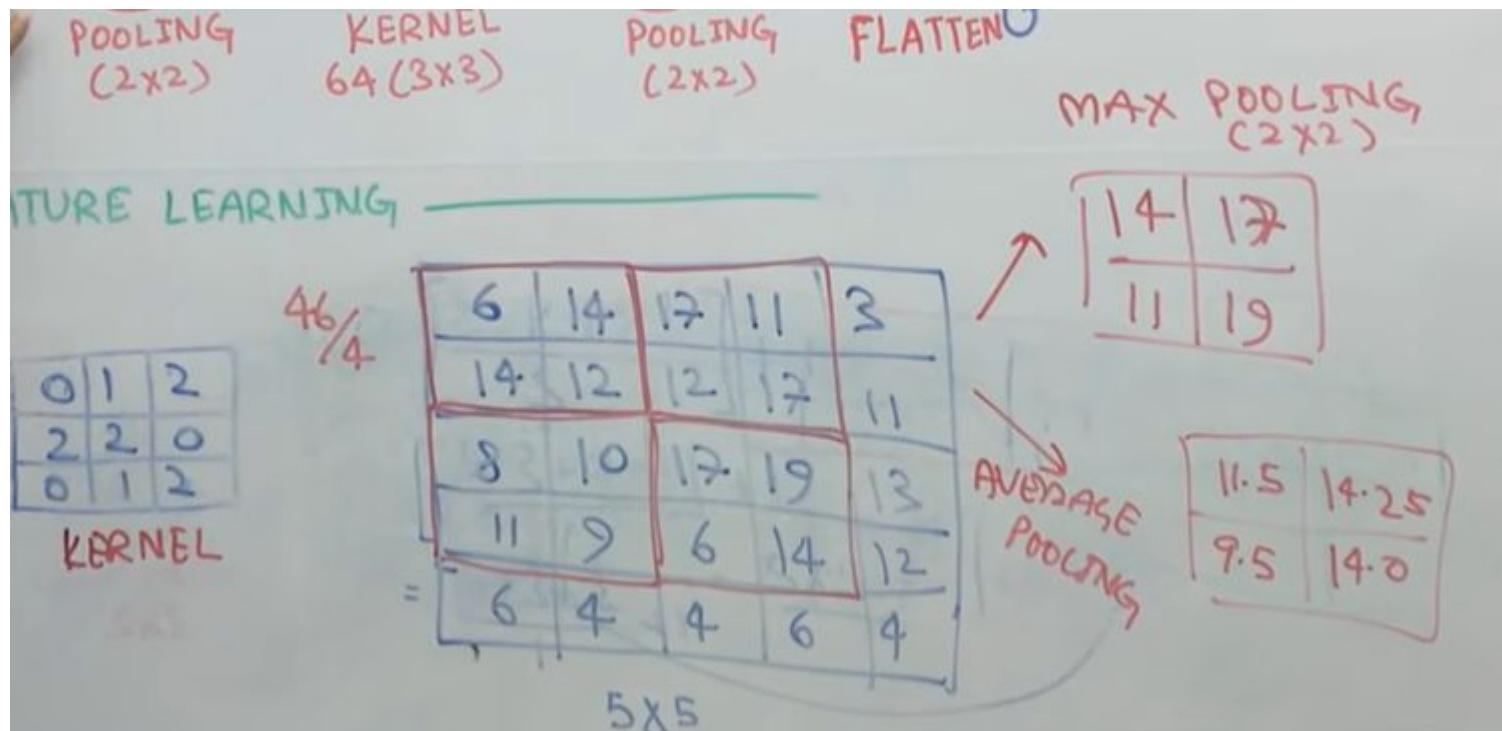
# Padding formula for resultant Matrix



# Pooling

- Pooling is required to down sample the detection of features in feature maps.
- Pooling layers provide an approach to down sampling feature maps by summarizing the presence of features in patches of the feature map. Two common pooling methods are average pooling and max pooling that summarize the average presence of a feature and the most activated presence of a feature respectively.

# Pooling types



# Benefits of pooling

Reduces dimensions & computation

Reduce overfitting as there are less parameters

Model is tolerant towards variations, distortions

## Convolution

- Connections sparsity reduces overfitting
- Conv + Pooling gives location invariant feature detection
- Parameter sharing

## ReLU

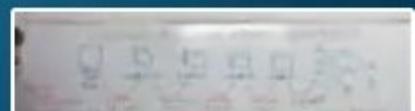
- Introduces nonlinearity
- Speeds up training, faster to compute

## Pooling

- Reduces dimensions and computation
- Reduces overfitting
- Makes the model tolerant towards small distortion and variations

# Flattening

- Once the pooled featured map is obtained, the next step is to flatten it.
- It involves transforming the entire pooled feature map matrix into a single column which is then fed to the neural network for processing.



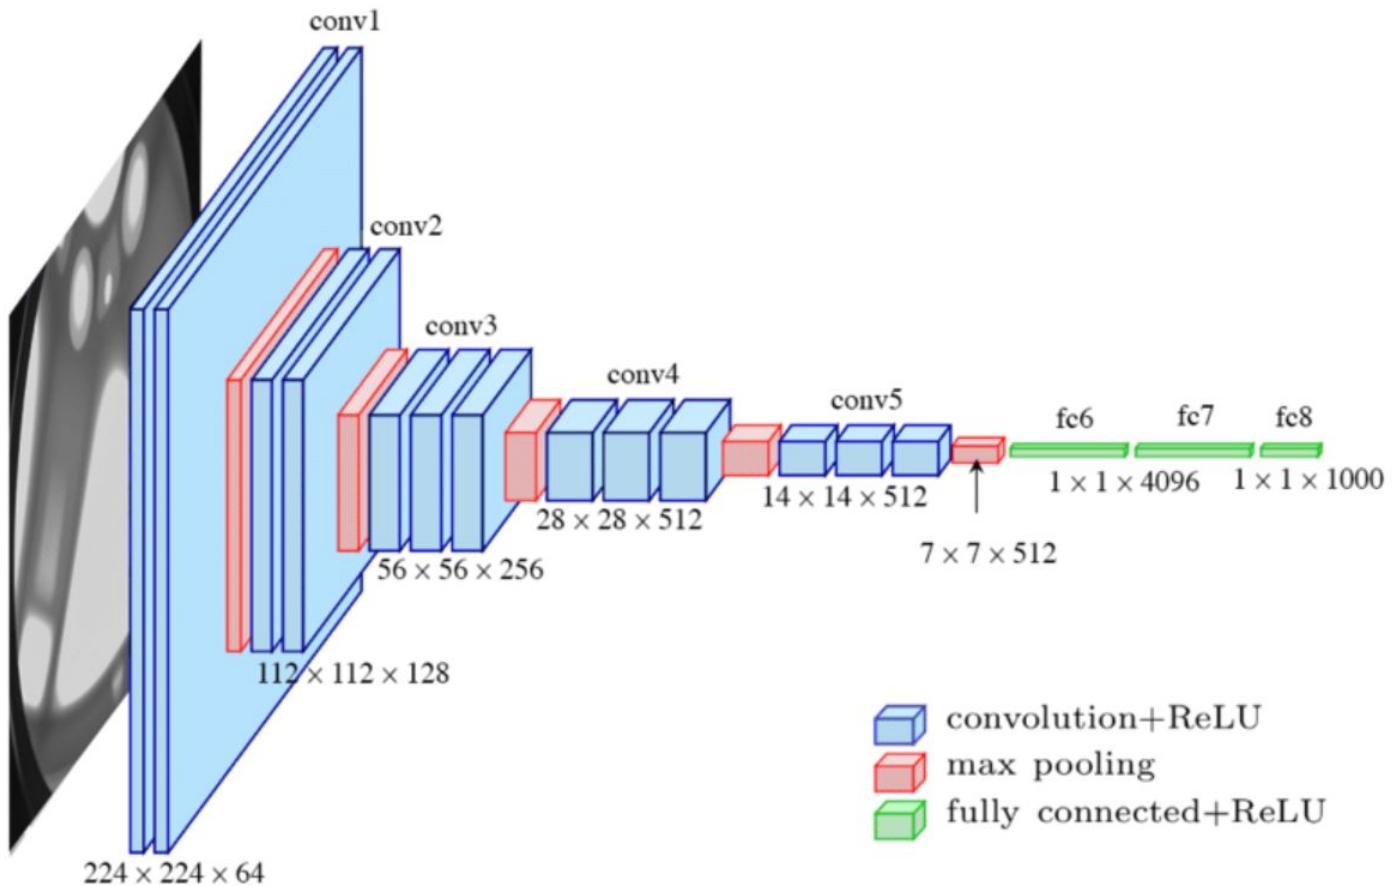
# Convolutional Neural Network: Components of CNN Architecture,

- <https://youtu.be/Y1qxl-Df4Lk>
- <https://youtu.be/zfiSAzpy9NM>

# Properties of CNN

	AlexNet	GoogLeNet	VGG-19
Model	Sequential	Inception modul	Sequential
Depth	8	22	19
Number of layers	25	144	47
Number of filters	96	64	64
Filter size	11x11	7x7	3x3
Stride	4x4	1x1	1x1
Dropout		50%	
Activation function		ReLU	
Output size		2	

# Architectures of CNN



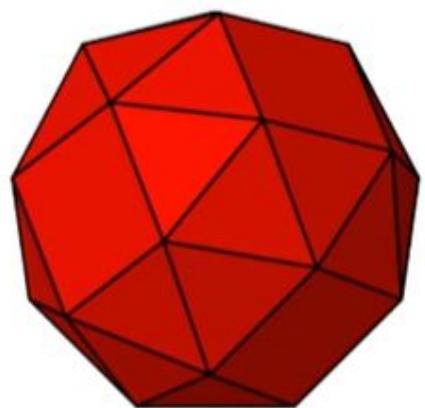
# Applications of CNN

- Image recognition / Retina/ Face Recognition
- Image captioning
- Image based Search Engine
- Image quality improvement
- Object recognition
- Colour generation from BW pictures
- Tumour reorganization

# Autoencoder: Introduction

- Data compression
- Unsupervised NN that uses ML for performing compression of data – dimensionality reduction
- Can learn non linear transformation

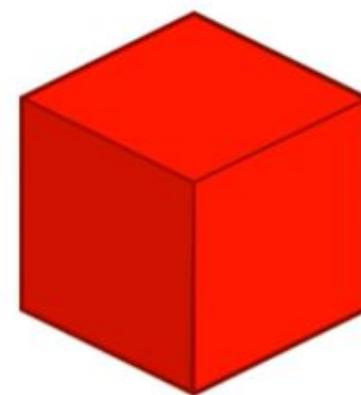
Multidimensional Data



Data represented best

Slow performance, High Precision

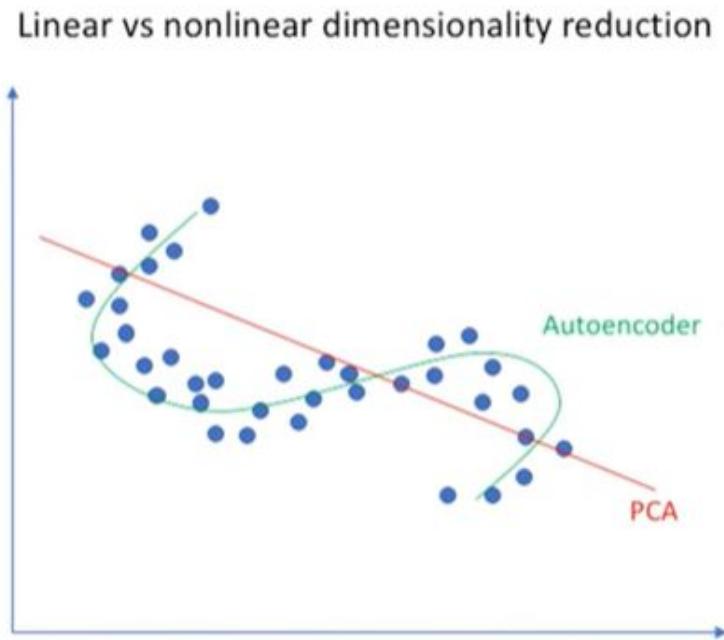
Low Dimensional Data



Reduced Precision

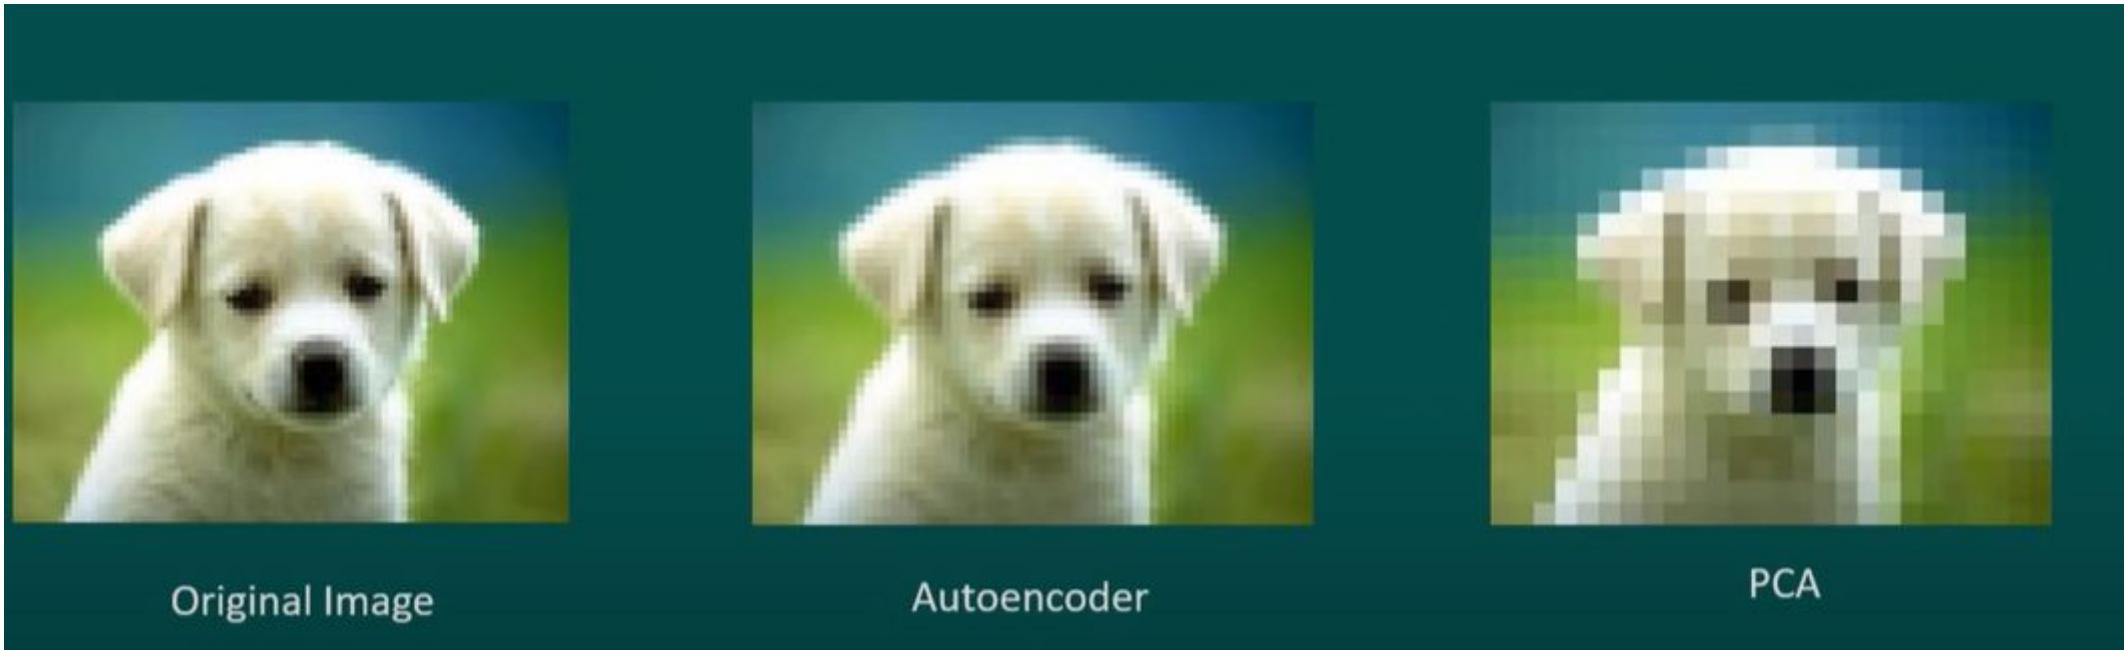
High Performance

# Need of Autoencoders when Principle Component analysis is there?



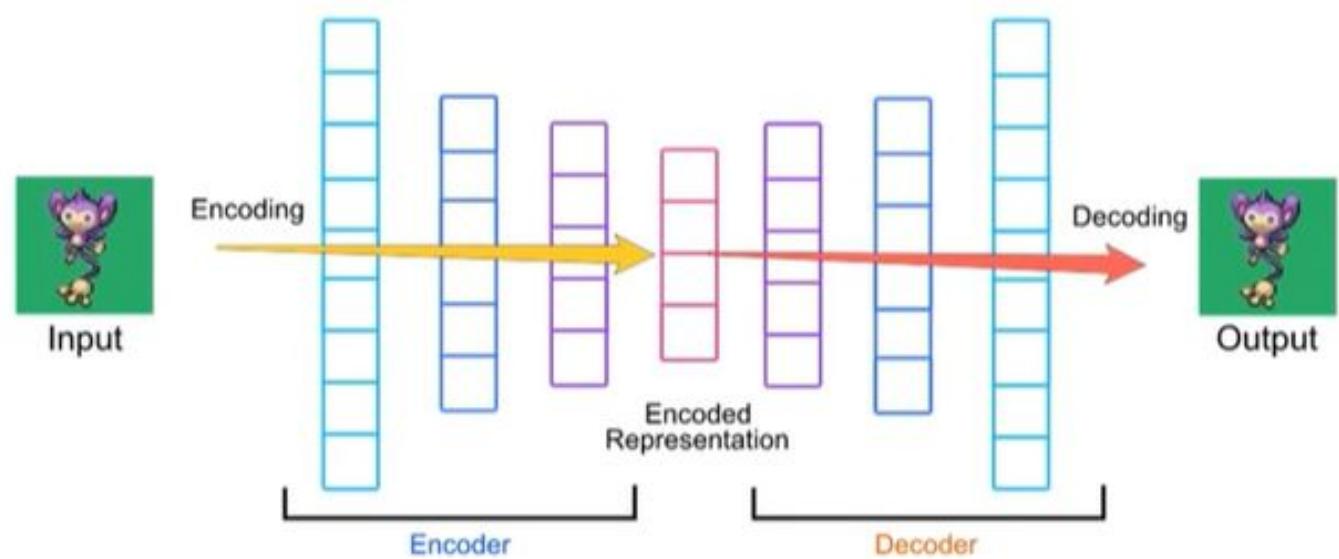
- Non-linear Transformations**  
Non-linear activation function and multiple layers
- Convolutional Layers**  
An autoencoder doesn't have to learn dense layers
- Higher Efficiency**  
More efficient to learn several layers with an autoencoder
- Multiple Transformations**  
It gives a representation as the output of each layer

Need of Autoencoders – dimension reduction and data visualization are the application of AECD

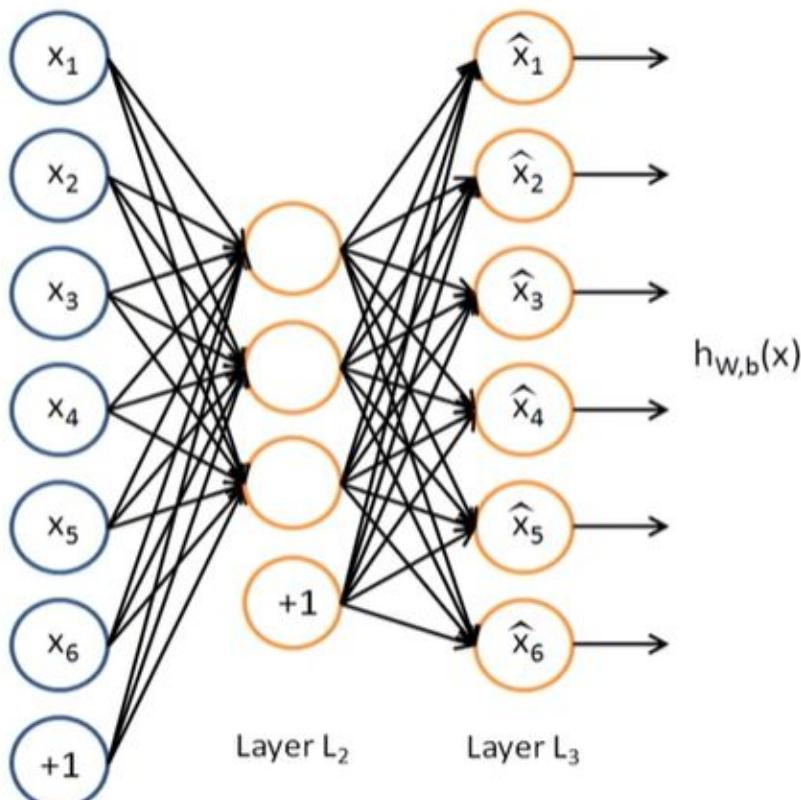


# Introduction to Autoencoders

An **autoencoder** neural network is an unsupervised Machine learning algorithm that applies backpropagation, setting the target values to be equal to the inputs.



# Autoencoder: Key Features

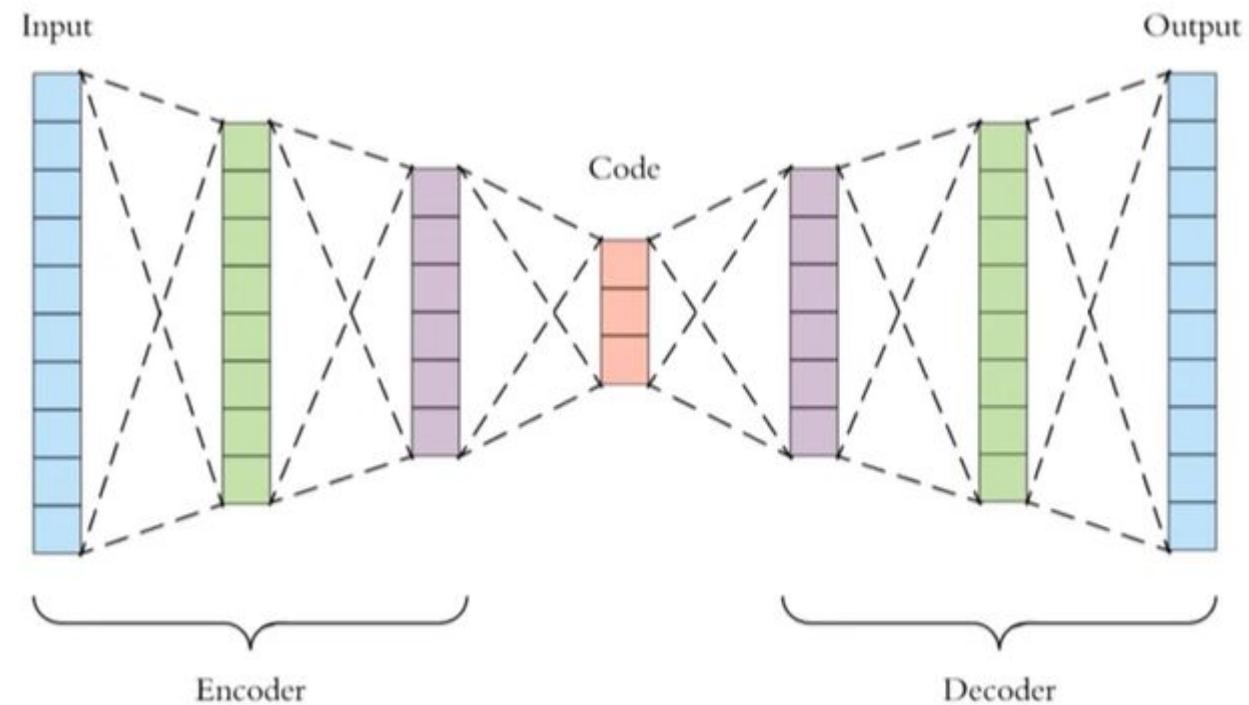
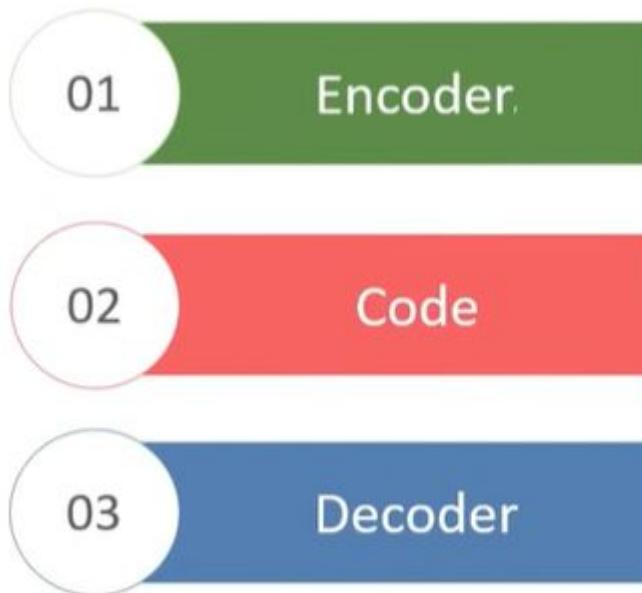


## Key Facts about Autoencoders

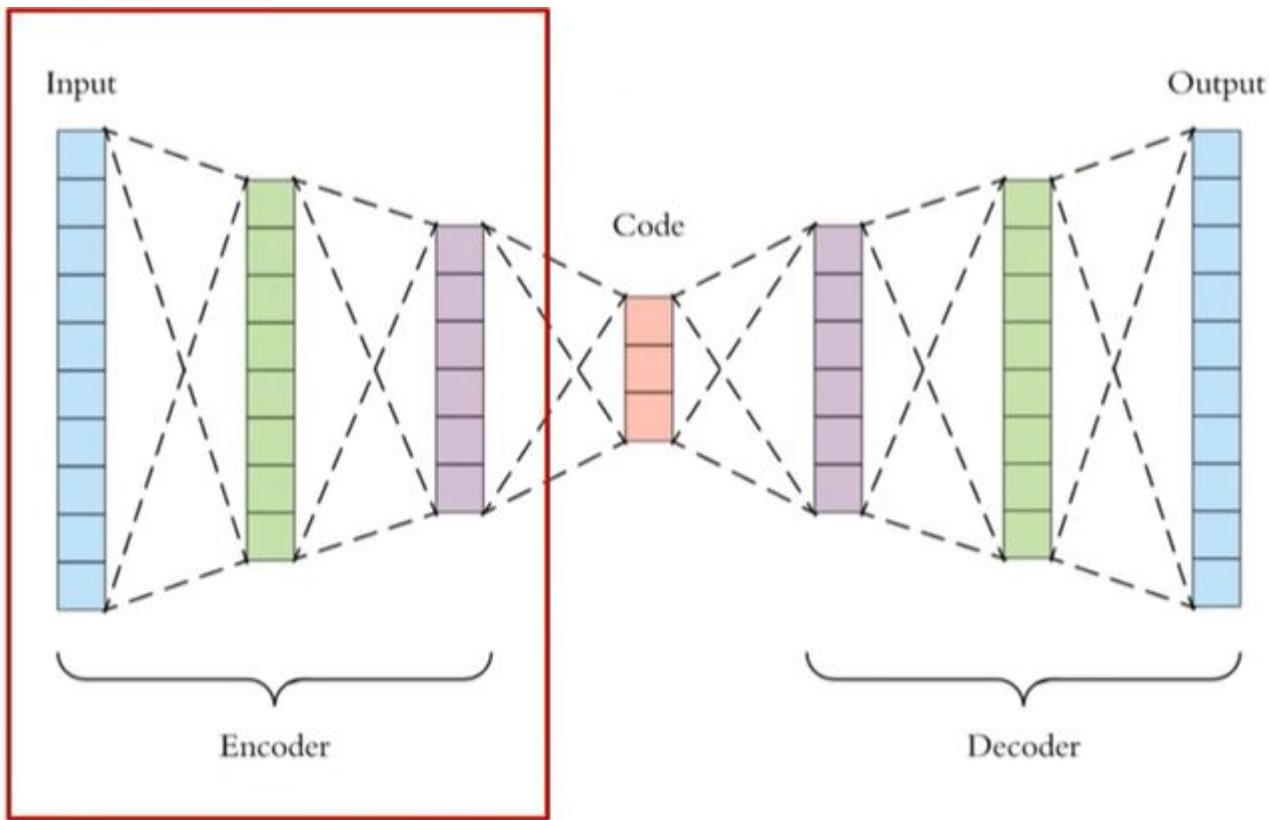
- It is an unsupervised ML algorithm similar to PCA
- It minimizes the same objective function as PCA
- It is a neural network
- The neural network's target output is its input

# Components of Autoencoder

## Components of Autoencoders



# Components of Autoencoder

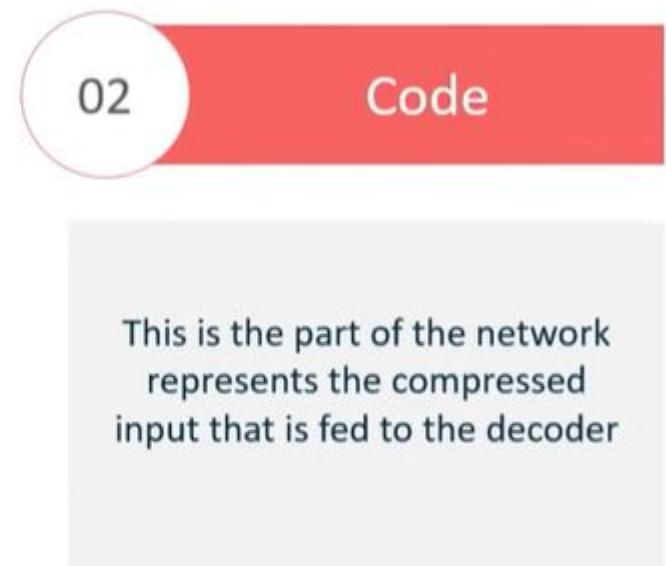
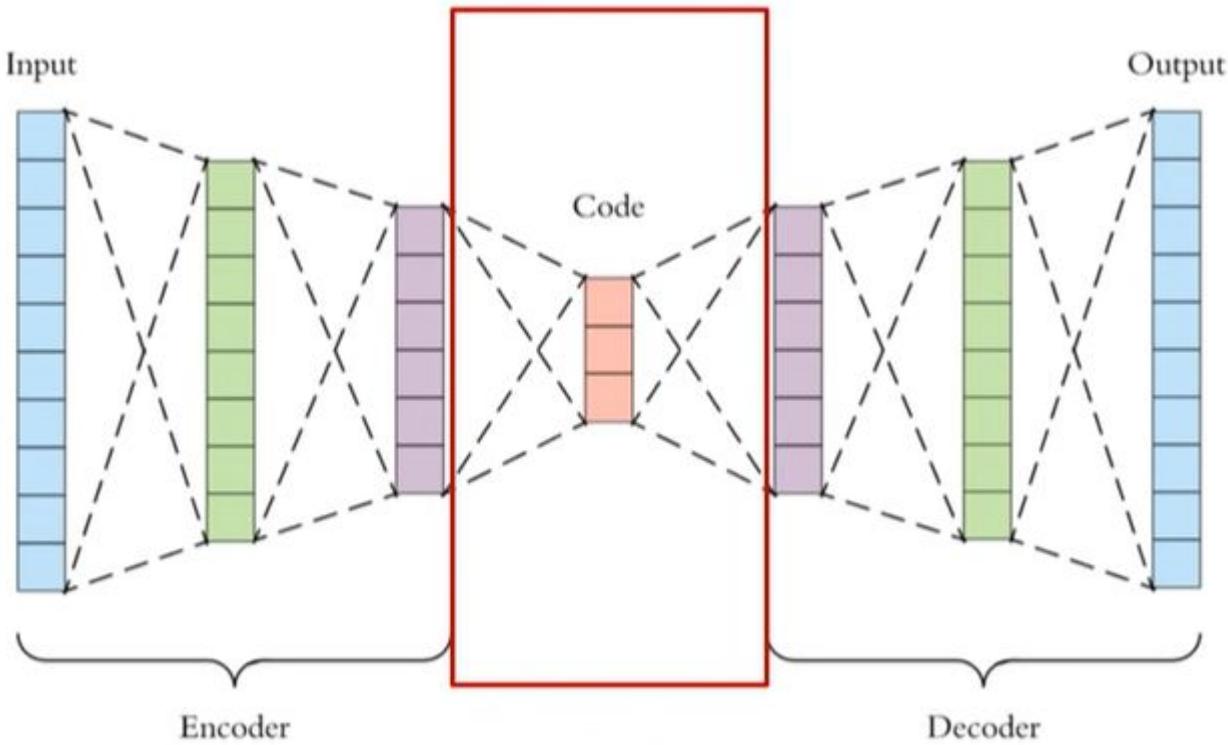


01

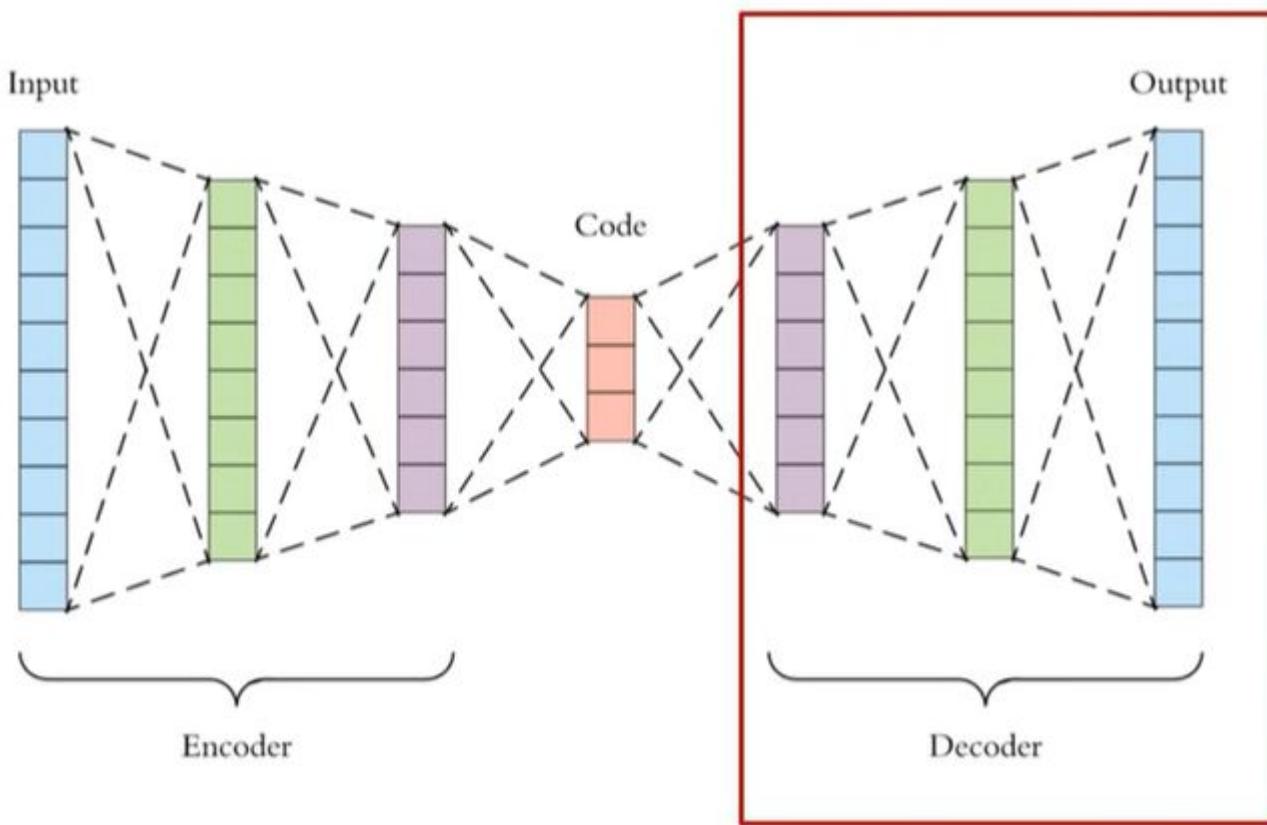
## Encoder

This is the part of the network that compresses the input into a latent space representation.

# Components of Autoencoder



# Components of Autoencoder



03

## Decoder

This part aims to reconstruct  
the input from the latent space  
representation

# Properties of Autoencoders



## Unsupervised 01

Autoencoders are considered an unsupervised learning technique since they don't need explicit labels to train on

## 02 Data-specific

Autoencoders are only able to meaningfully compress data similar to what they have been trained on

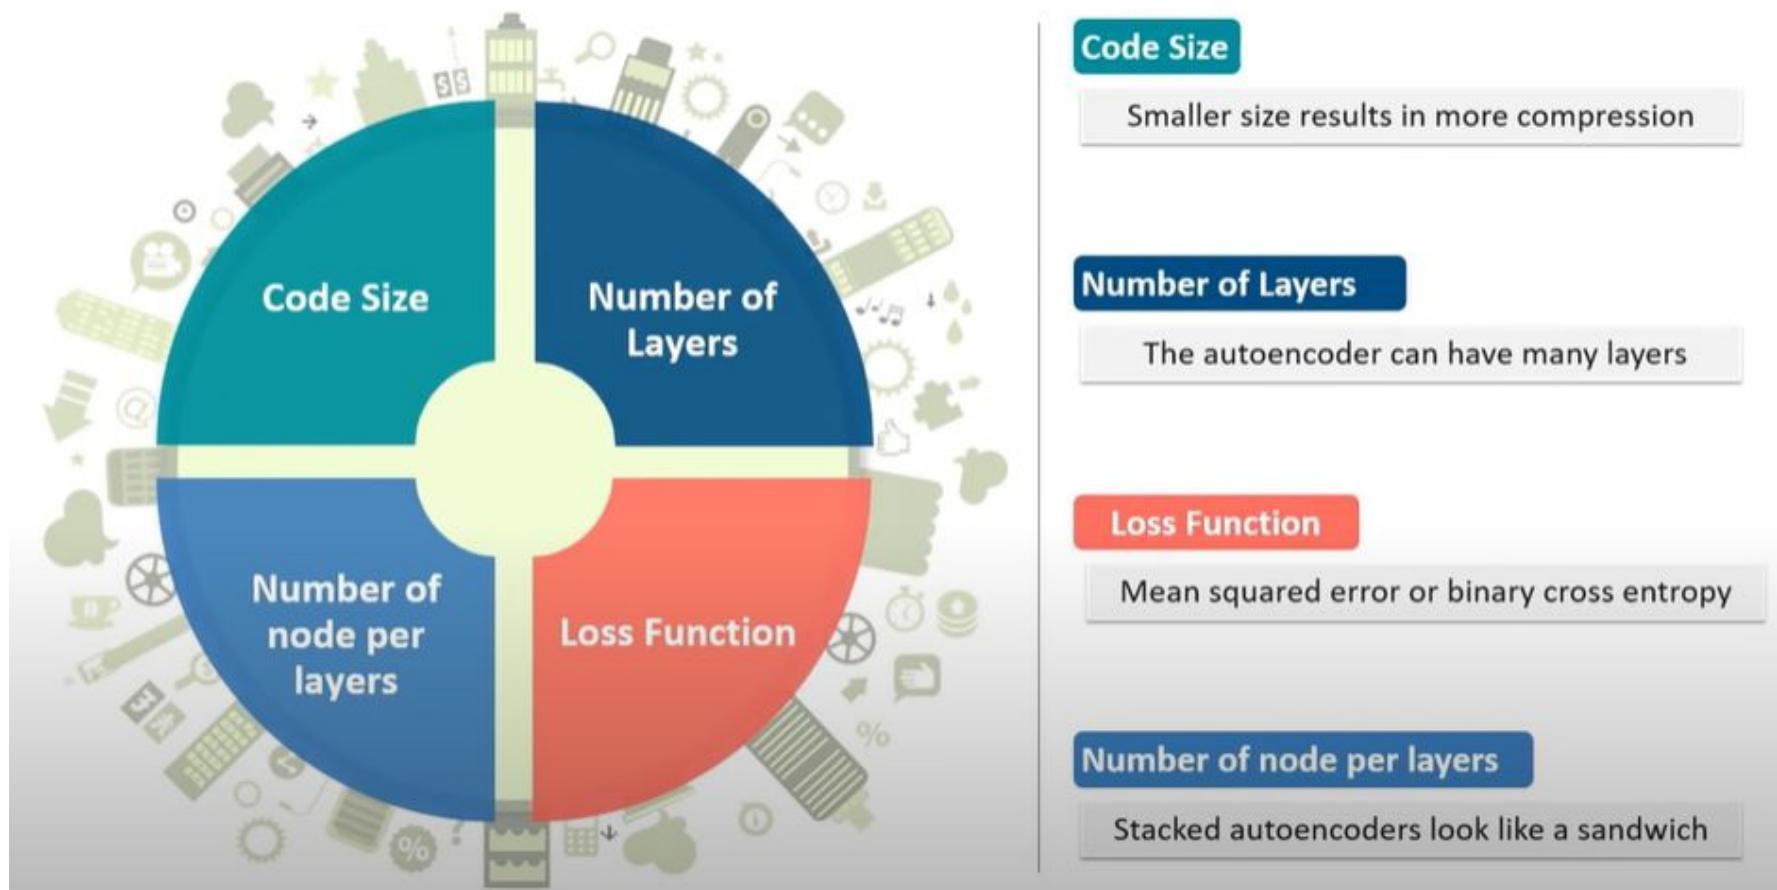


## 03 Lossy

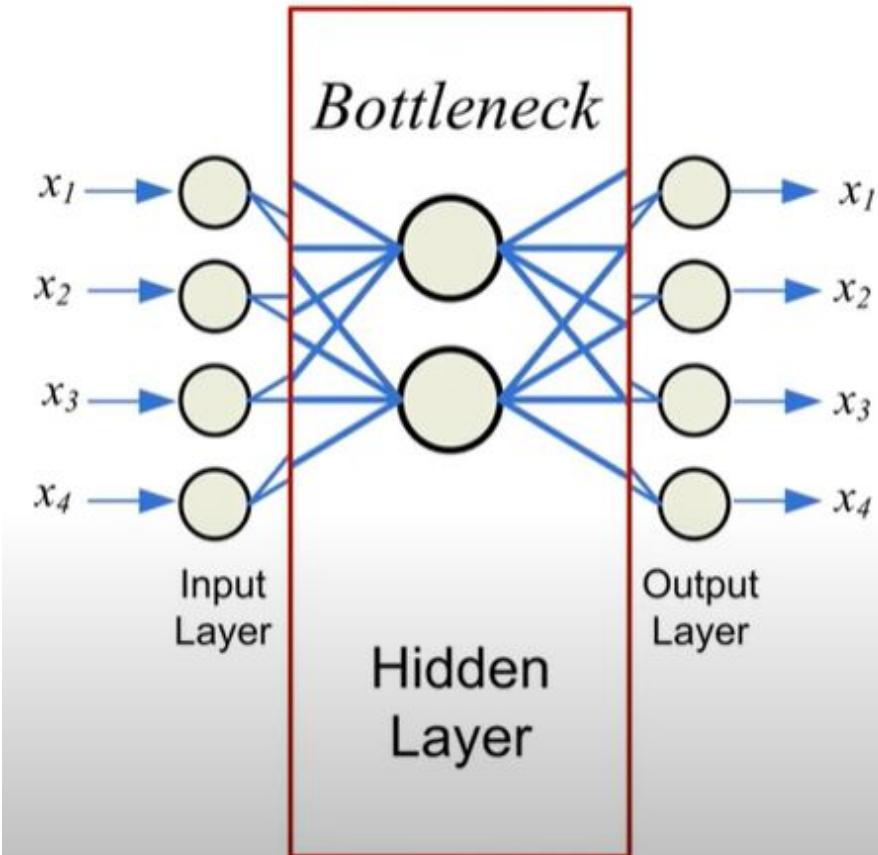
The output of the autoencoder will not be exactly the same as the input, it will be a close but degraded representation



# Hyperparameter – For Training AECD



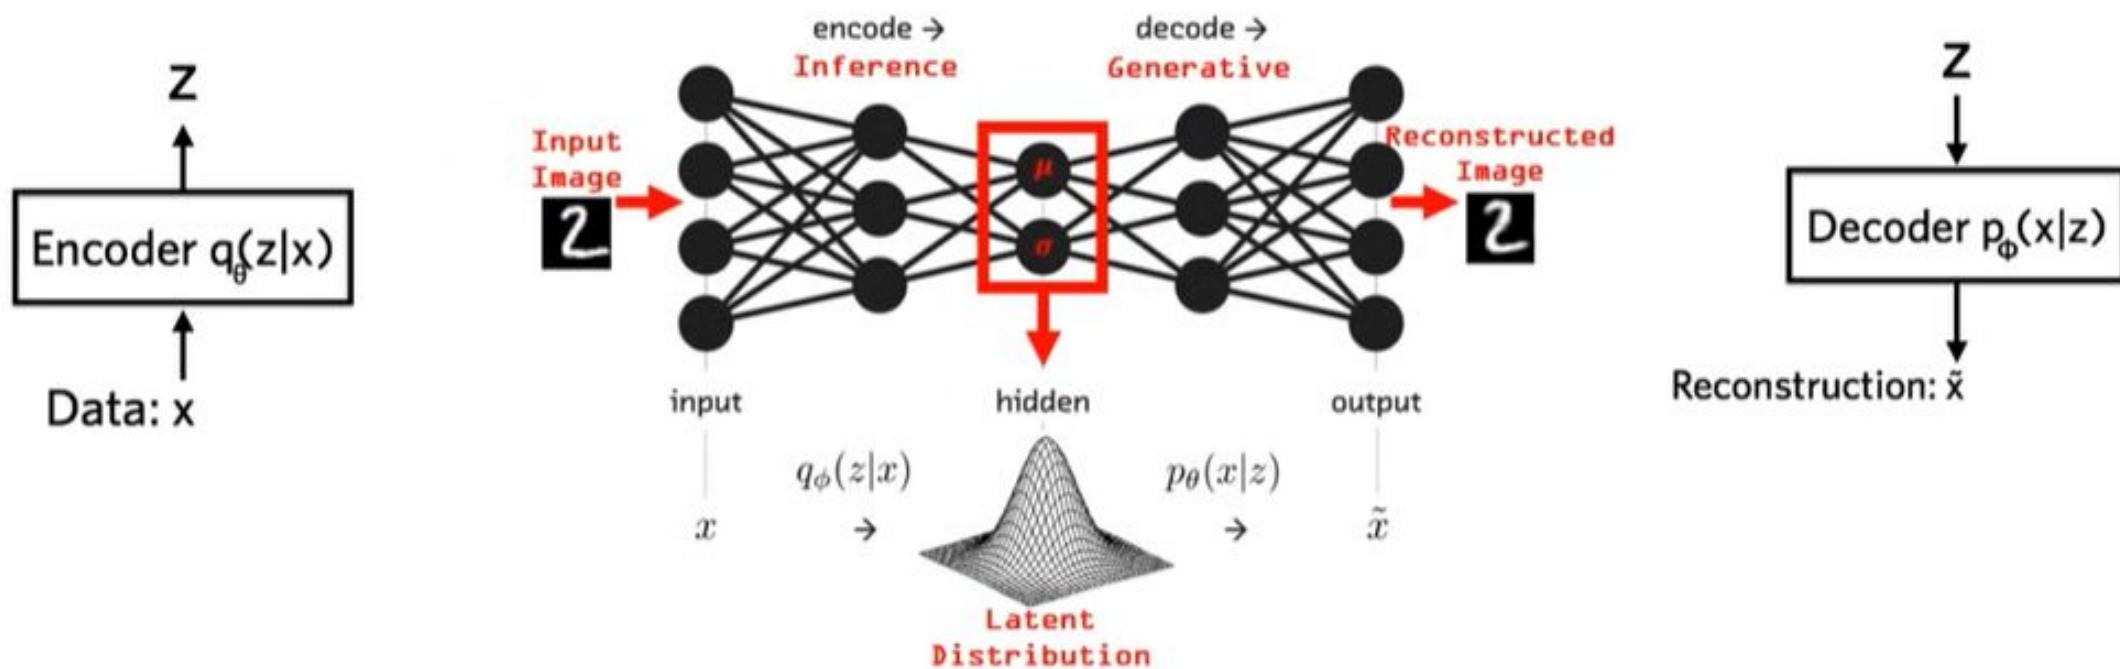
# Architecture of Autoencoders



Bottleneck approach is an approach to for deciding which aspects of observed data are relevant information and what aspects can be thrown away

- Compactness of representation, measured as the compressibility
- Representation retains about some behaviourally relevant variables

# Architecture of Autoencoders

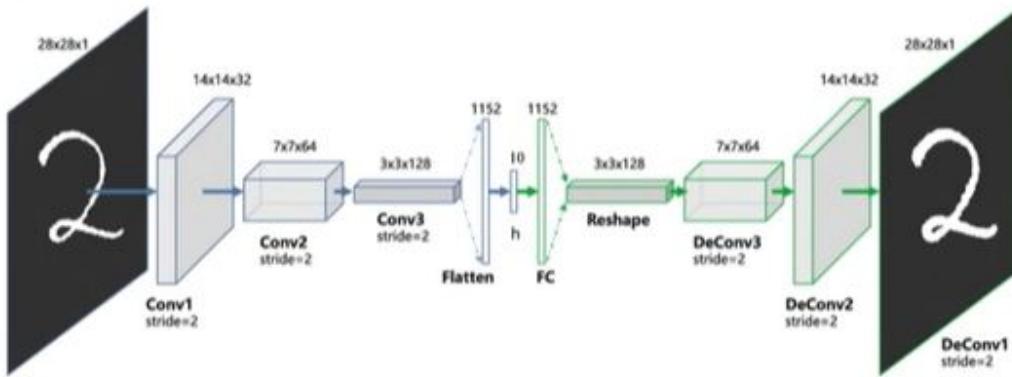


# Autoencoder: Types

- Convolutional AECD
- Latent AECD
- Sparse AECD
- Deep AECD
- Contractive AECD

# Convolutional AECD

## Convolution Autoencoders



Convolution Autoencoders use the convolution operator to learn to encode the input in a set of simple signals and then try to reconstruct the input from them.

$$f(t) * g(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau)d\tau$$

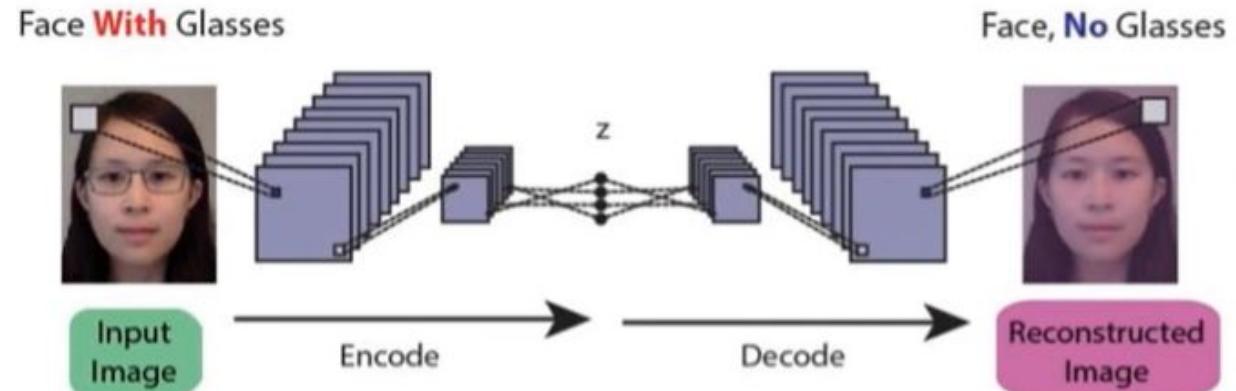
# Case studies of Convolutional AECD

## Use Case of CAE:

1 Image Reconstruction

2 Image Colorization

3 Advanced Applications

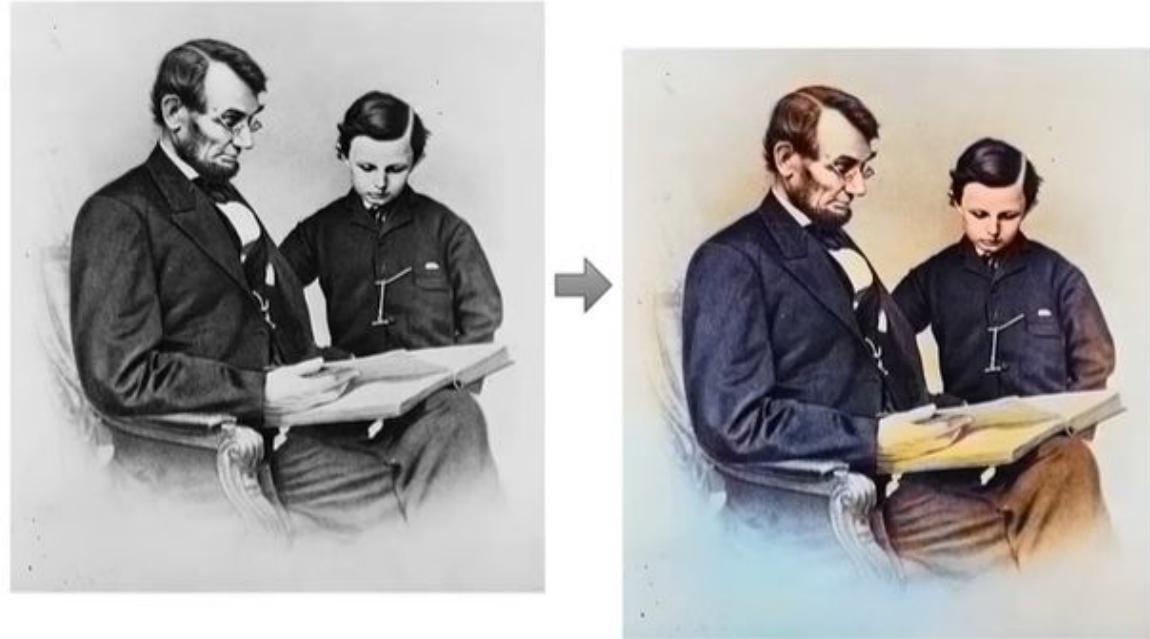
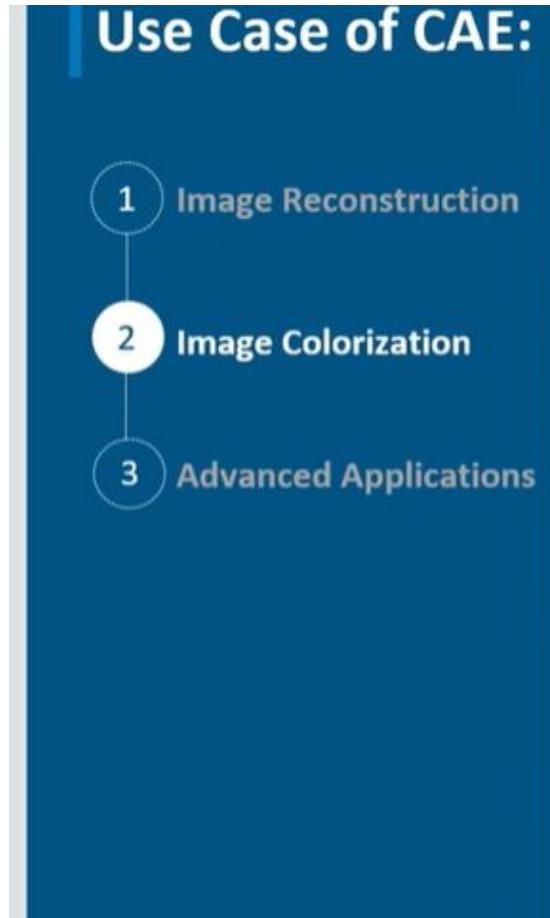


learns to **remove noise** or **reconstruct** missing parts

Noisy Version is converted to clean version

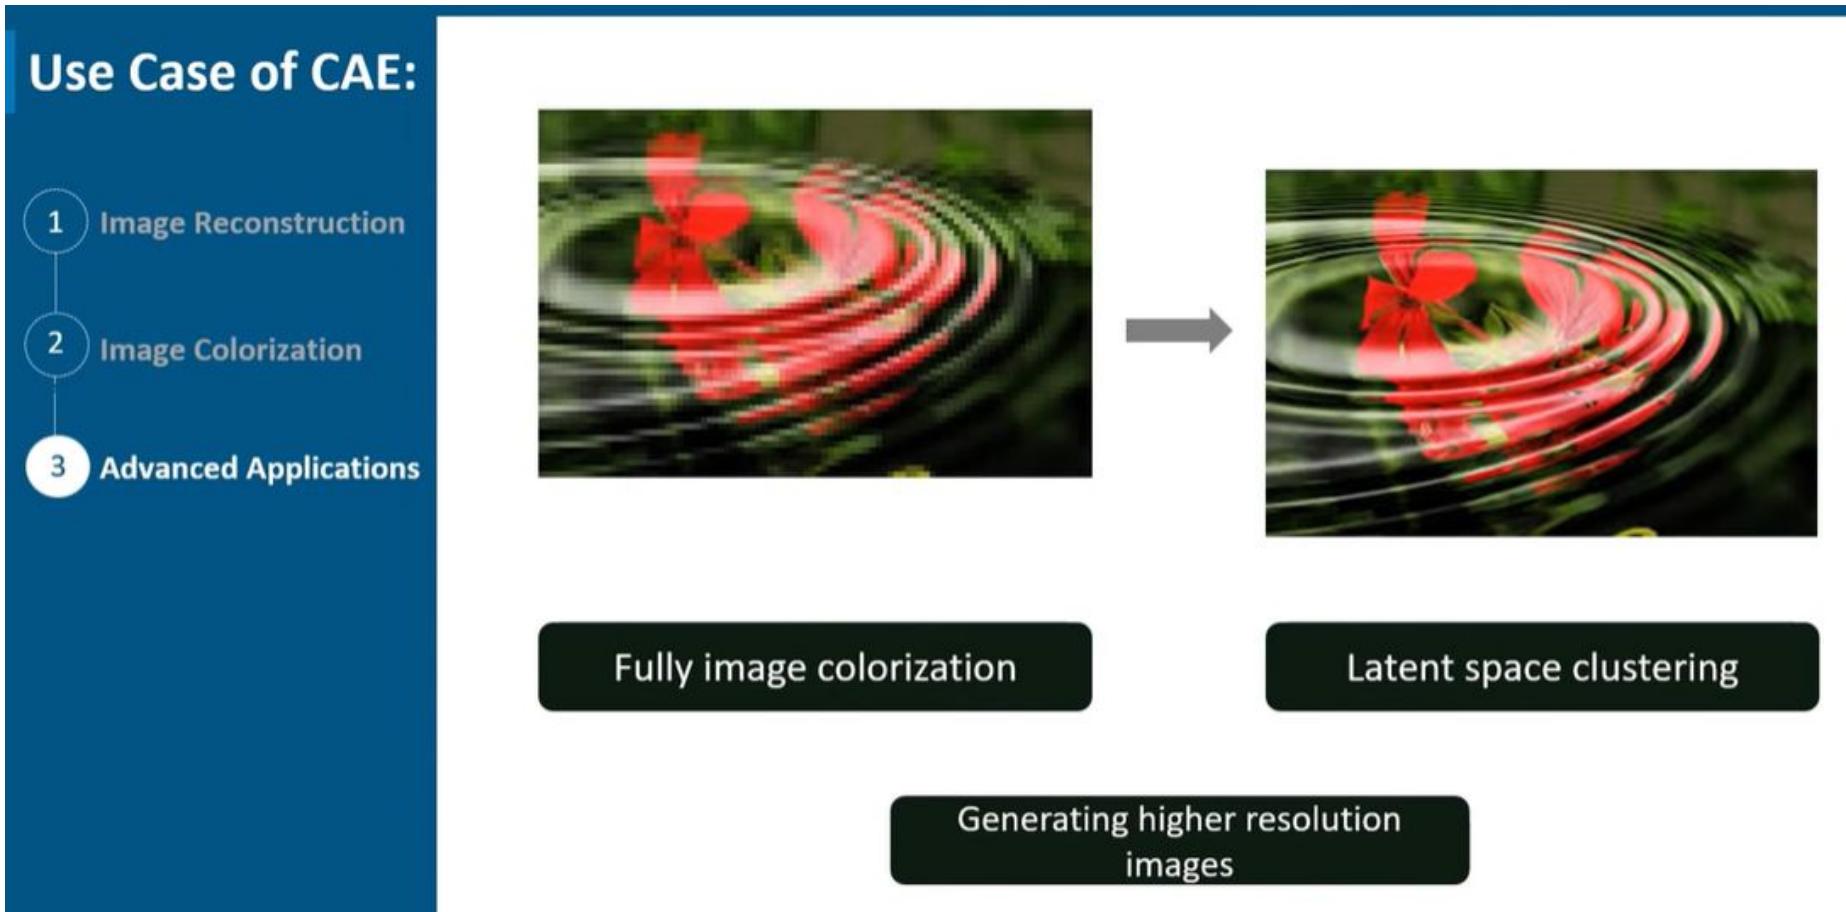
the network fills the gaps in the image

# Case studies of Convolutional AECD



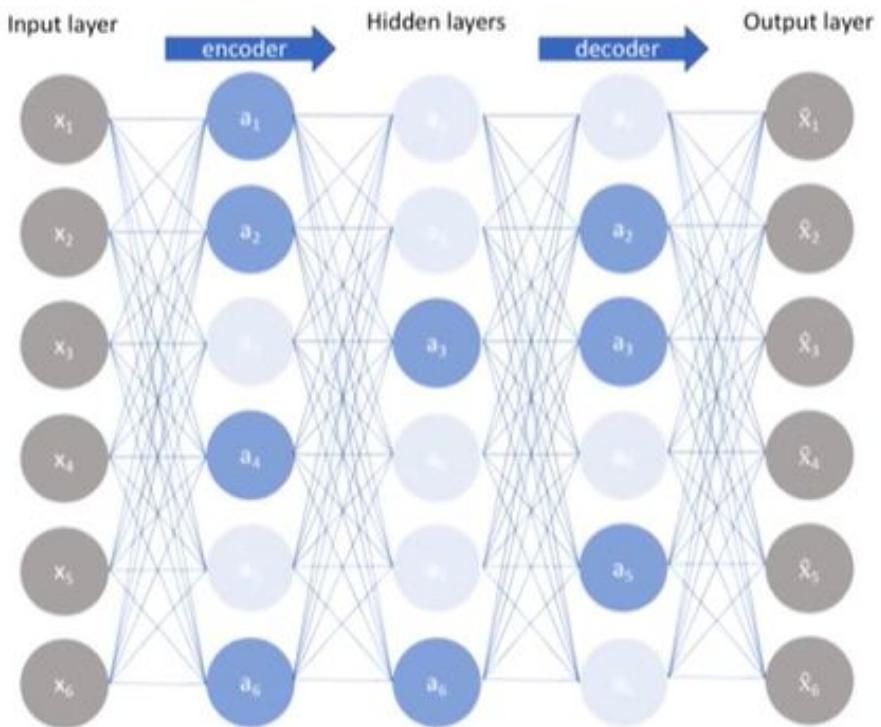
- maps **circles** and **squares** from an image to the same image but with Colors
- Purple is formed sometimes because of **blend** of colors, where network hesitates between circle or square.

# Case studies of Convolutional AECD



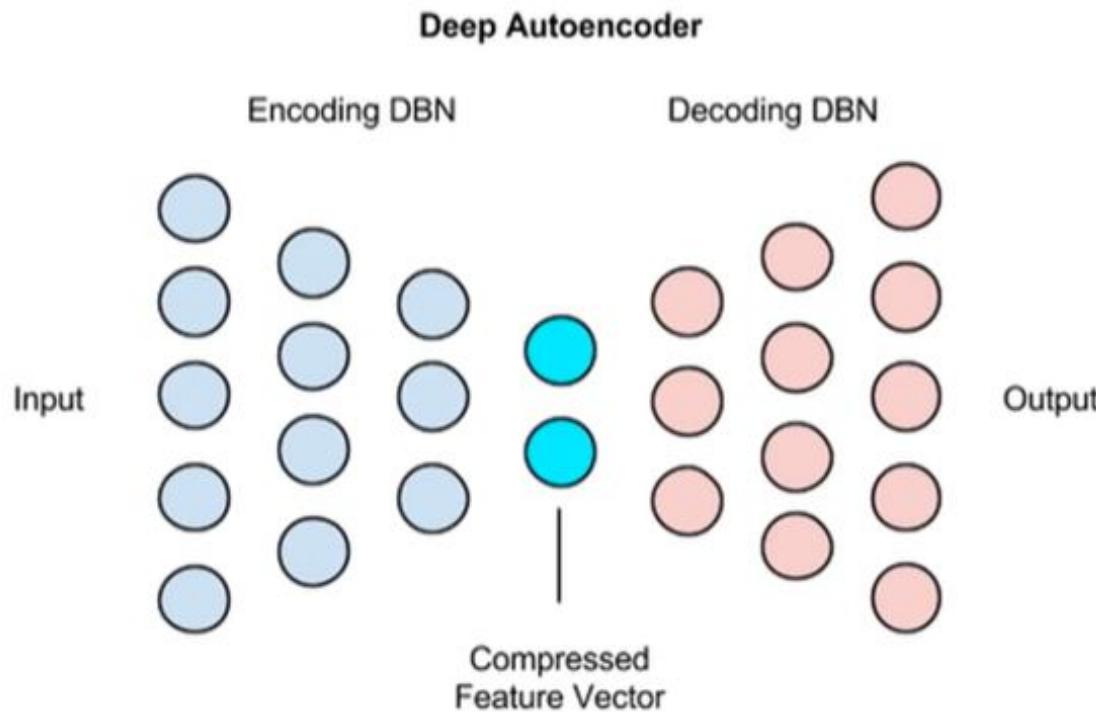
# Sparse Autoencoders

---



**Sparse autoencoders** offer us an alternative method for introducing an information bottleneck without requiring a reduction in the number of nodes at our hidden layers

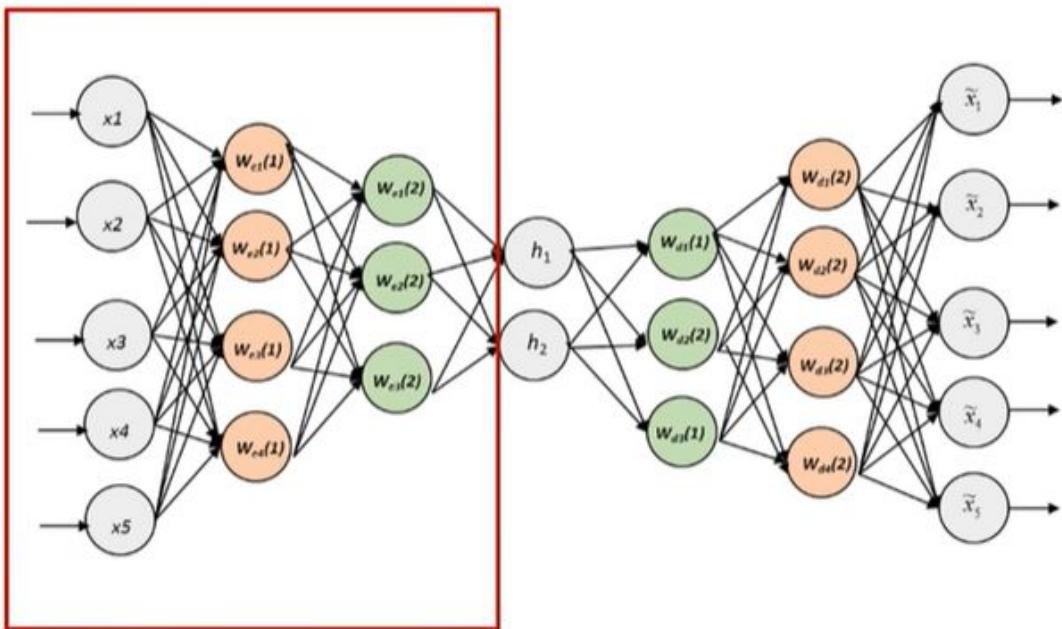
# Deep Autoencoders



A **deep autoencoder** is composed of two, symmetrical deep-belief networks-

- First four or five shallow layers representing the encoding half of the net
- second set of four or five layers that make up the decoding half.

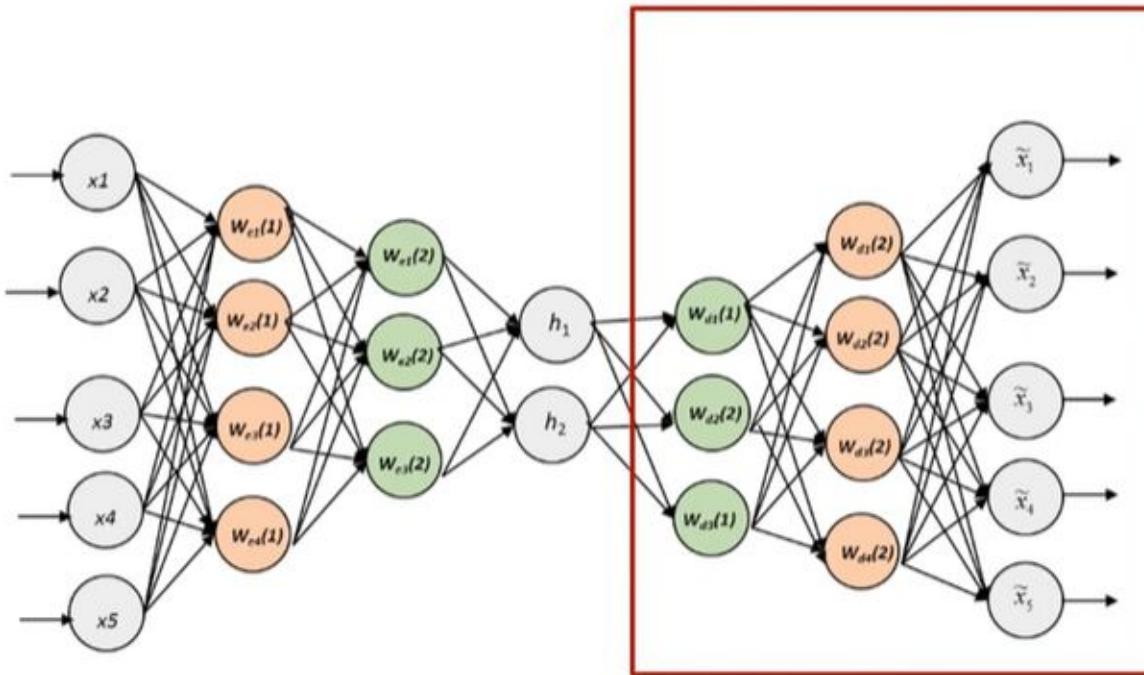
# Deep Autoencoders



The first layer of the Deep Autoencoder learns first-order features in the raw input such as edges in an image

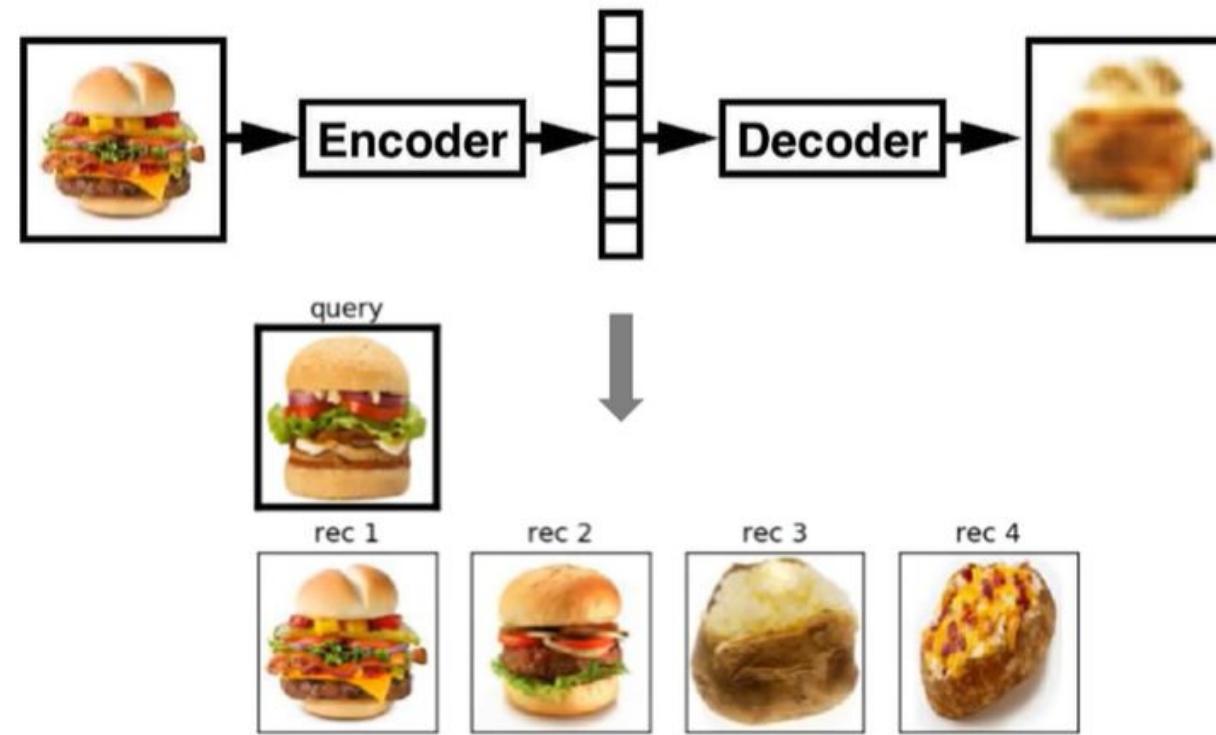
# Deep Autoencoders

The second layer learns second-order features corresponding to patterns in the appearance of first-order features



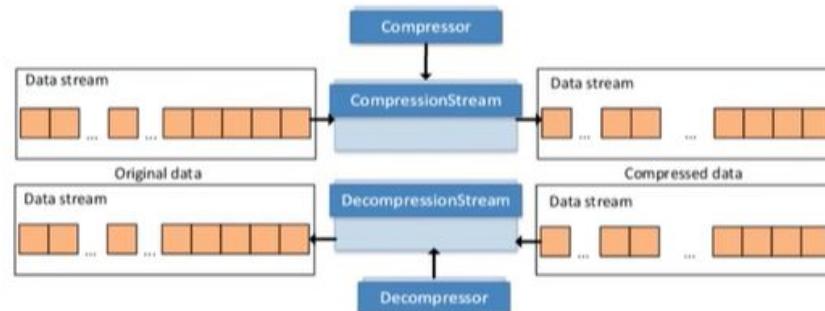
## Use Case: Deep Autoencoders

- 1 Image Search
- 2 Data Compression
- 3 Topic Modeling & Information Retrieval



## Use Case: Deep Autoencoders

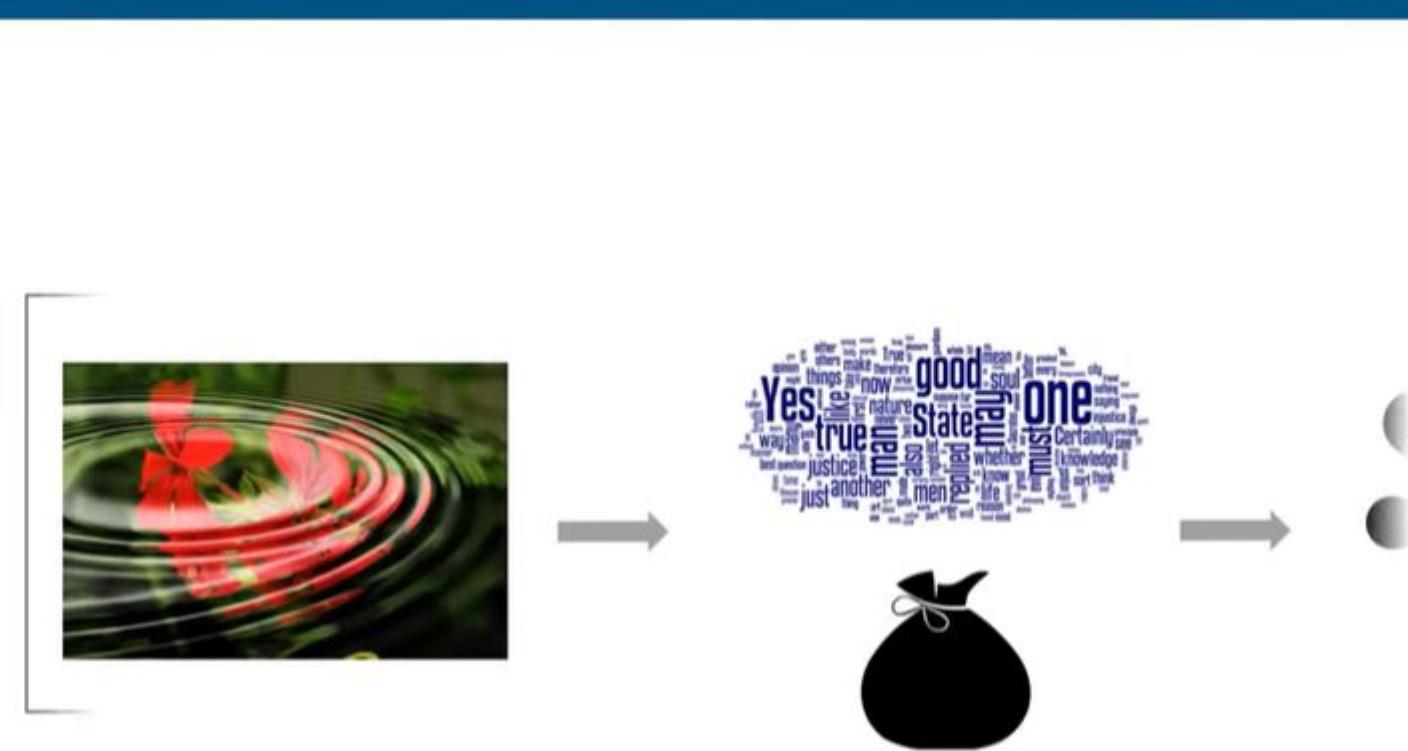
- 1 Image Search
- 2 Data Compression
- 3 Topic Modeling & Information Retrieval



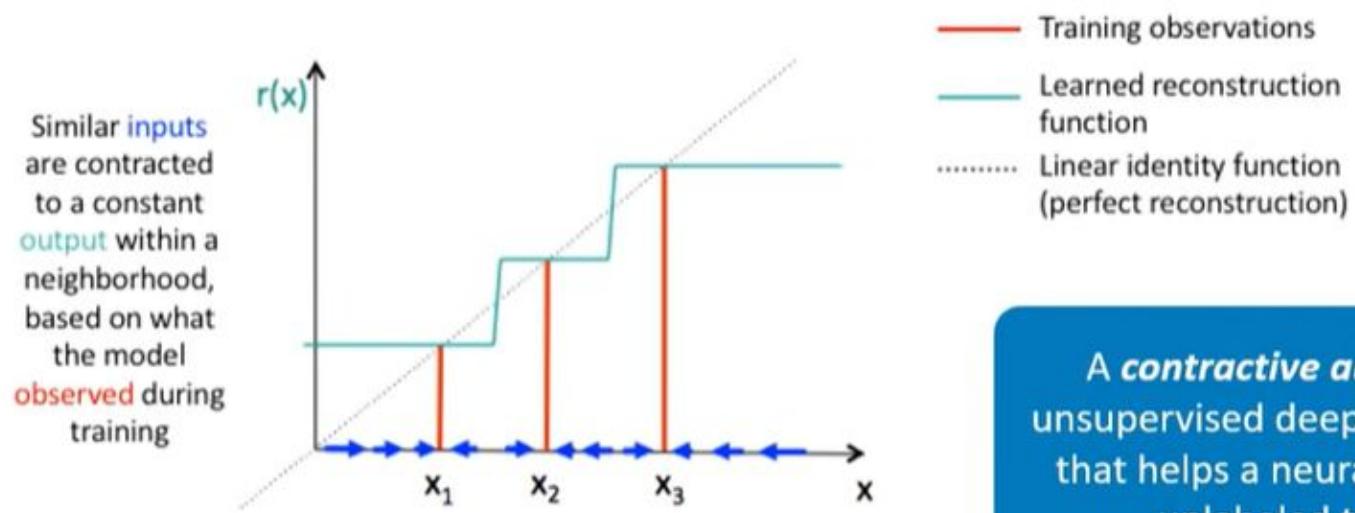
68.34K

## Use Case: Deep Autoencoders

- 1 Image Search
  - 2 Data Compression
  - 3 Topic Modeling & Information Retrieval



# Contractive Autoencoders



A **contractive autoencoder** is an unsupervised deep learning technique that helps a neural network encode unlabeled training data.

# References

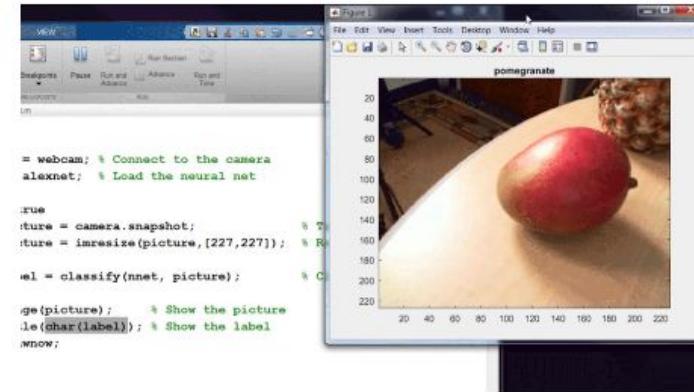
- [https://www.youtube.com/watch?v=nTt\\_ajul8NY&t=616s](https://www.youtube.com/watch?v=nTt_ajul8NY&t=616s)

# Applications of Deep Learning.

## Use AlexNet to Recognize Objects with Your Webcam

Use MATLAB, a simple webcam, and a deep neural network to identify objects in your surroundings.

[» Learn more](#)



## Example: Object Detection Using Deep Learning

In addition to [object recognition](#), which identifies a specific object in an image or video, deep learning can also be used for [object detection](#). [Object detection](#) algorithms like YOLO can recognize and locate the object in a scene, and can locate multiple objects within the image.



# Top Applications of Deep Learning Across Industries 2022

1. [Self Driving Cars](#)
2. [News Aggregation and Fraud News Detection](#)
3. [Natural Language Processing](#)
4. [Virtual Assistants](#)
5. [Entertainment](#)
6. [Visual Recognition](#)
7. [Fraud Detection](#)
8. [Healthcare](#)
9. [Personalisation's](#)
10. [Detecting Developmental Delay in Children](#)
11. [Colourisation of Black and White images](#)
12. [Adding sounds to silent movies](#)
13. [Automatic Machine Translation](#)
14. [Automatic Handwriting Generation](#)
15. [Automatic Game Playing](#)
16. [Language Translations](#)
17. [Pixel Restoration](#)
18. [Photo Descriptions](#)
19. [Demographic and Election Predictions](#)
20. [Deep Dreaming](#)

- 
- 1 Self Driving Cars  
2 Entertainment  
3 Visual Recognition  
4 Virtual Assistants  
5 Fraud Detection
- 6 Natural Language Processing  
7 News Aggregation and Fraud News Detection  
8 Detecting Developmental Delay in Children  
9 Colourisation of Black and White images  
10 Adding sounds to silent movies
- Healthcare 11  
Personalisations 12  
Automatic Machine Translation 13  
Automatic Handwriting Generation 14  
Demographic & Election Predictions 15
- 16 Automatic Game Playing  
17 Language Translations  
18 Pixel Restoration  
19 Photo Descriptions  
20 Deep Dreaming

# Self-learning Topics: Restricted Boltzmann Machine (RBM).

# References

- <https://www.mygreatlearning.com/blog/deep-learning-applications/>